

Open Source Monitoring for IBM MQ

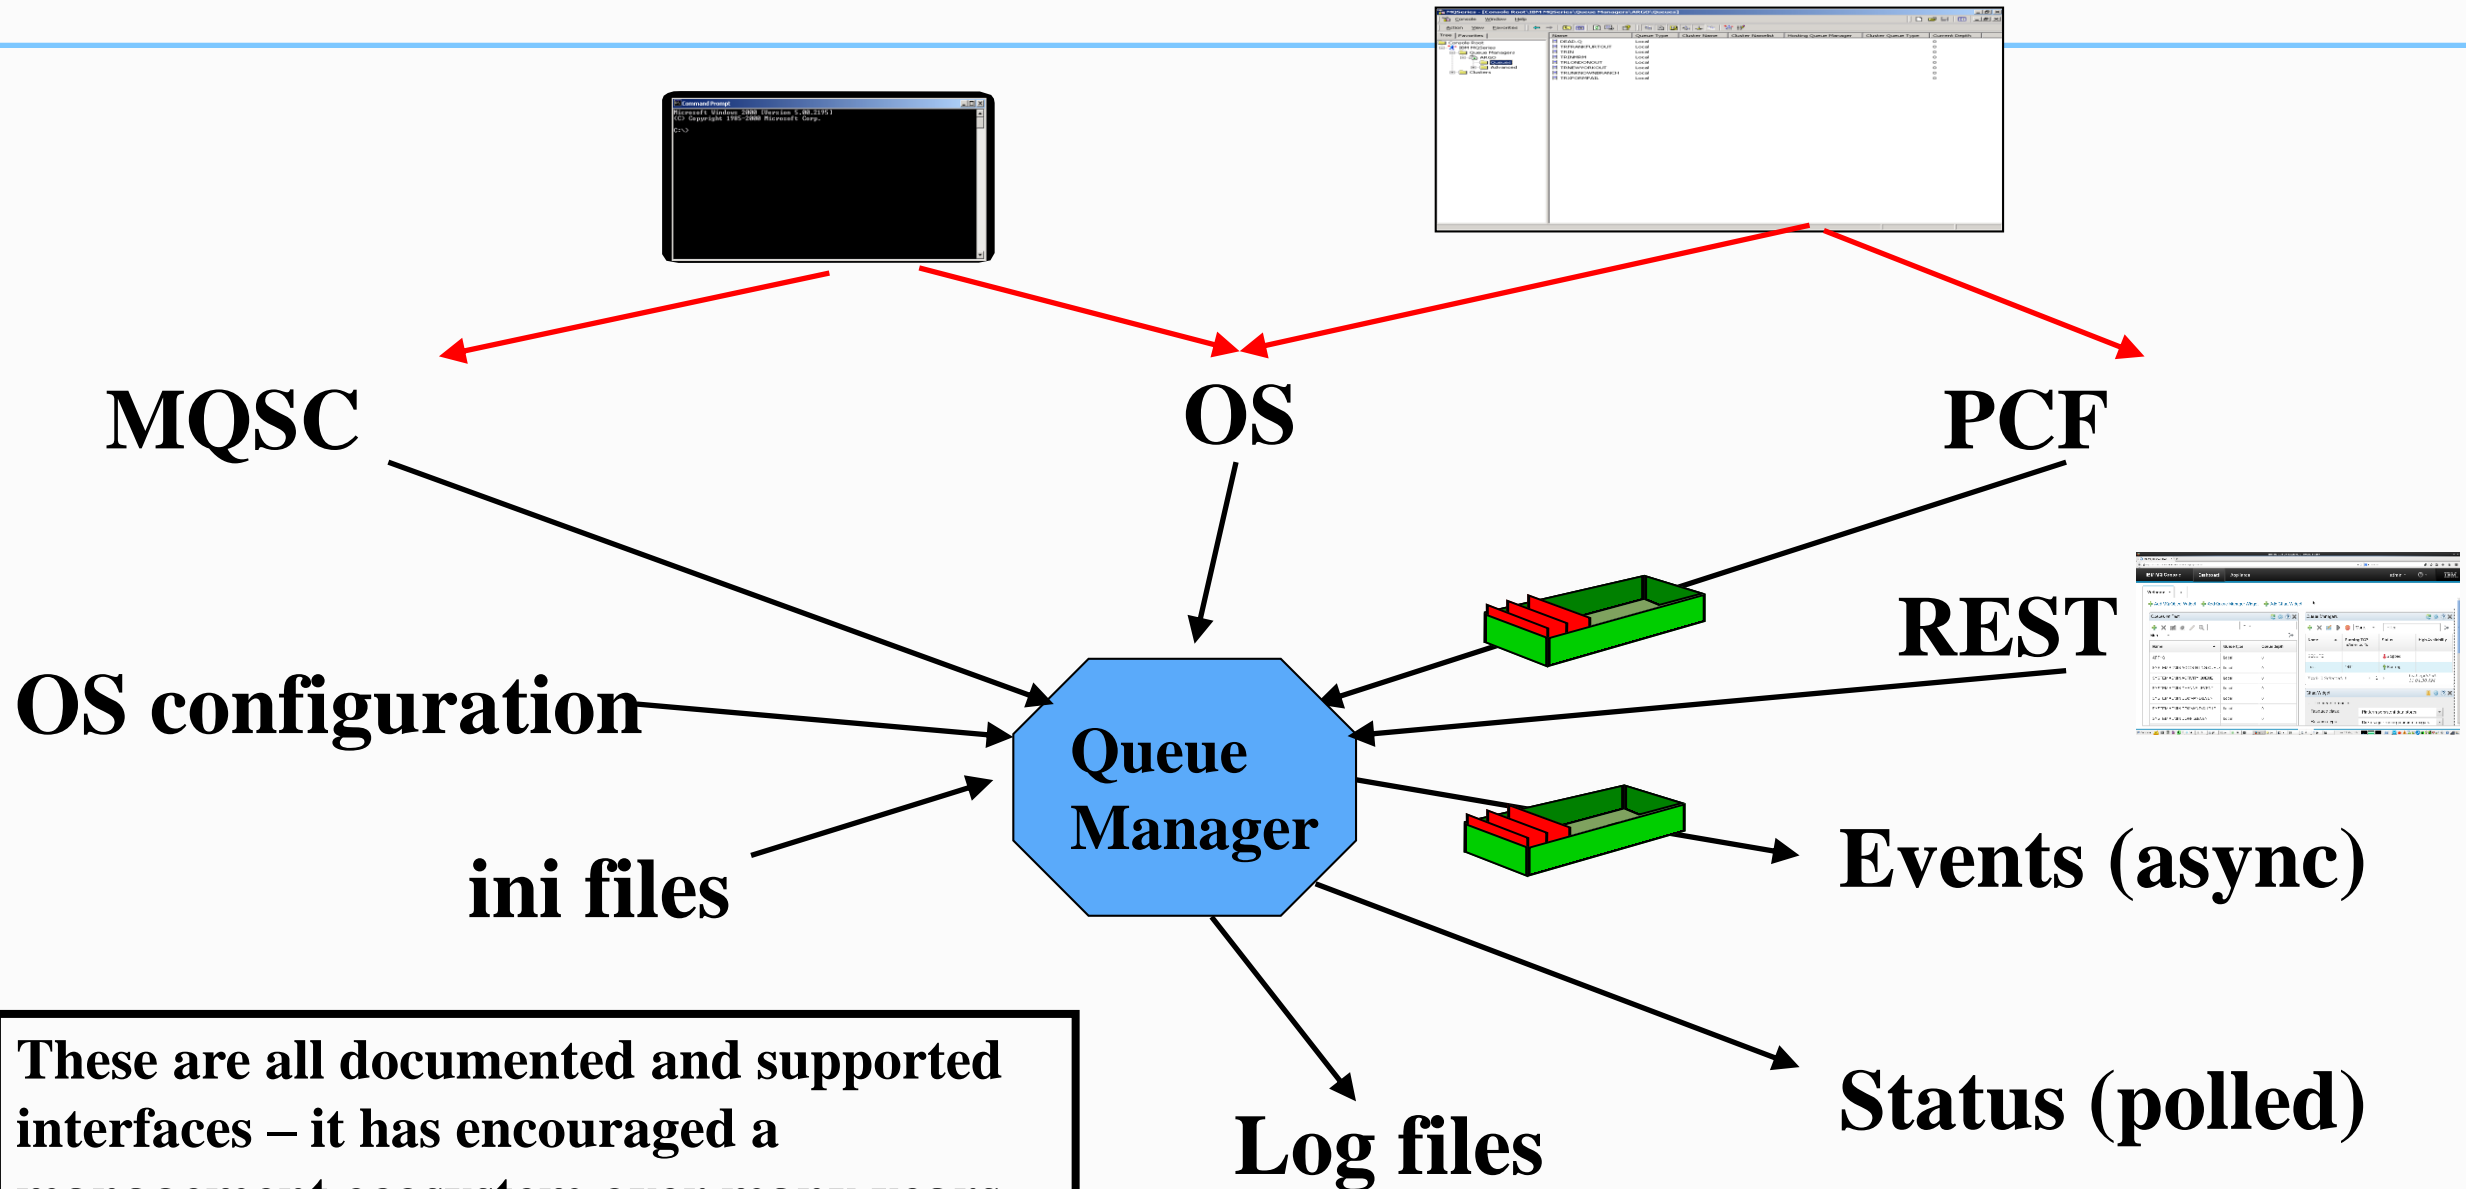
Mark Taylor

marke_taylor@uk.ibm.com

IBM Hursley



MQ Administration



These are all documented and supported interfaces – it has encouraged a management ecosystem over many years

IBM MQ - MQSC

Old Example (1995): AIX smit panels

- Command line interface
- V8 enhanced runmqsc
 - Make it world-executable
 - Enable direct client-connection
- MQSC intended for human consumption
 - Parsable by eye, less easy in programs
 - For example, **DESCR('This is 'a' description with quote & paren(')**
 - No guaranteed ordering in runmqsc, two-column output
- Despite awkwardness, basis for many script-based admin tools
 - echo "DISPLAY Q(X) IPPROCS" | runmqsc QM1
- Same commands – different front-end (CSQUTIL) – for z/OS

IBM MQ - PCF

- A "self-describing" MQ message used for administrative operations
- Your programs can send commands and get responses using PCF
 - Equivalent to "DISPLAY QSTATUS" or "ALTER CHANNEL"
- MQ emits events in PCF format
 - "Queue is getting full"
- PCF intended for programs – usually C or Java
 - Can tell exactly what the parameter is for, its length and value
 - But cannot easily be scripted
- Approximately one-one mapping between MQSC commands and PCF
- Remember that PCF invented before formats like JSON or XML
 - And there are many MQ apps that are built on PCF

An event message

**** Message length - 300 of 300 bytes ***

```
00000000: 0000 0007 0000 0024 0000 0003 0000 0063 '.....$......c'
00000010: 0000 0001 0000 0001 0000 0000 0000 096C '.....1'
00000020: 0000 0002 0000 0014 0000 0010 0000 1F41 '.....A'
00000030: 0000 0004 0000 0004 0000 0020 0000 0BE5 '.....å'
00000040: 0000 0333 0000 000C 6D65 7461 796C 6F72 '...3...metaylor'
00000050: 2020 2020 0000 0003 0000 0010 0000 03F3 '.....ó'
00000060: 0000 0001 0000 0004 0000 0044 0000 0BE7 '.....D...ç'
00000070: 0000 0333 0000 0030 5638 3030 335F 4120 '...3...0V8003_A '
00000080: 2020 2020 2020 2020 2020 2020 2020 2020 ' '
00000090: 2020 2020 2020 2020 2020 2020 2020 2020 ' '
000000A0: 2020 2020 2020 2020 0000 0003 0000 0010 '.....'
000000B0: 0000 03FD 0000 005A 0000 0014 0000 0010 '...ý...Z.....'
000000C0: 0000 1F42 0000 0004 0000 0004 0000 0018 '...B.....'
000000D0: 0000 0BFB 0000 0000 0000 0001 5800 0000 '...û.....X...'
000000E0: 0000 0003 0000 0010 0000 03F8 0000 0001 '.....ø.....'
000000F0: 0000 0006 0000 0024 0000 0BF9 0000 0000 '.....$....ù....'
00000100: 0000 0001 0000 0008 6D65 7461 796C 6F72 '.....metaylor'
00000110: 0000 0000 0000 0005 0000 0018 0000 045C '.....\'
00000120: 0000 0002 0000 000B 0000 0009 '.....'
```

An event message

**** Message length - 300 of 300 bytes ***

00000000:	0000	0007	0000	0024	0000	0003	0000	0063	'.....\$......c'
00000010:	0000	0001	0000	0001	0000	0000	0000	096C	'.....1'
00000020:	0000	0002	0000	0014	0000	0010	0000	1F41	'.....A'
00000030:	0000	0004	0000	0004	0000	0020	0000	0BE5	'.....å'
00000040:	0000	0333	0000	000C	6D65	7461	796C	6F72	'...3...metaylor'
00000050:	2020	2020	0000	0003	0000	0010	0000	03F3	'.....ó'
00000060:	0000	0001	0000	0004	0000	0044	0000	0BE7	'.....D...ç'
00000070:	0000	0333	0000	0030	5638	3030	335F	4120	'...3...0V8003_A '
00000080:	2020	2020	2020	2020	2020	2020	2020	2020	'.....'
00000090:	2020	2020	2020	2020	2020	2020	2020	2020	'.....'
000000A0:	2020	2020	2020	2020	0000	0003	0000	0010	'.....'
000000B0:	0000	03FD	0000	005A	0000	0014	0000	0010	'...ý...Z.....'
000000C0:	0000	1F42	0000	0004	TYPE (cfst) LEN (24)				'...B.....'
000000D0:	PARM (MQCA...)		CCSID (0)		LEN (1)		DATA		'...û.....X...'
000000E0:	0000	0003	0000	0010	0000	03F8	0000	0001	'.....ø.....'
000000F0:	0000	0006	0000	0024	0000	0BF9	0000	0000	'.....\$....ù.....'
00000100:	0000	0001	0000	0008	6D65	7461	796C	6F72	'.....metaylor'
00000110:	0000	0000	0000	0005	0000	0018	0000	045C	'.....\'
00000120:	0000	0002	0000	000B	0000	0009			'.....'

Event formatting C sample in V8.0.0.4

- No sample previously shipped to format all "standard" events
 - Authorisation, queue full, service interval, command/config etc
 - Other product samples **are** available for acct/stats, activity reports
 - Several SupportPacs but product only has out-of-date source code in the KC
- The **amqsevt** program formats events into readable English-ish text
 - Option to stay with full MQI constant name instead of making it look nice
 - Uses MQCB to read from multiple event queues. No polling required
 - Can connect as client to any remote queue manager including z/OS
 - Source code included
- Includes C header file to help convert MQI numbers to strings
 - Similar to Java MQConstants.lookup() capability for all sets of constants

```
printf("Error is %s\n",MQRC_STR(2035));
```

An event message decoded

Event Type : Command Event
Reason : Command MQSC
Event created : 2015/06/03 13:28:20.51 GMT
Correlation ID : 414D512056383030335F412020202020556F00F120001E05

COMMAND CONTEXT

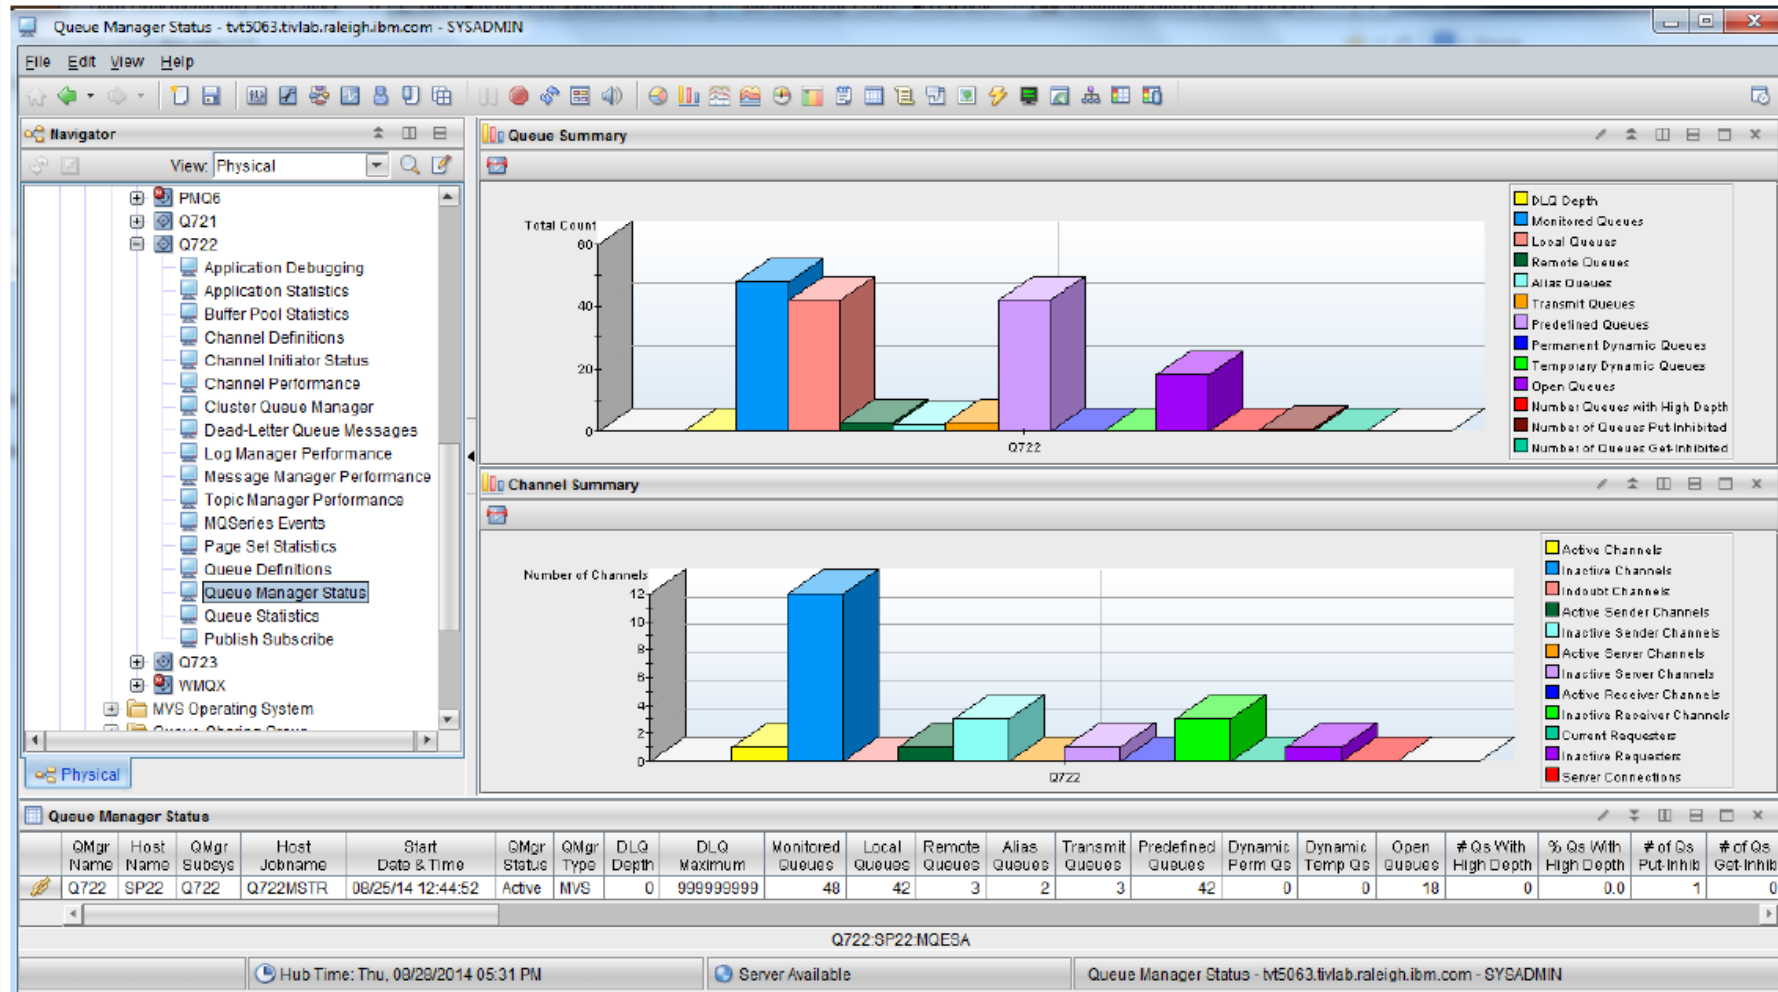
Event User Id : metaylor
Event Origin : Console
Event Queue Mgr : V8003_A
Command : Set Auth Rec

COMMAND DATA

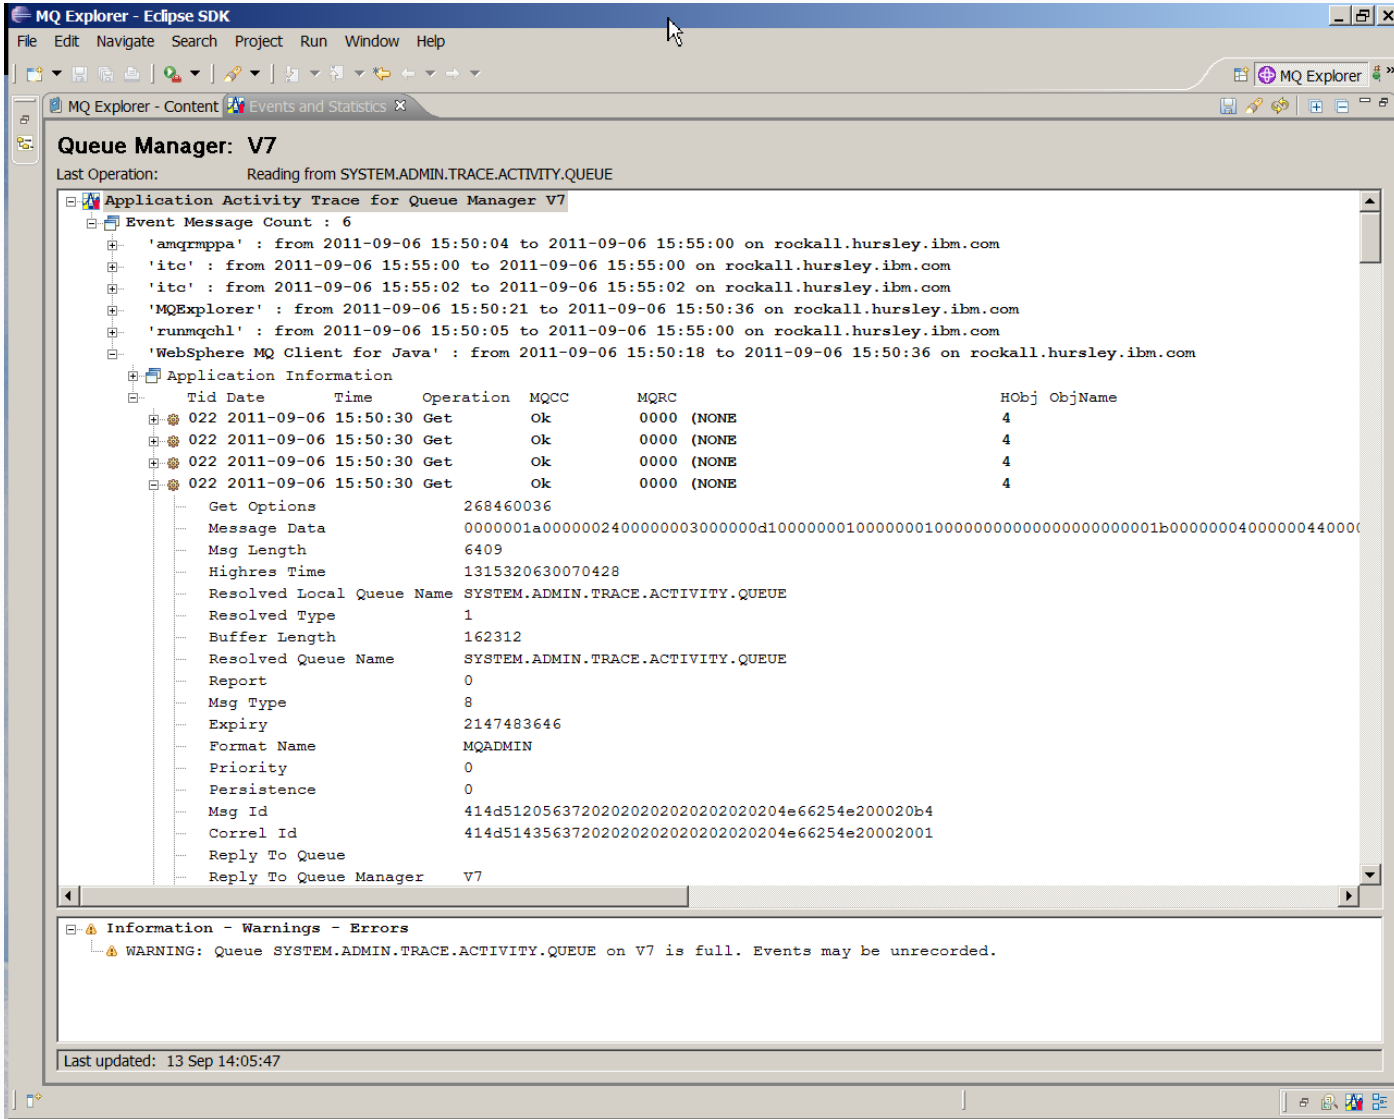
Auth Profile Name : X
Object Type : Queue
Principal Entity Names: metaylor
Auth Add Auths : Output
: Input

Third-party solutions

- Many vendor products – this screenshot from ITCAM/Omegamon

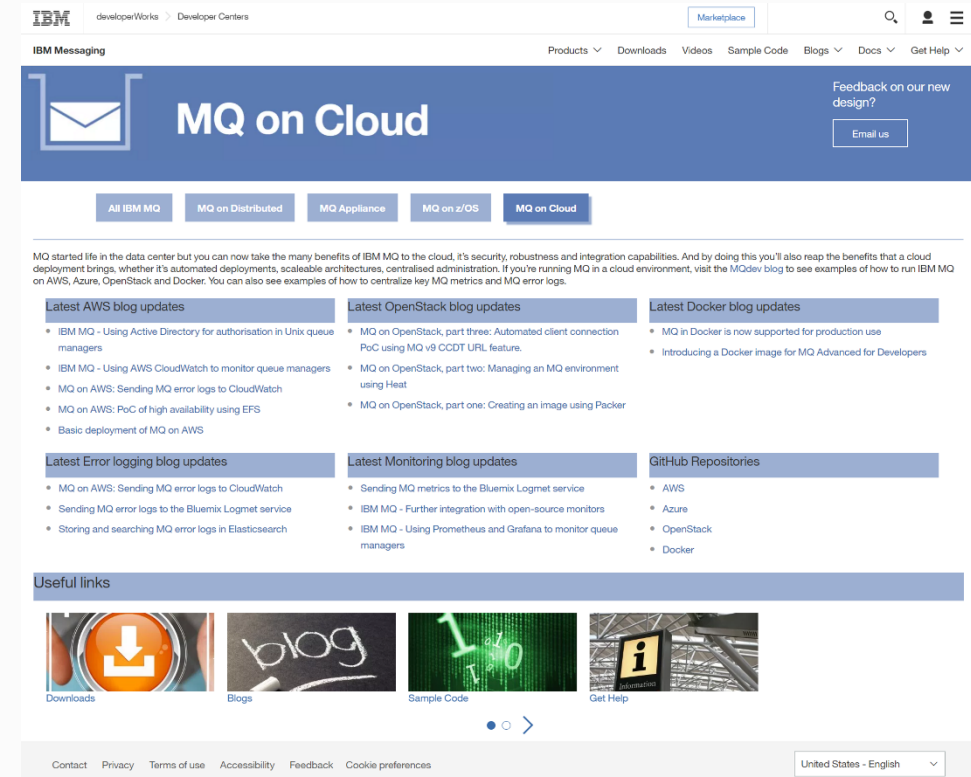


Application Activity inside MQ Explorer using MSOP



Many people now using different tools

- Because they are using those tools for other products
- And because MQ is being used in more environments
- Therefore MQ has to be able to be integrated with them
- You are unlikely to buy Tivoli if other parts of your infrastructure are being monitored via Grafana



<https://developer.ibm.com/messaging/mq-on-cloud/>



Decided to demonstrate MQ monitoring integration

- Using the V9 resource statistics data
- Feeding a variety of monitoring tools
- And doing it in public – Github, blog articles etc
 - See github.com/ibm-messaging/mq-golang
 - Video at youtube.com/watch?v=Pi_jHCiqTgU
- Other integration aspects – availability, security, deployment – also demonstrated

System Monitoring with V9

- More statistics available via a pub/sub model
- Includes CPU and Disk usage
 - As well as many MQ statistics
 - Not full replacement for accounting/statistics events but many key values
- Subscribe to meta-topic to learn which classes of statistics are available
 - **\$SYS/MQ/INFO/QMGR/<qmgr>/Monitor/METADATA/CLASSES**
 - Then subscribe to specific topics
 - See amqsrua sample program
- Distributed platforms only
- User applications can generate their own monitoring data in this style
 - The MQ Bridge to Salesforce contributes statistics

System Monitoring Example

```
$ amqsrua -m V9000_A
```

```
CPU : Platform central processing units
```

```
DISK : Platform persistent data stores
```

```
STATMQI : API usage statistics
```

```
STATQ : API per-queue usage statistics
```

```
Enter Class selection
```

```
==> CPU
```

```
SystemSummary : CPU performance - platform wide
```

```
QMgrSummary : CPU performance - running queue manager
```

```
Enter Type selection
```

```
==> SystemSummary
```

```
Publication received PutDate:20160411 PutTime:10465573
```

```
User CPU time percentage 0.01%
```

```
System CPU time percentage 1.30%
```

```
CPU load - one minute average 8.00
```

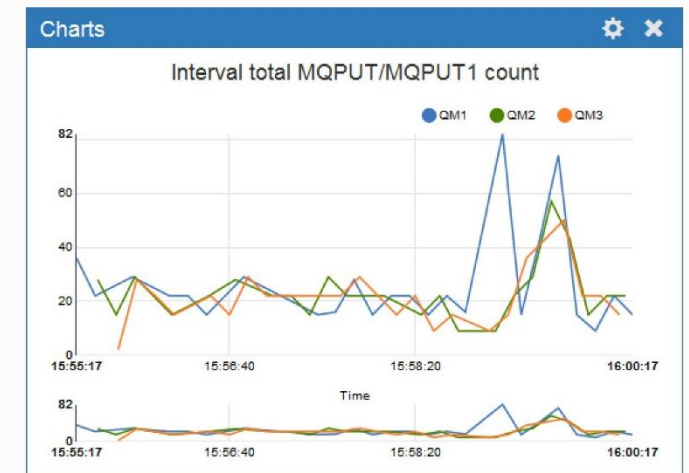
```
CPU load - five minute average 7.50
```

```
CPU load - fifteen minute average 7.30
```

```
RAM free percentage 2.02%
```

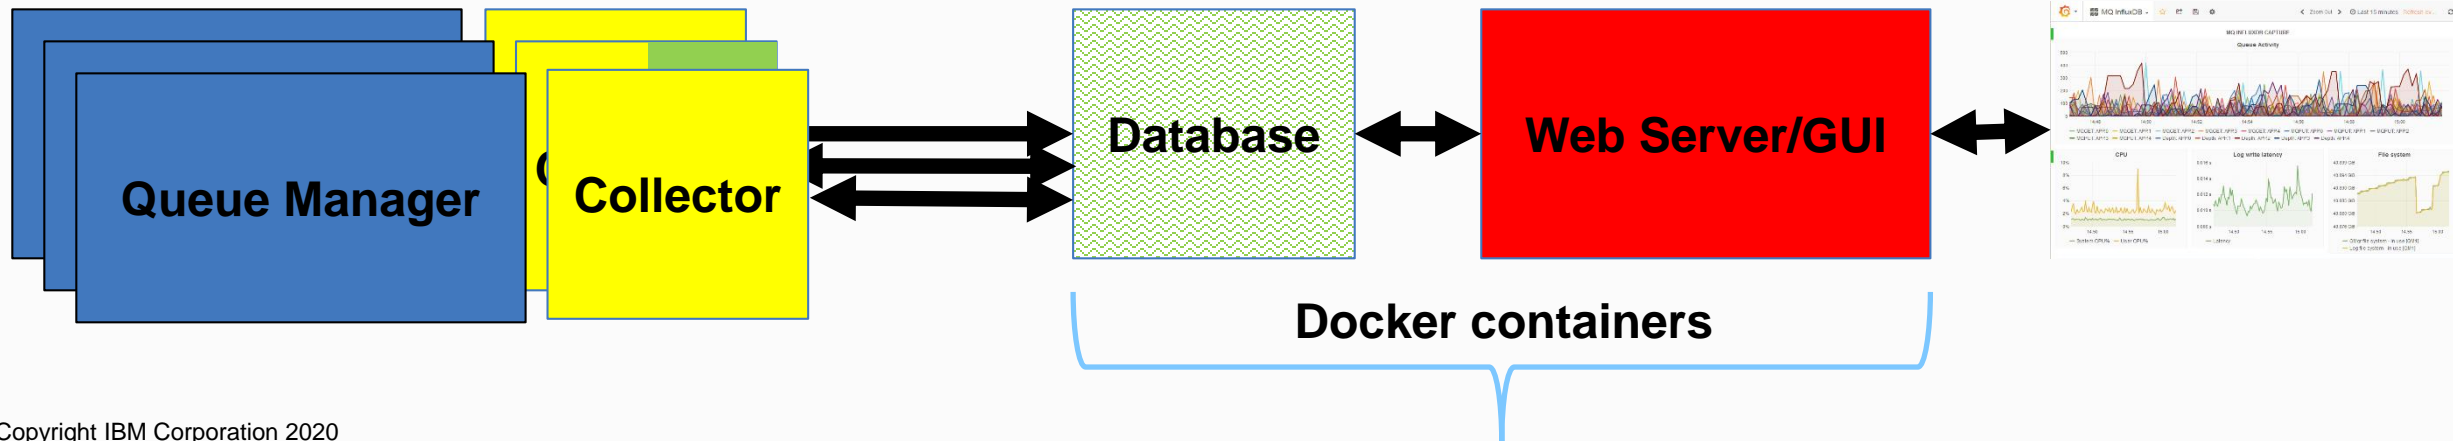
```
RAM total bytes 8192MB
```

This capability underpins
the charting in the MQ
Console UI



Monitoring Architecture

- Architecture is split – database and user interface
 - The database is usually a "time-series" DB, not traditional SQL
 - Designed and optimised for {timestamp, metric, value} storage and queries
- These databases include Prometheus, InfluxDB, OpenTSDB
- Collection architecture may have intermediate layers – collectd



Started with Prometheus

- Seemed to be one of the most popular
- Which does have its own limited GUI
- Model is "pull" – calls a collector program at intervals via http
 - Most other DBs are "push" where collector sends to DB at interval
- Standard API for getting data to Prometheus is in Go
 - And we had no Go API for MQ ...

The Go API for MQ

- So first off, I had to create a new language binding
 - Based on full MQI rather than a "simplified" version
 - But not all function implemented
 - Trying to make it look natural to Go programmers

```
if err == nil {  
    putmqmd := ibmmq.NewMQMD()  
    pmo      := ibmmq.NewMQPMO()  
    pmo.Options = ibmmq.MQPMO_SYNCPOINT | ibmmq.MQPMO_NEW_MSG_ID  
    putmqmd.Format = "MQSTR"  
    msgData := "Hello from Go"  
    buffer := []byte(msgData)  
    err = qObject.Put(putmqmd, pmo, buffer)  
    if err != nil {  
        fmt.Println(err)  
    }  
}
```

Collector configurations

- Collector subscribes to all data for qmgr (cpu, disk etc) and nominated queues
 - Command line parameters name the queues with wildcards
- Started via MQ Service definition and shell script
- Can connect as client to remote queue managers including MQ appliance
 - Any system that supports the resource statistics
 - One collector instance per queue manager

```
/usr/local/bin/mqgo/mq_prometheus -ibmmq.queueManager=QM1  
-ibmmq.monitoredQueues=APP.*,MYQ.*  
-ibmmq.httpListenPort=9157  
-log.level=error
```

Grafana

- Although Prometheus has a GUI it is not very sophisticated
- Instead, prefer to use Grafana as visualisation tool
 - Supports many different backend databases
 - Understands the metric names, query capabilities etc of each

Add data source

Config Dashboards

Name My data source name ⓘ Default ☐

Type Prometheus

Http settings

Url http://localhost:9090 ⓘ

Access proxy ⓘ

Http Auth Basic Auth ☐ With Credentials ☐

Add data source

Name My data source name ⓘ Default ☐

Type CloudWatch

CloudWatch details

Credentials profile name default ⓘ

Default Region ⓘ

Custom Metrics namespace Namespace1, Namespace2 ⓘ

Assume Role ARN arn:aws:iam:* ⓘ

Add data source

Config Dashboards

Name My data source name ⓘ Default ☐

Type Graphite

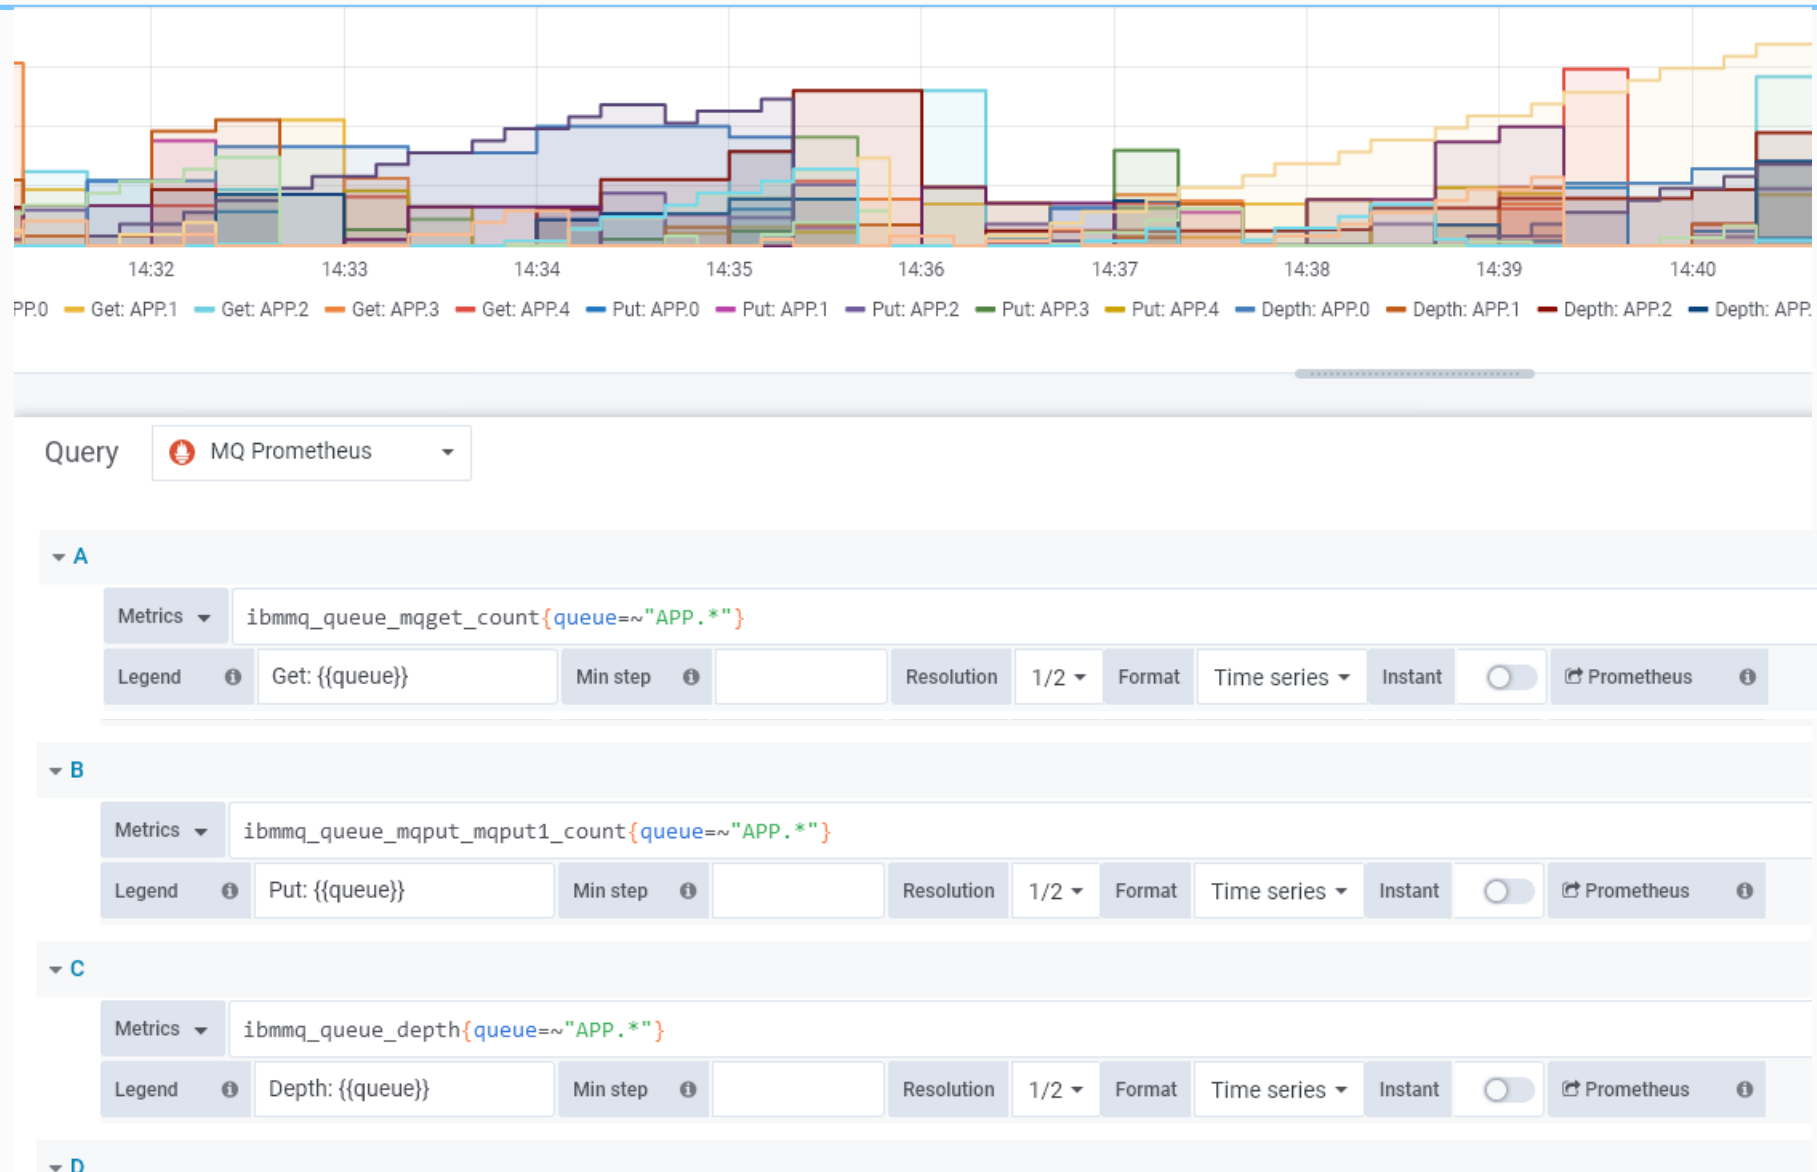
Http settings

Url http://localhost:8080 ⓘ

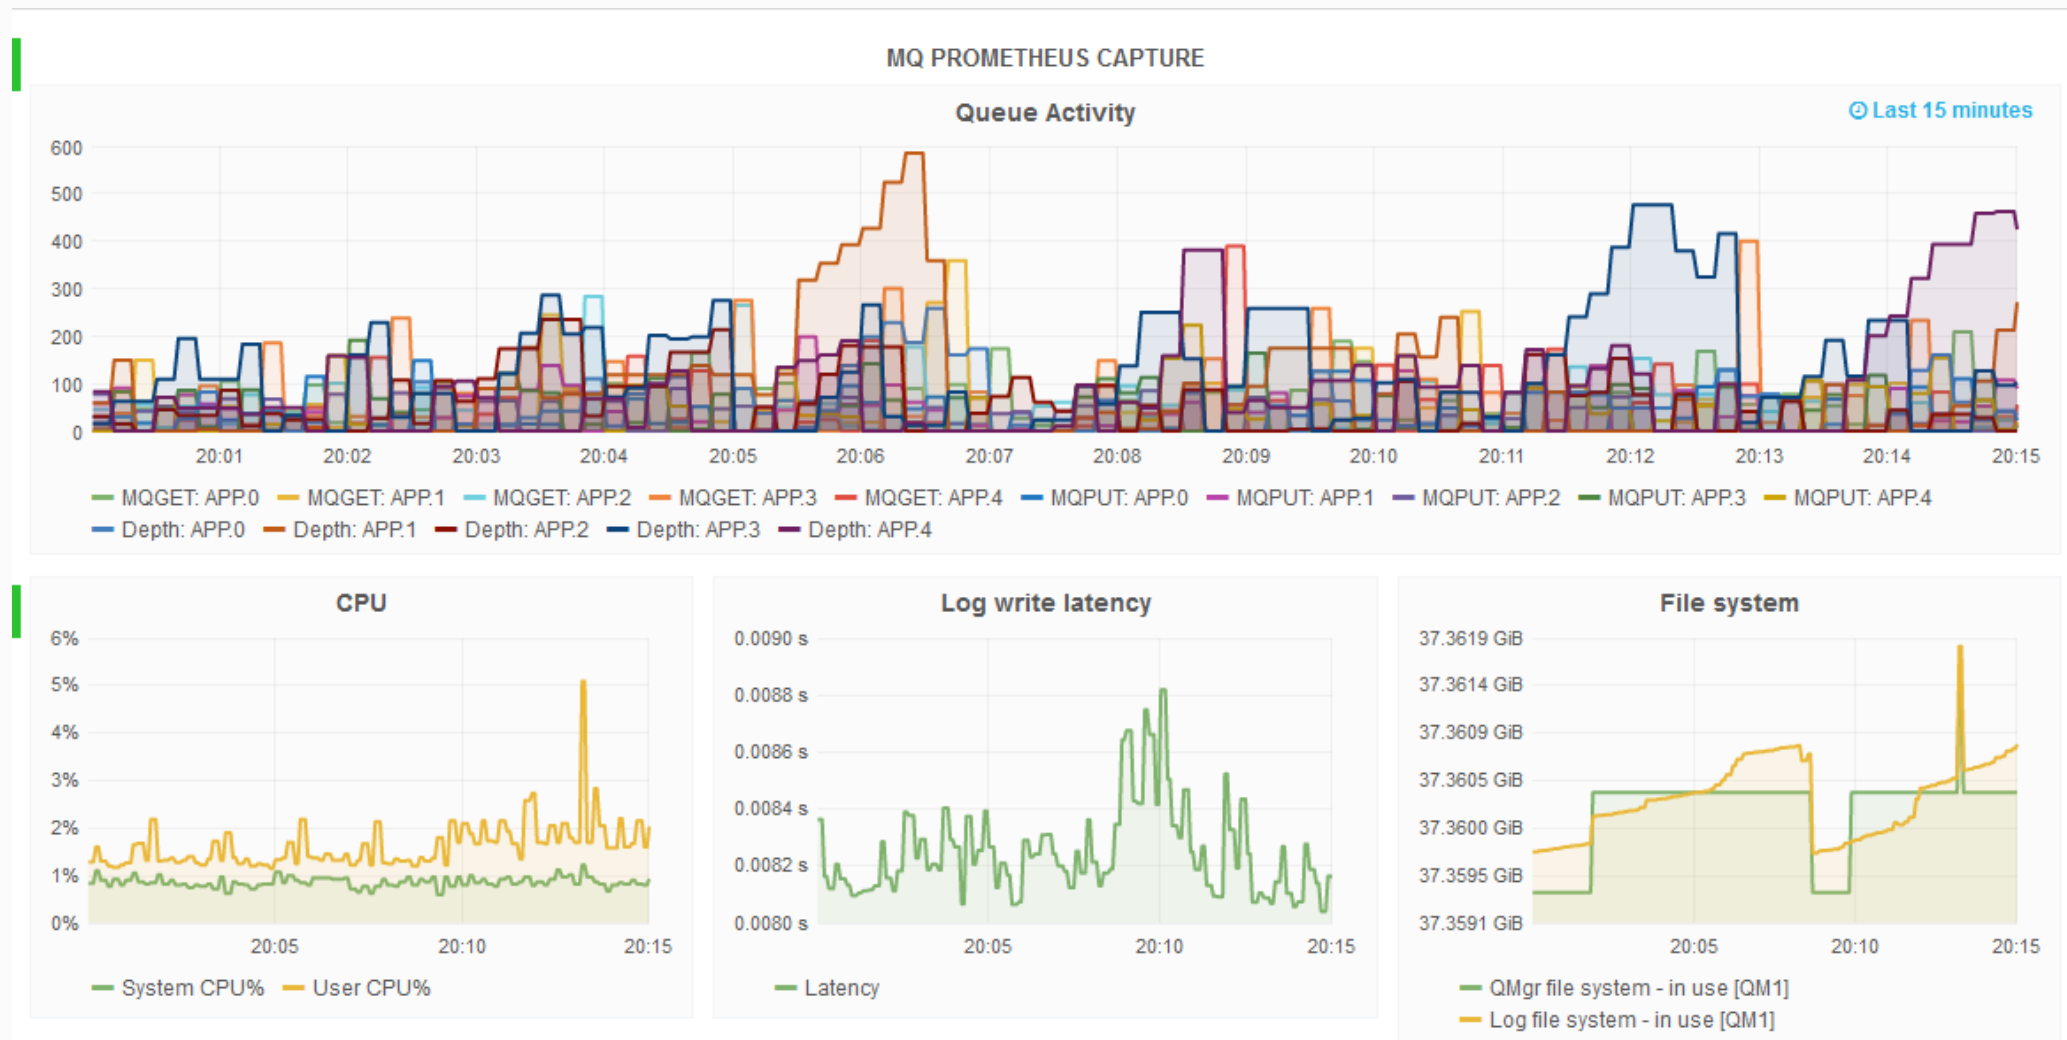
Access proxy ⓘ

Http Auth Basic Auth ☐ With Credentials ☐

Accessing queue stats from Prometheus in Grafana



Grafana dashboard

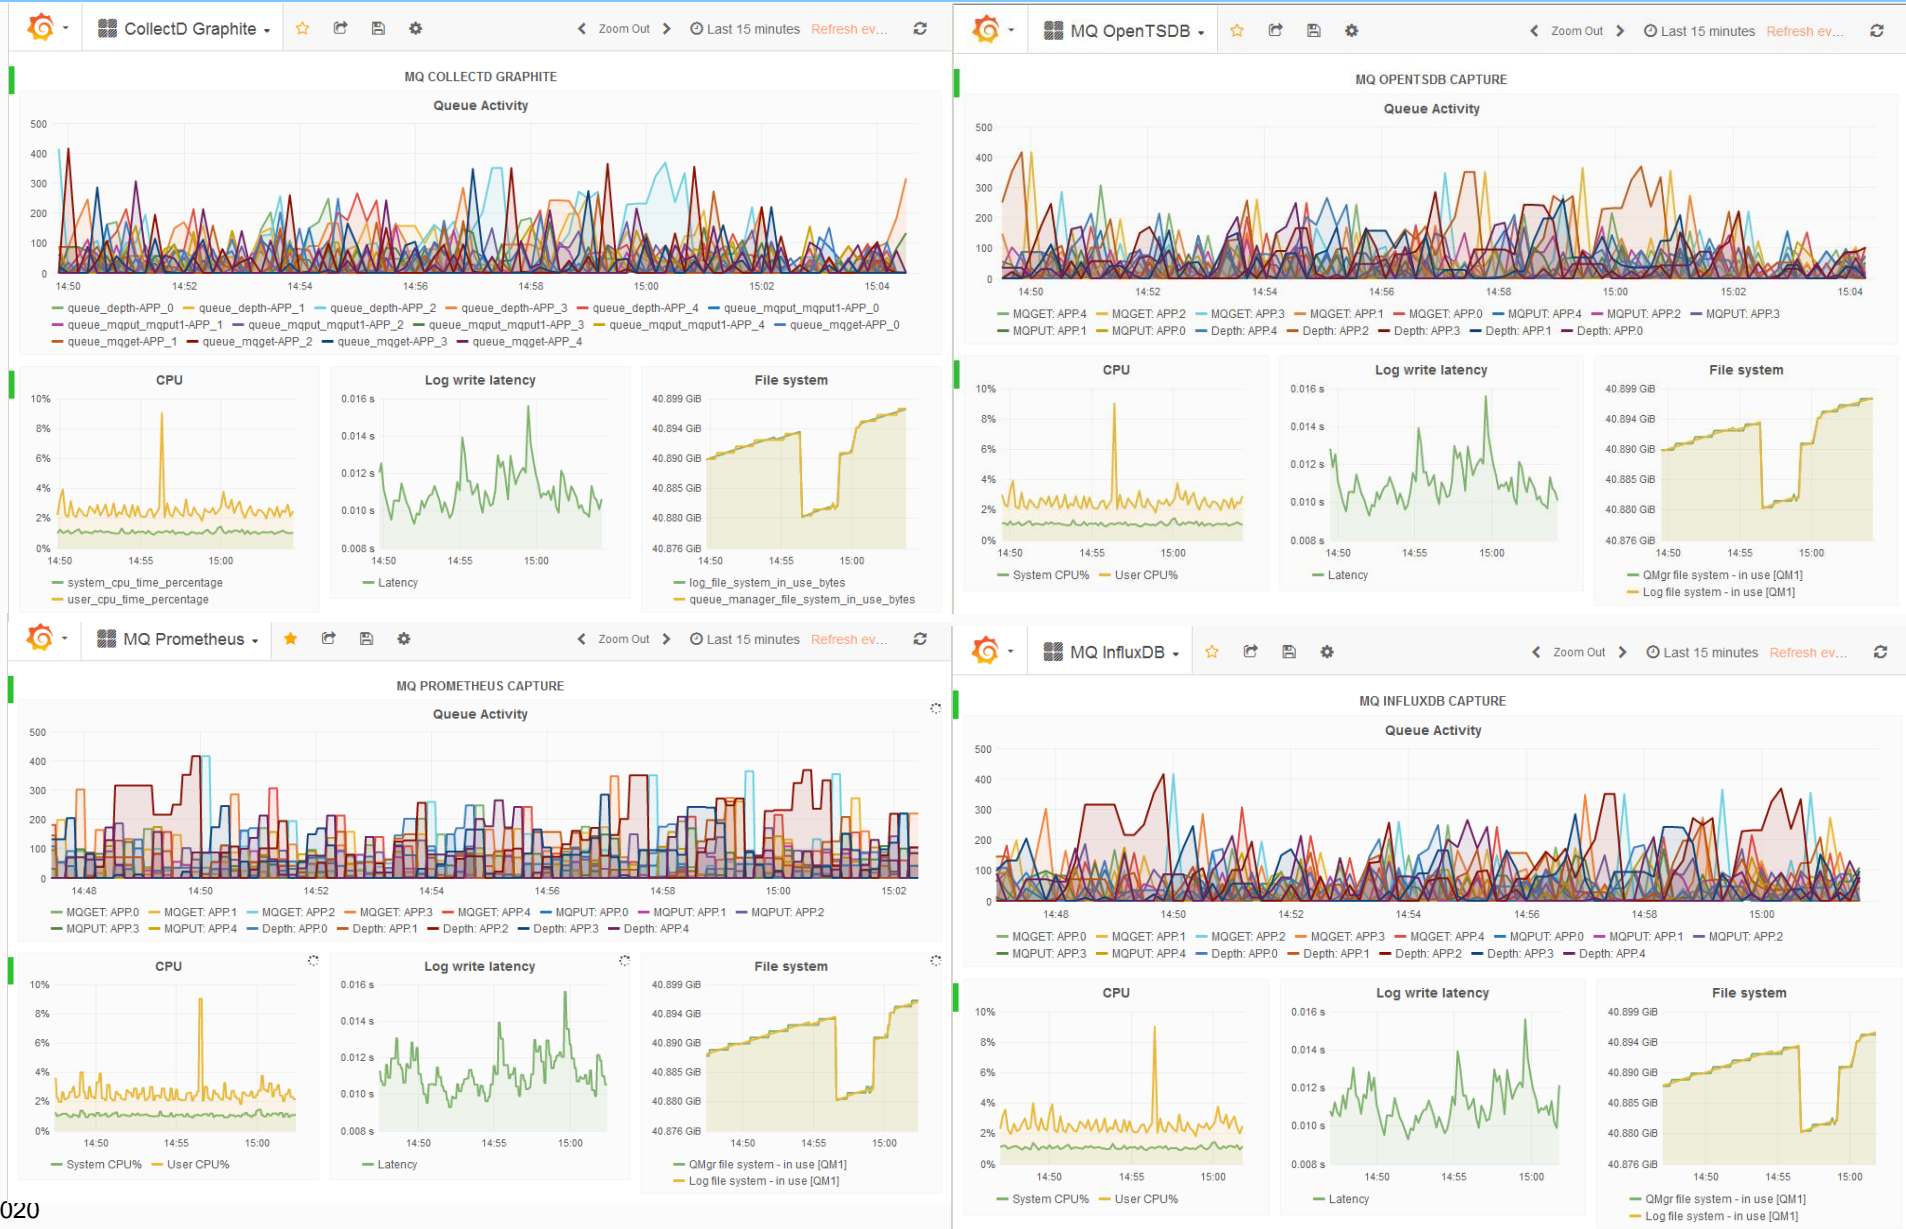


Then added more variants

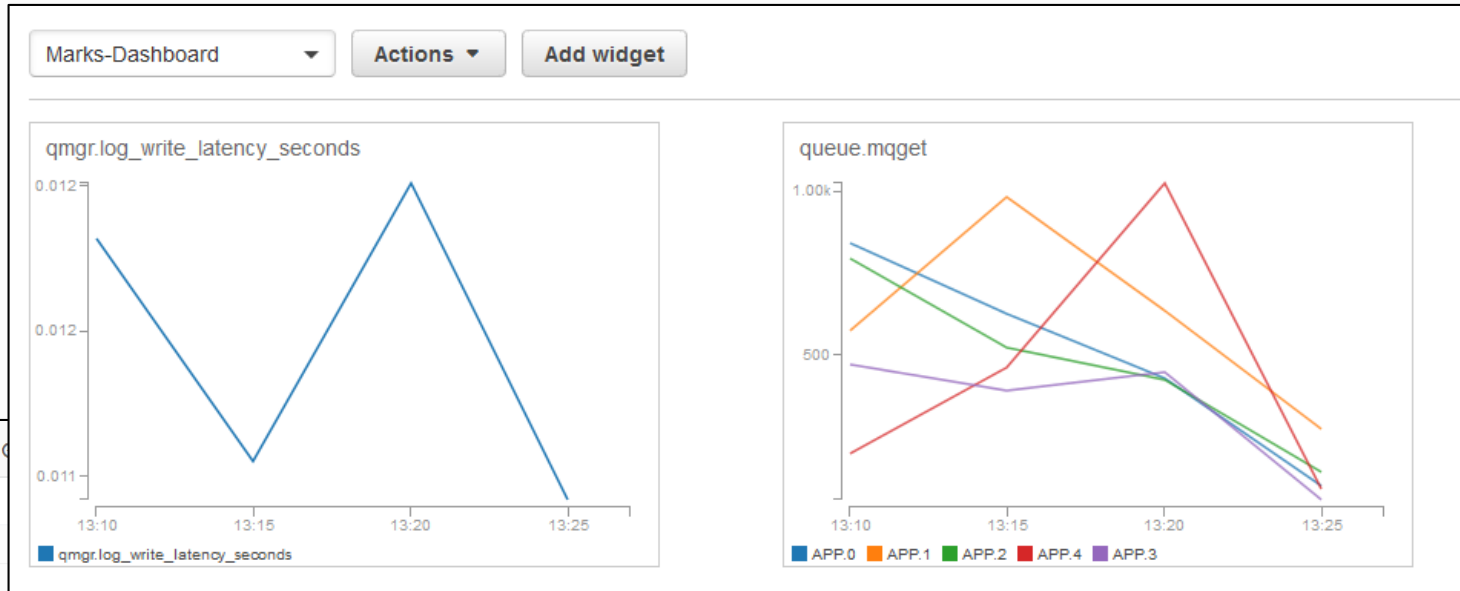
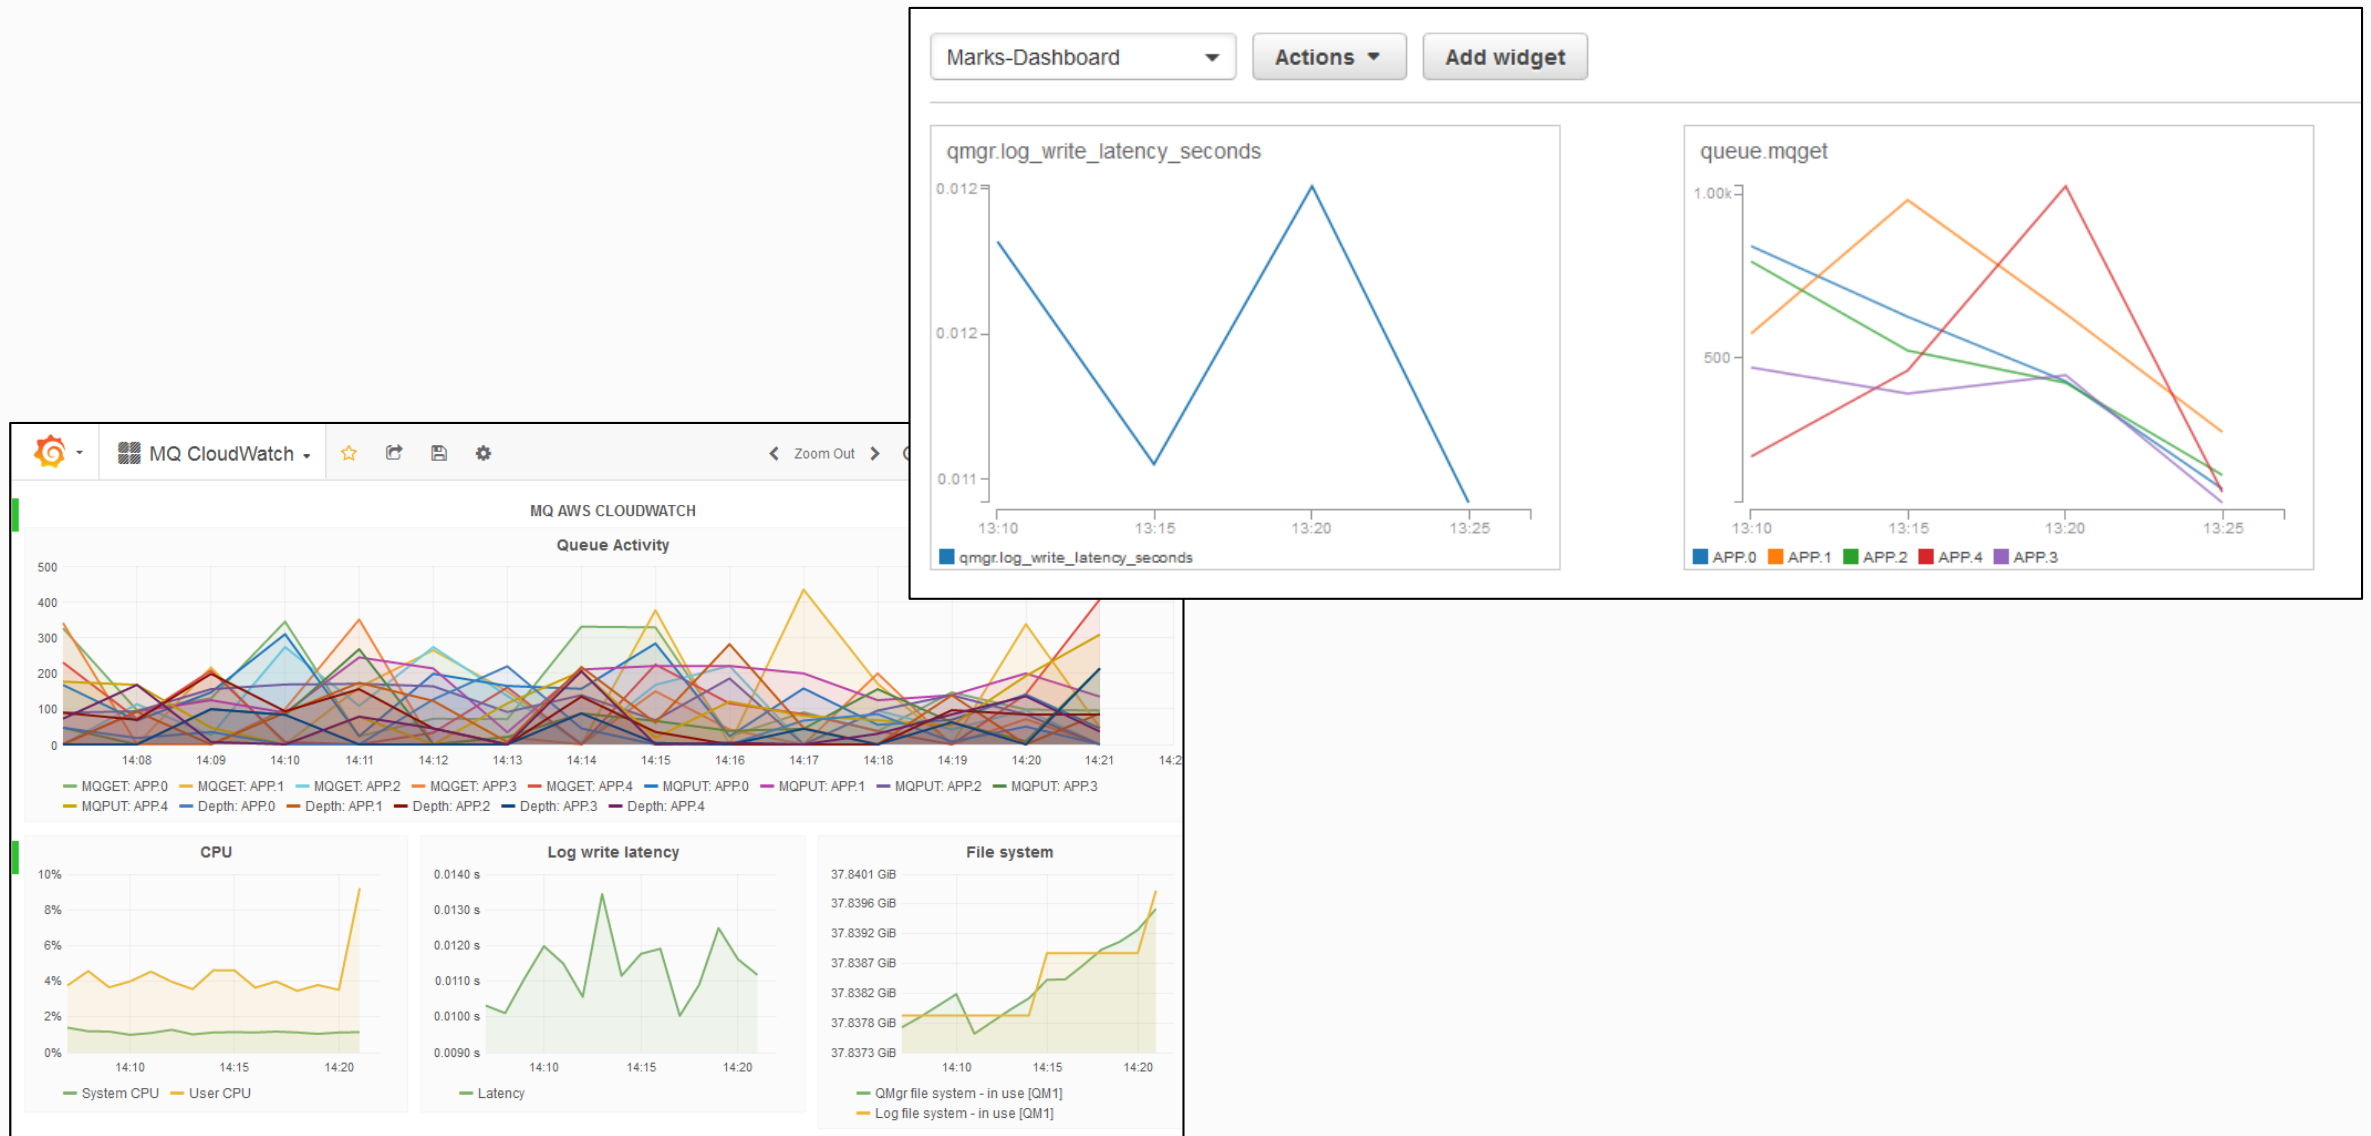
- Rapidly added support for Influx, OpenTSDB
 - Different collectors with slightly different parameters
- Graphite is another database, but fed via collectd
- Also added an AWS collector for CloudWatch
- Generic JSON formatting

```
{ "collectionTime" : {  
    "timeStamp" : "2016-11-07-T15:00:55Z"  
    "epoch" : 1478527255    },  
  "points" : [  
    { "queueManager" : "QM1", "ramTotalBytes" : 15515735206 },  
    { "queueManager" : "QM1", "userCpuTimePercentage" : 1.33 }  
  ]  
}
```

Four equivalent Grafana dashboards



AWS Cloudwatch



Requirement: Monitor channels

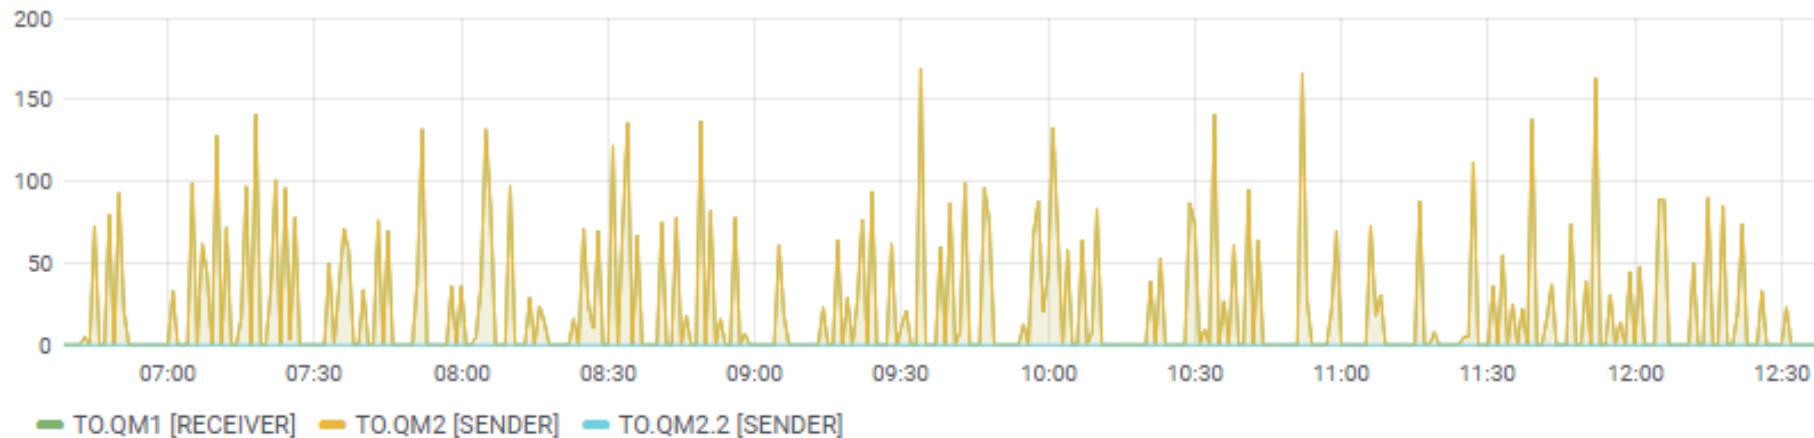
Monitoring - channels

Channels ▾

Channel Status

Channel Name ▲	Connection Name	Local QMgr	Remote QMgr	Type	State	Status
TO.QM1	127.0.0.1	QM1	QM2	RECEIVER	Running	3
TO.QM2	127.0.0.1(1415)	QM1	QM2	SENDER	Running	3
TO.QM2.2	localhost(1415)	QM1		SENDER	Transition	5

Channel Messages

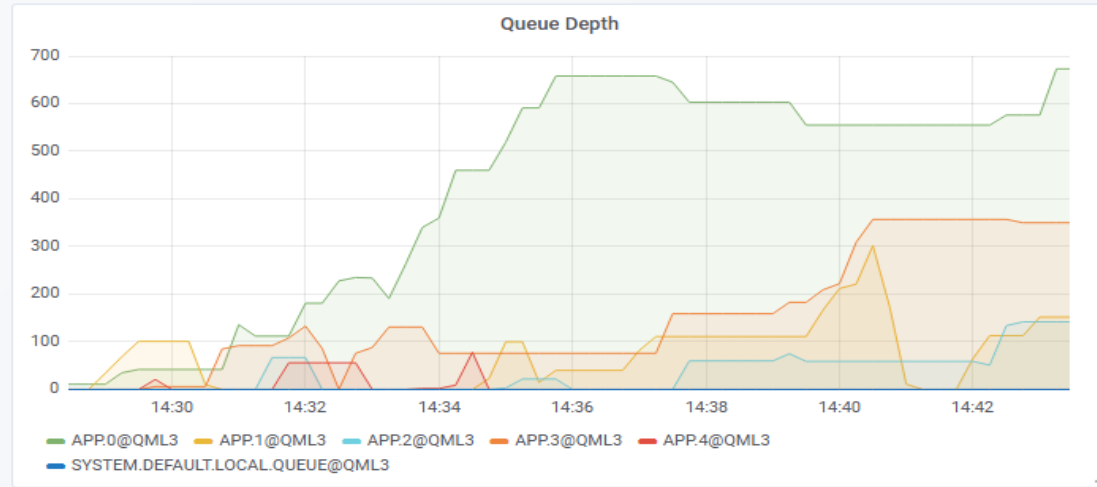


Requirement: Monitor z/OS

Monitoring for z/OS

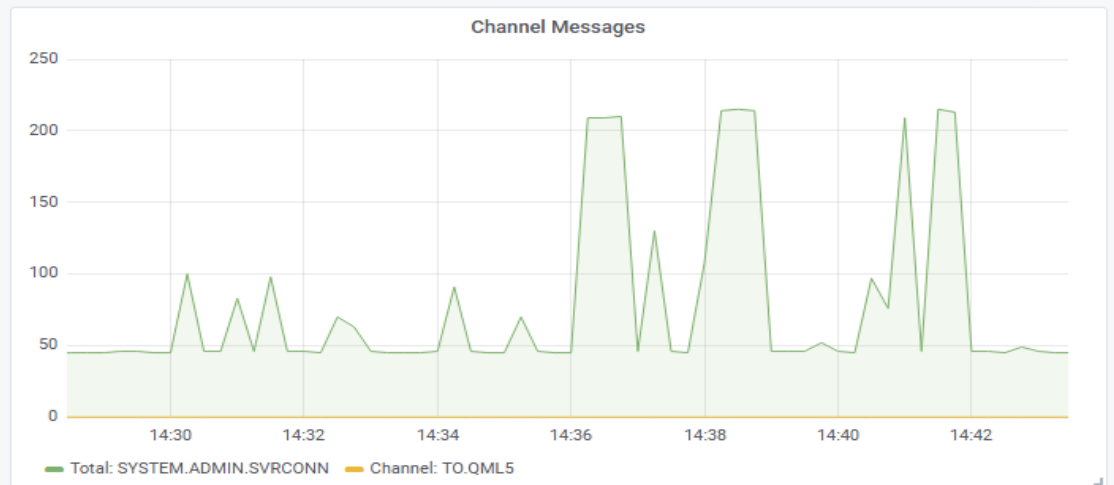
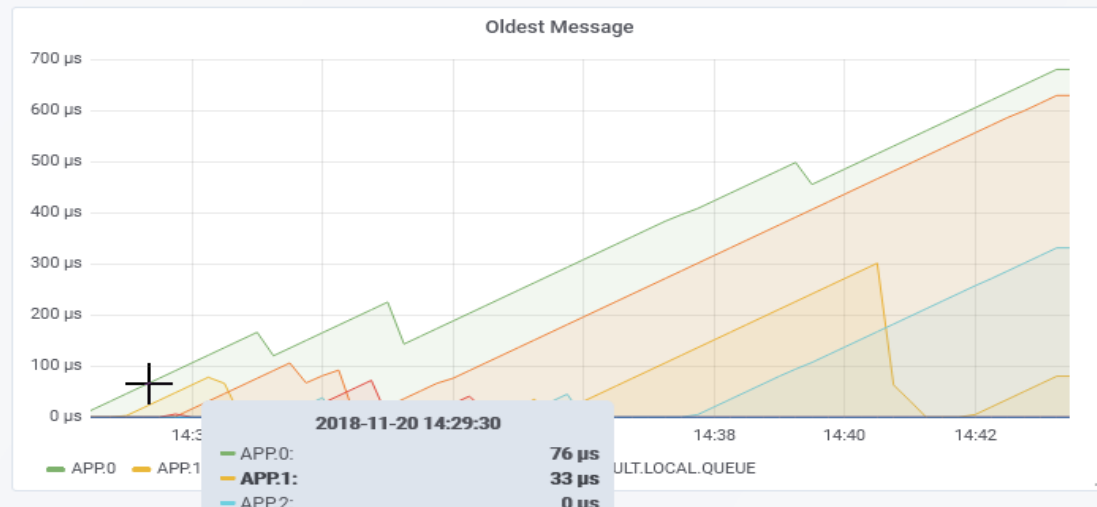
z/OS status - QML3

Last 15 minutes Refresh every 10s



Channel Status

Channel Name	Connection Name	Local QMgr	Type	State	Status
SYSTEM.ADMIN.SVRCONN	192.168.0.121	QML3	SVRCONN	Running	3
SYSTEM.ADMIN.SVRCONN	192.168.0.126	QML3	SVRCONN	Running	3
SYSTEM.ADMIN.SVRCONN	192.168.0.141	QML3	SVRCONN	Running	3
SYSTEM.ADMIN.SVRCONN	192.168.0.99	QML3	SVRCONN	Running	3
SYSTEM.ADMIN.SVRCONN	192.168.215.140	QML3	SVRCONN	Running	3
SYSTEM.ADMIN.SVRCONN	192.168.215.204	QML3	SVRCONN	Running	3

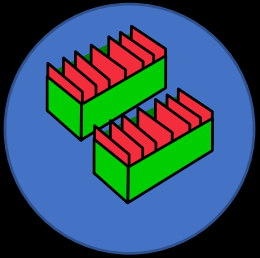


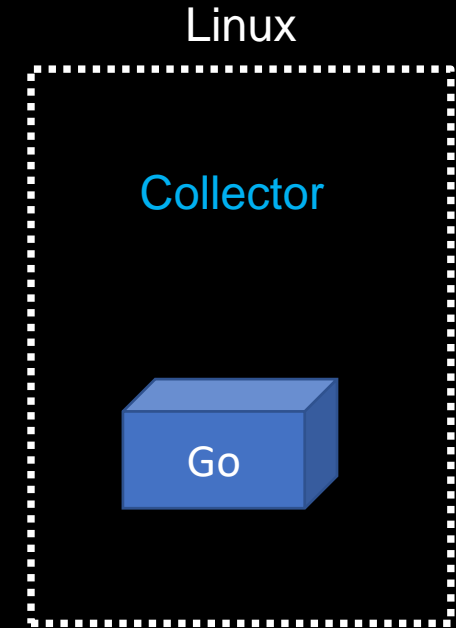
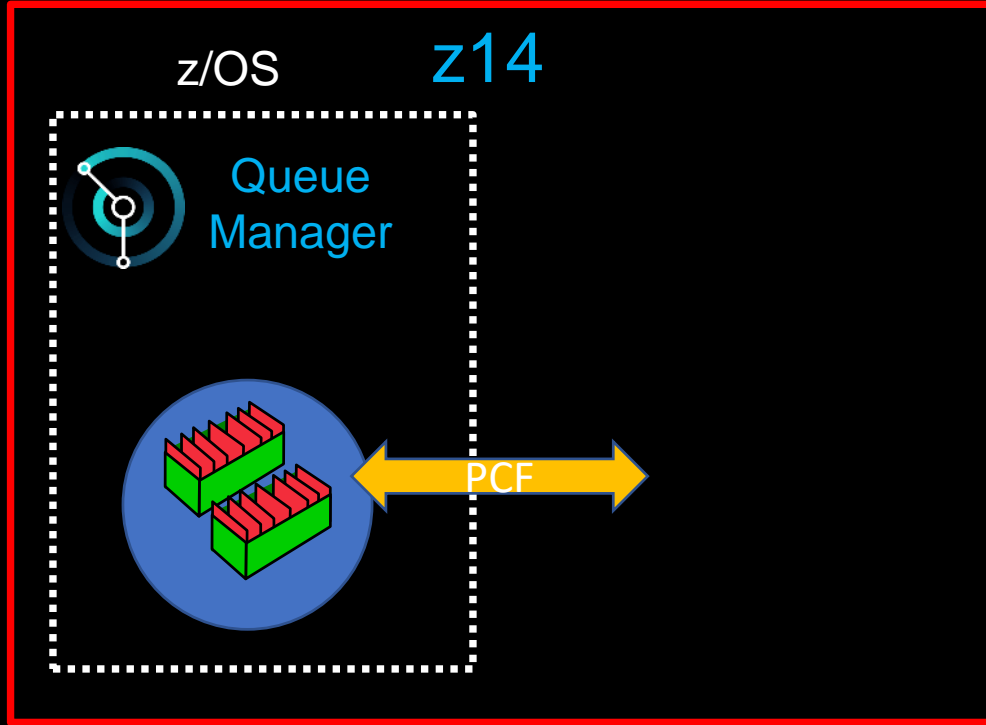
z/OS

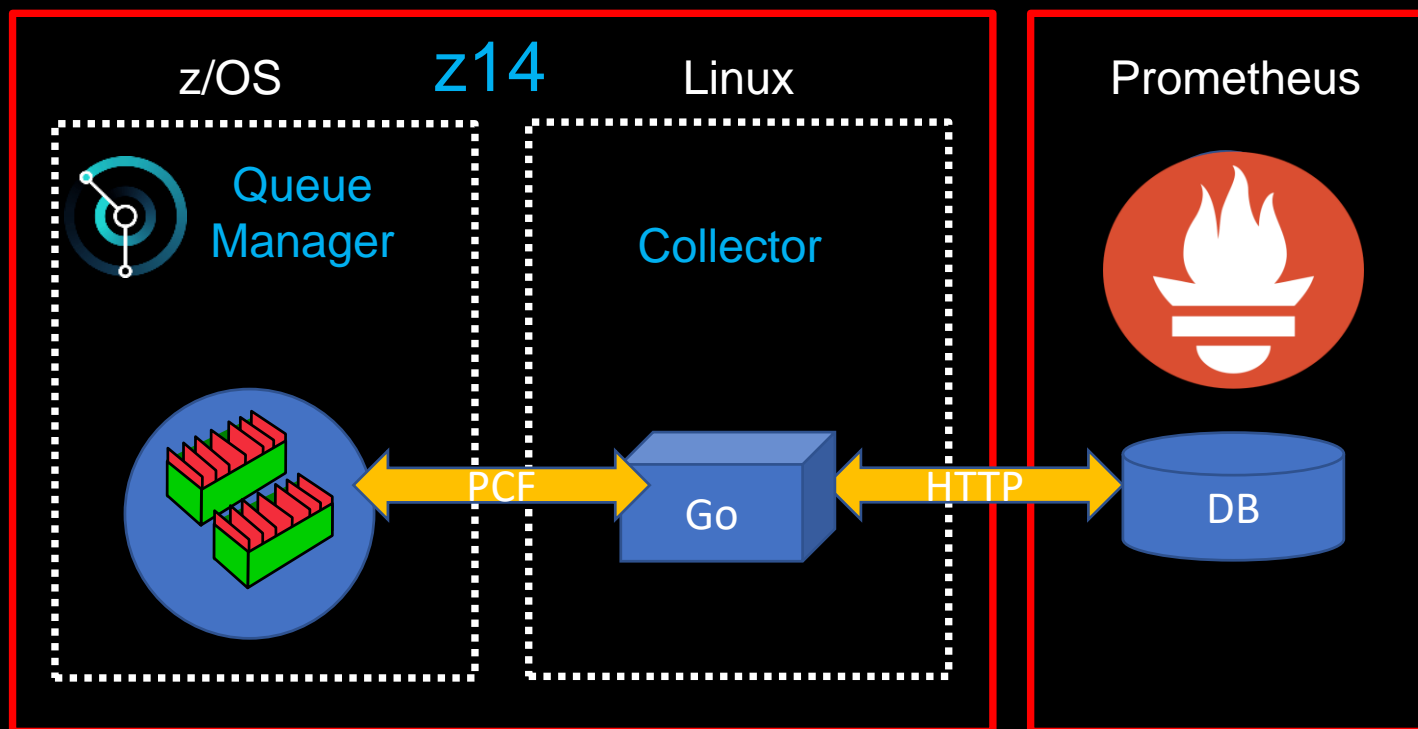
z14

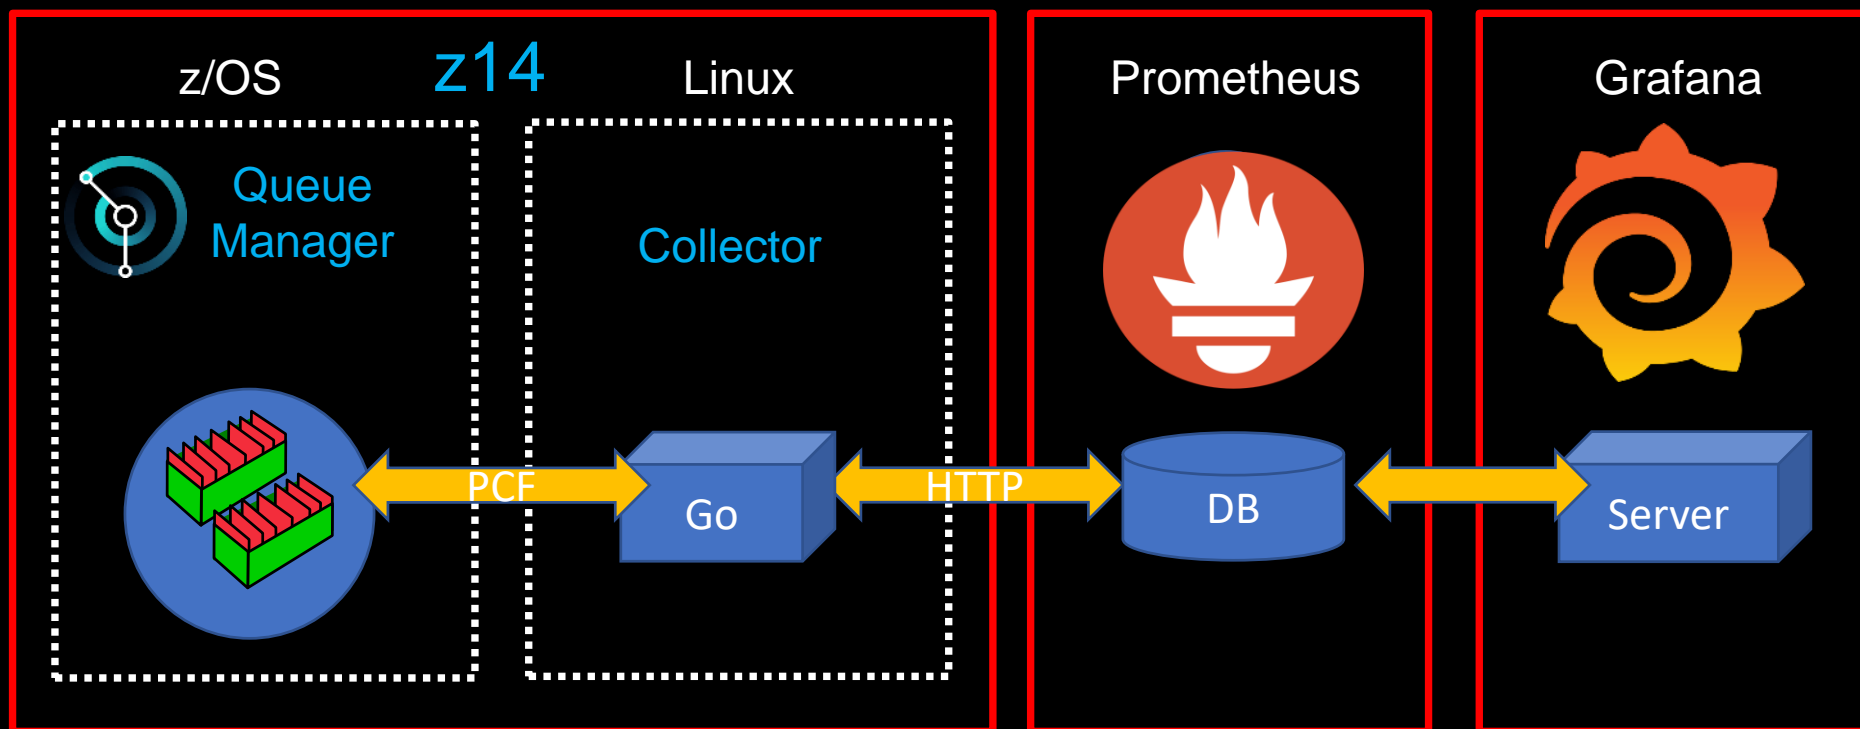


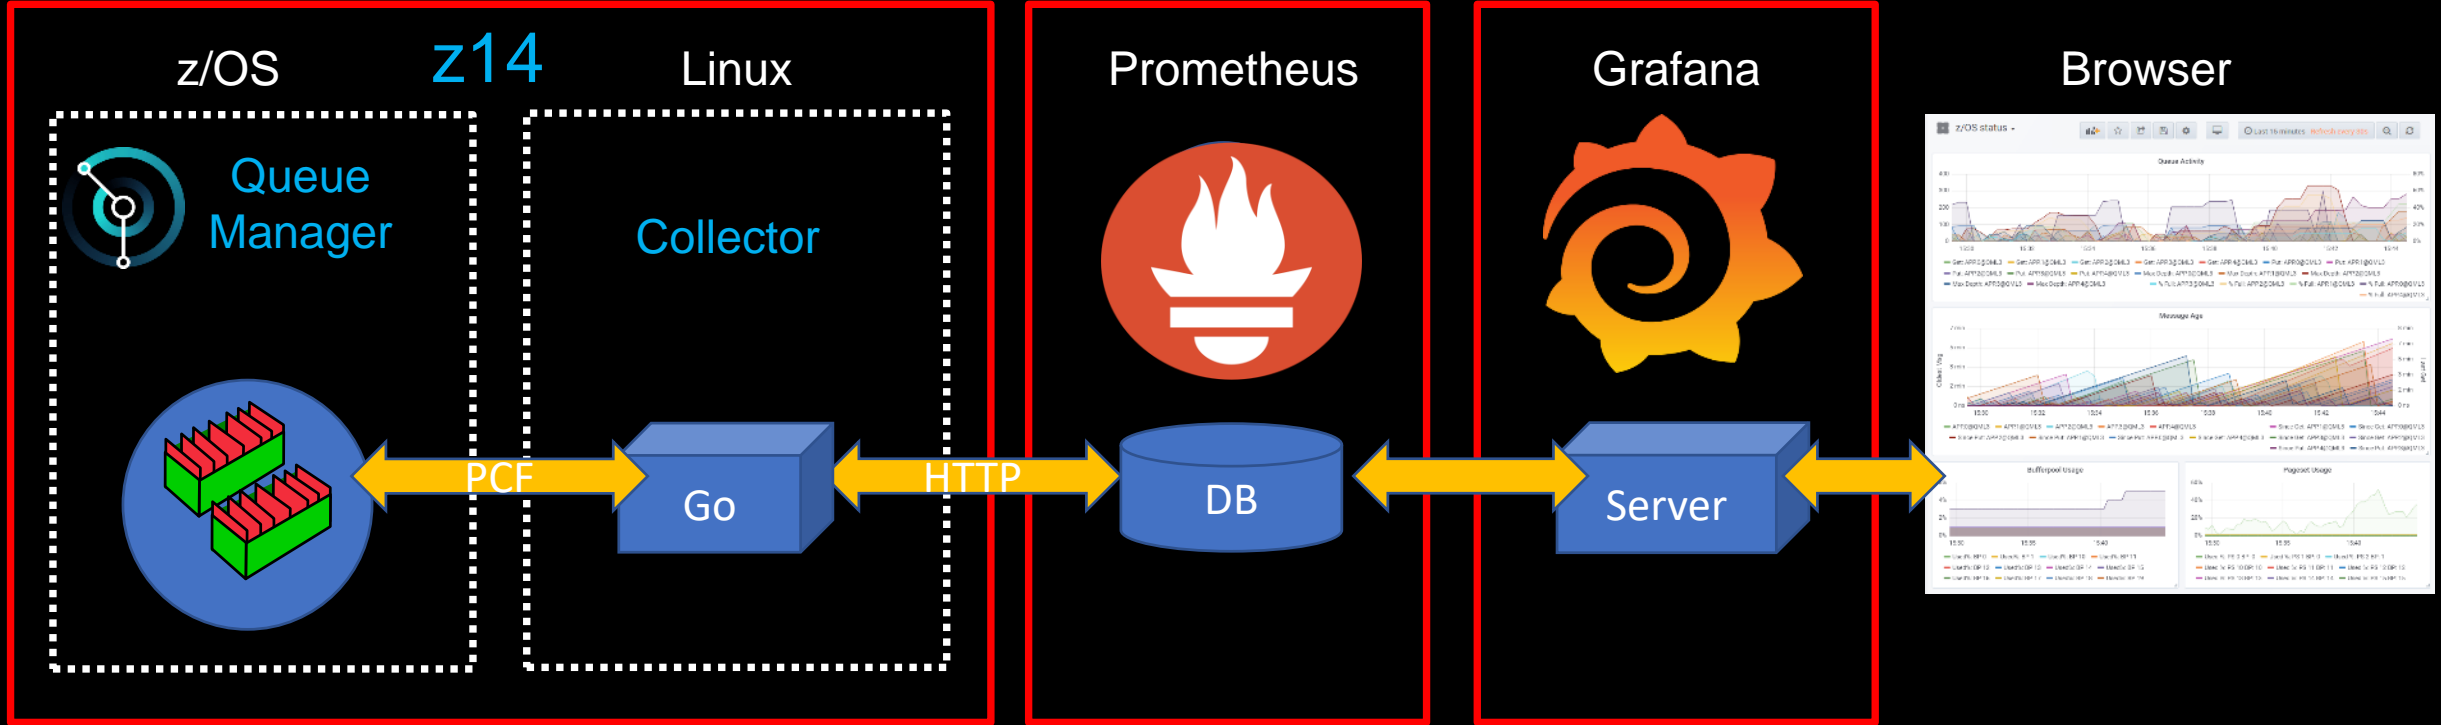
Queue
Manager



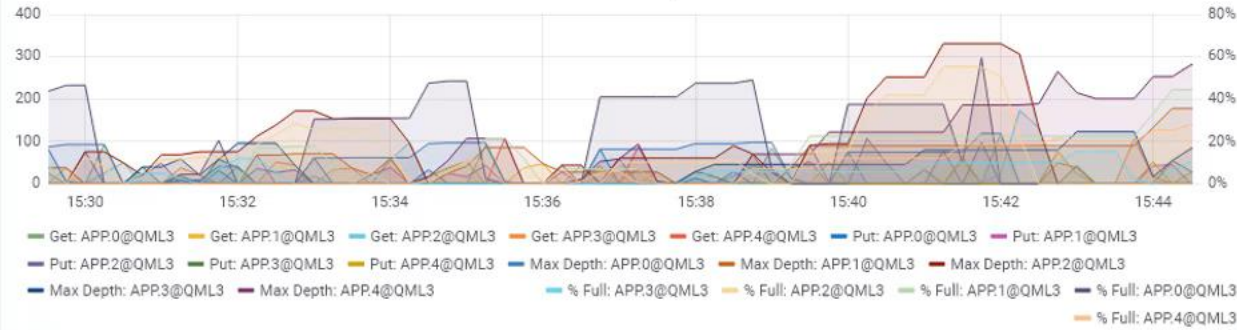




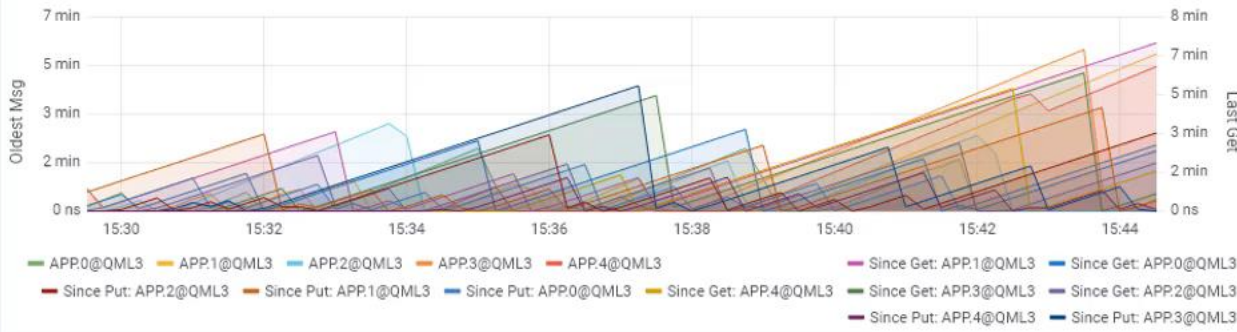




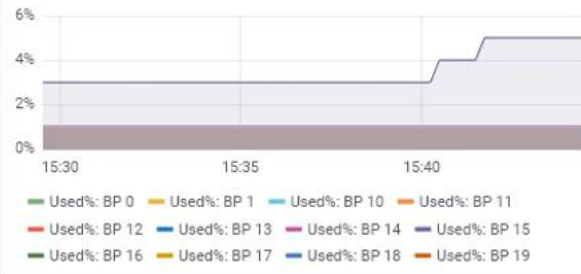
Queue Activity



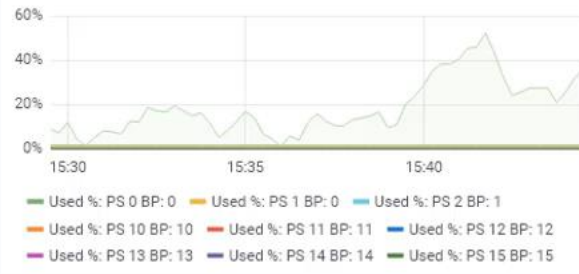
Message Age



Bufferpool Usage



Pageset Usage



Channel Status

Channel Name	Connection Name	Local QMgr	Type	State	Status
SYSTEM.ADMIN.SVRCONN	192.168.0.254	QML3	SVRCONN	Running	3
SYSTEM.ADMIN.SVRCONN	192.168.17.254	QML3	SVRCONN	Running	3
TO.QML5	192.168.17.252	QML3	CLUSSDR	Transition	5

Channel Messages



Unique Channels



Unique Channels

46

BP/PS Usage

Bufferpool	Page Set	Queue Manager	PS Used	BP Used
3	4	QML3	35.6%	5.0%
0	0	QML3	1.6%	1.0%
0	1	QML3	1.0%	1.0%
2	3	QML3	0.3%	1.0%
1	2	QML3	0.2%	1.0%

No longer reliant on resource publications

- The monitors now extract information from STATUS commands
 - Queues, Channels, QMgr, Topics, Subscriptions
- And USAGE for Bufferpools and Pagesets
- Giving key information for ALL platforms including z/OS

New – June 2020

- All collectors brought to same functional level
- Can use YAML file as alternative to command line options

```
prometheus:
  port: 9157
  metricsPath: "/metrics"
  namespace: ibmmq
global:
  useObjectStatus: true
  useResetQStats: false
  logLevel: INFO
  rediscoverInterval: 1h
connection:
  queueManager: QM1
  replyQueue: SYSTEM.DEFAULT.MODEL.QUEUE
objects:
  queues:
    - APP.*
    - "!SYSTEM.*"
  queueSubscriptionSelector:
    - PUT
    - GET
```

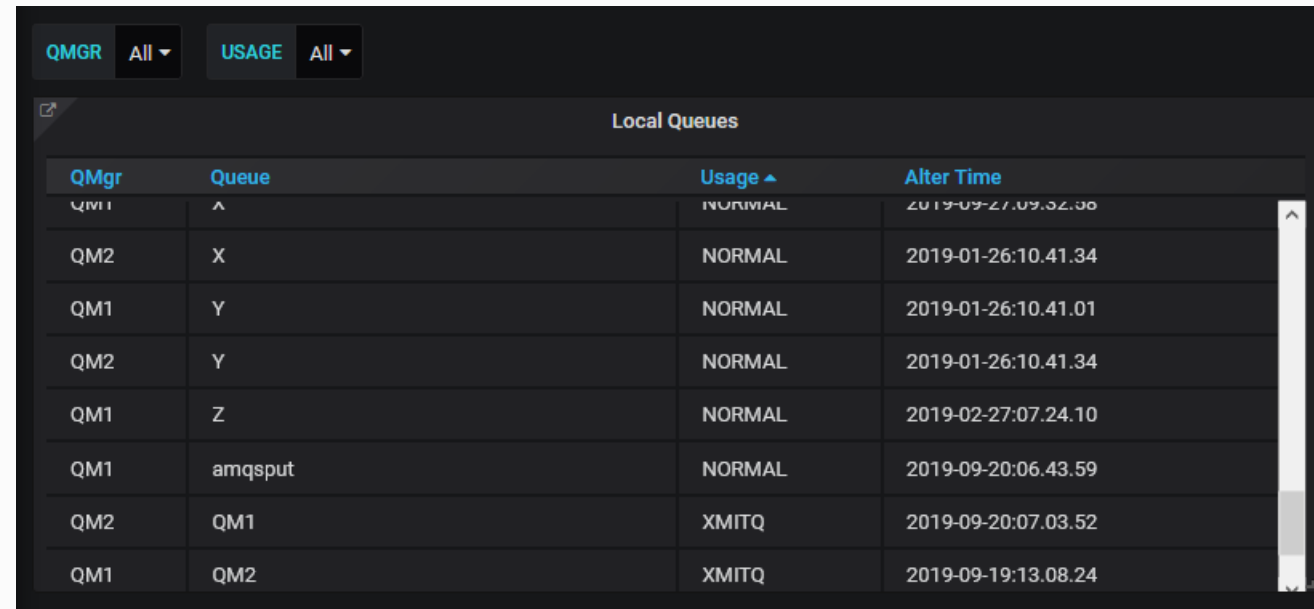
Additional info

- GitHub source: <https://github.com/ibm-messaging/mq-metric-samples>
- z/OS monitoring video (2019): <https://youtu.be/dCTTZWh5NDw>
- Original features video (2016): https://youtu.be/Pi_jHCiqTgU
- RFEs that (if implemented) would help get more data from z/OS:
 - 135074: Queue Statistics
 - 134864: Publish SMF

Connecting metrics to configuration

- Grafana can also display information from some SQL databases
- It is possible to populate tables with MQ configuration information
 - For example, which queues are associated with particular bufferpools and pagesets
- The MQ REST API gets responses in JSON to be directly inserted
 - Databases are able to do SELECT on JSON-formatted columns

See <https://marketaylor.synology.me/?p=532>



The screenshot shows a web interface for monitoring MQGR USAGE. It features a table titled 'Local Queues' with four columns: 'QMGr', 'Queue', 'Usage', and 'Alter Time'. The table contains eight rows of data. The first row is partially obscured by a header row. The 'Usage' column shows values like 'NORMAL', 'XMITQ', and 'XMITQ'. The 'Alter Time' column shows timestamps in YYYY-MM-DD:HH:MM:SS format.

QMGr	Queue	Usage	Alter Time
QM1	^	NORMAL	2019-09-27:09:52.30
QM2	X	NORMAL	2019-01-26:10:41.34
QM1	Y	NORMAL	2019-01-26:10:41.01
QM2	Y	NORMAL	2019-01-26:10:41.34
QM1	Z	NORMAL	2019-02-27:07:24.10
QM1	amqsput	NORMAL	2019-09-20:06:43.59
QM2	QM1	XMITQ	2019-09-20:07:03.52
QM1	QM2	XMITQ	2019-09-19:13:08.24

QMGR

All ▾

USAGE

XMITQ ▾

Local Queues

QMgr	Queue	Usage ▴	Alter Time
QM2	QM1	XMITQ	2019-09-20:07.03.52
QM1	QM2	XMITQ	2019-09-19:13.08.24
QM1	QM2.RETRY	XMITQ	2019-09-30:14.49.41
QM1	QM2.STOPPED	XMITQ	2019-06-11:09.53.22

Query

MySQL ▾

▼ A




```
SELECT doc->>"$.queueManager" as 'QMgr',
       doc->>"$.a.queue" as 'Queue',
       doc->>"$.a.usage" as 'Usage',
       concat(doc->>"$.a.altdate", ":", doc->>"$.a.alttime") as "Alter Time"
FROM MQCFG.QLOCAL
WHERE doc->>"$.a.queue" NOT LIKE 'SYSTEM.%' AND (
  CASE WHEN $USAGE = 'ALL' THEN
    doc->>"$.a.usage" LIKE '%'
  ELSE
    doc->>"$.a.usage" = $USAGE
  END \ AND /
```

Format as

Table ▾

Show Help ▶

Generated SQL ▶



Adding resource statistics to your own applications

- Article showing how to publish similar statistics from your own applications
 - And therefore have monitors such as these showing status
 - Even if your apps are connecting to a z/OS queue manager
- Based on the MQ Salesforce Bridge code
 - Shows how to construct the PCF metadata describing your resources
- See <https://marketaylor.synology.me/?p=380>

What are differences? Which is best?

- Differences are generally in
 - The names and formats of metrics ("ibmmq_queue_mqget")
 - Naming for individual resources such as the queue name
 - Query capabilities to select and display chosen metrics
 - Can you use wildcards on object names
 - Creating labels on graphs
 - Can it be automatic based on the query?
 - Alerting capabilities
- The best is going to be whatever you are already using!
 - But I found the Prometheus/Grafana combination to be flexible and usable
- No easy way to report as string (eg "STARTED", "STOPPED" status)
 - Have to do a mapping via an integer or label

Squashing Metrics

Type	State	Status
RECEIVER	Running	3
SENDER	Running	3
SENDER	Transition	5

ibmmq_channel_status_squash+0	
Legend format	legend format
ibmmq_channel_status+0	
Legend format	legend format

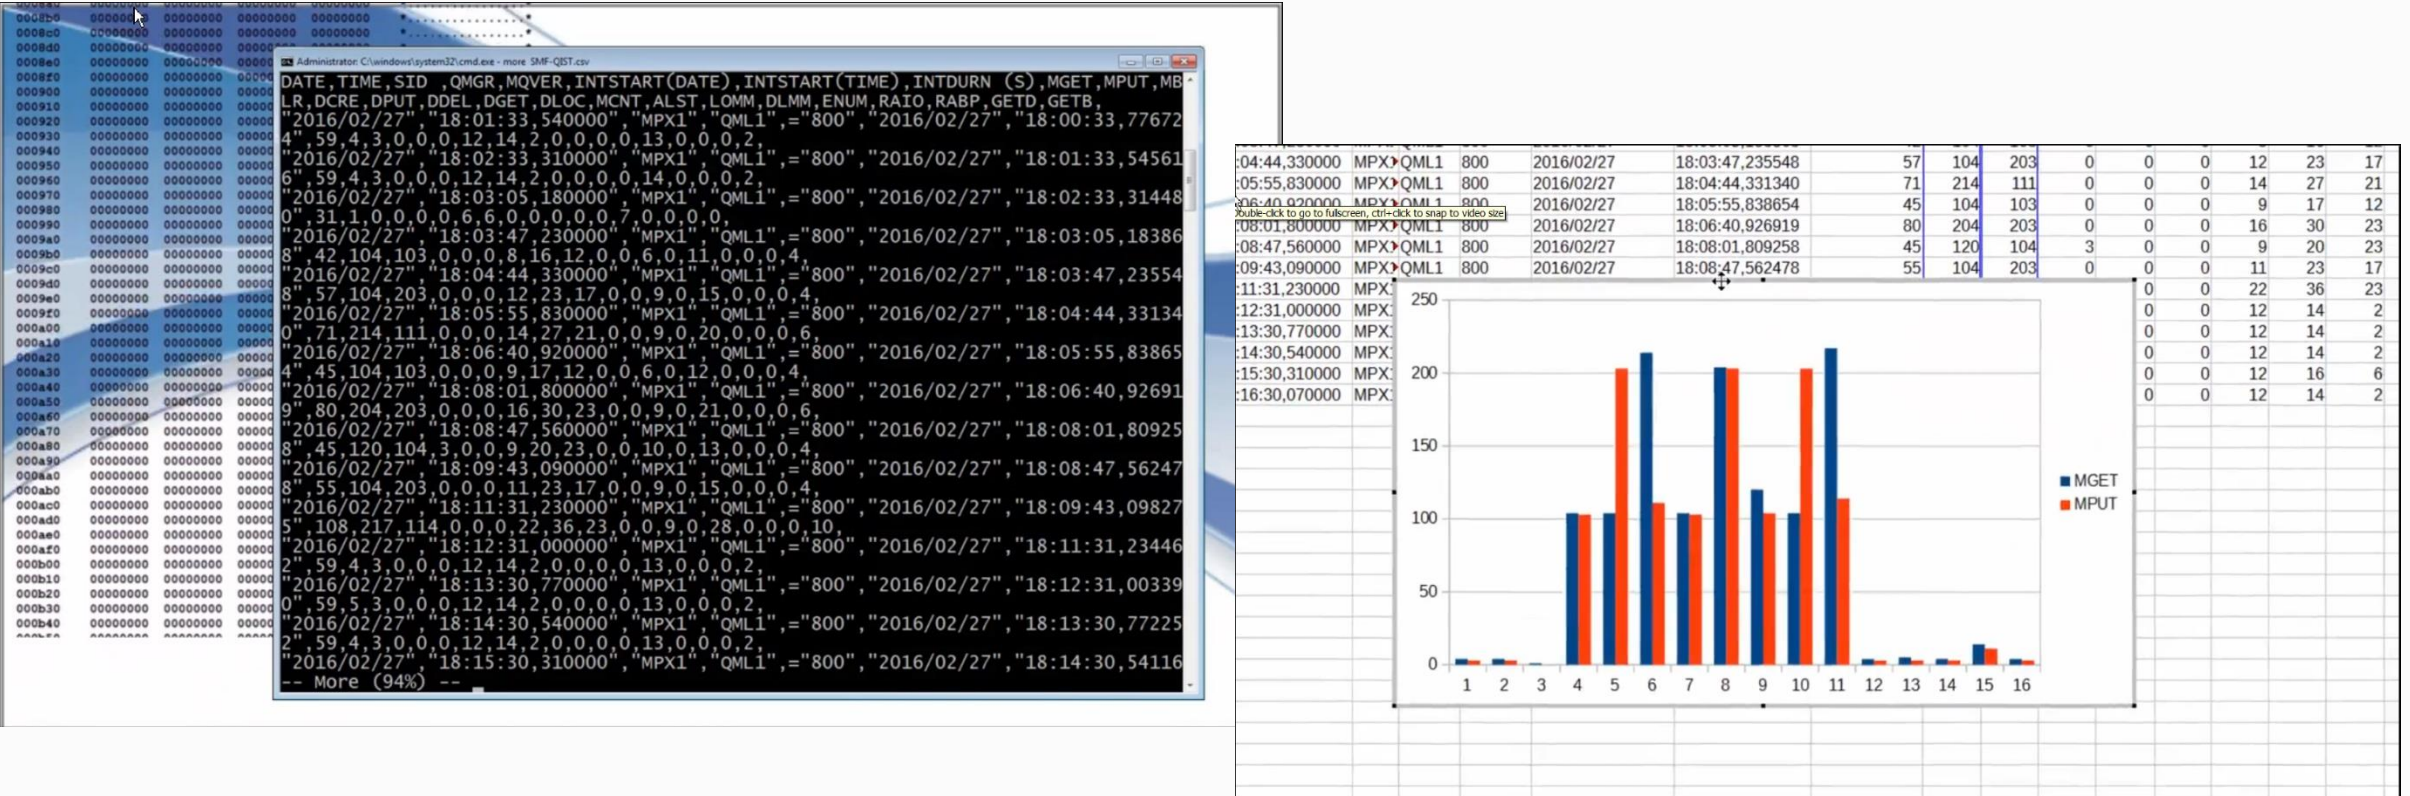
Value #A	Options	Type	Value Mappings	Thresholds
channel	Apply to columns named Value #A	Type String	Type Value to text	Thresholds 1,2
qmgr	Column Header State	Sanitize HTML	0 Stopped	Color Mode Cell
Value #B	Render value as link	Preserve Formatting	1 Transition	Colors Invert
Value #C			2 Running	
type			+	

Latest Go/Prometheus status

- All repositories under <https://github.com/ibm-messaging>
- **mq-golang** has the core MQI and PCF packages
 - Some sample code to demonstrate use of most functions
 - Assumes you already know the MQI principles from another language
- **mq-metric-samples** has Prometheus, Cloudwatch etc monitor programs
 - Along with a "vendor" tree
- Both have scripts to compile programs in Docker containers for ease of build

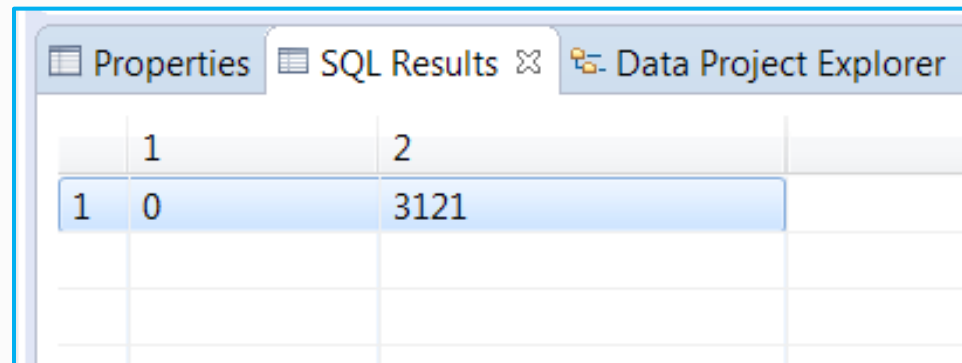
Similar resource data available on z/OS but via SMF

- mqsmfcsv ... open source tool to format MQ z/OS SMF records for easy import to spreadsheets and databases
 - <http://github.com/ibm-messaging/mq-smf-csv>



Example queries

- What was my largest message size retrieved for this queue?
 - `SELECT MAX(Get_Max_Msg_Size) from MQSMF.WQ where (Base_Name='LYNS.TEST.QUEUE');`
 - Result was 11,189 (application people insisted it was 3,800)
- How many MQPUTs and MQPUT1s were completed?
 - `SELECT SUM (Put_Count), SUM (Put1_Count) from MQSMF.WQ where (Base_Name = 'LYNS.TEST.QUEUE');`
 - Results:



The screenshot shows a software interface with three tabs: 'Properties', 'SQL Results', and 'Data Project Explorer'. The 'SQL Results' tab is active, displaying a table with two columns and one row of data. The first column contains the value '1' and the second column contains the value '3121'.

1	2
1	3121

And we can now do it in JSON

- `mqsmfcsv -i <input file> -f json`

```
{
  "recordType" : 116,
  "recordSubType" : 0,
  "structure" : "QMAC",
  "date" : "2015/11/23",
  "time" : "11:00:00.020000",
  "lpar" : "H019",
  "qmgr" : "MQPC",
  "mqVersion" : "800",
  "authorisationId" : "IMS      ",
  "correlId" : "F0F2F3F6C2C3F1E4C4D6C340",
  "connectionName" : "PRDC      ",
  "operatorId" : "PLN1231 ",
  "applicationType" : "IMS MPP/BMP",
  "accountingToken" :
  "0000000000000000000000000000000000000000000000000000000000000000",
  "networkId" :
  "D7D9C4C340404040044E0A0800000001",
  ...
}
```

Processing other MQ events

- Already shown amqsevt as shipped in MQ V8
- It now also supports JSON output option
 - Included from V9.1
- Can be used to feed JSON consumers such as splunk

MQ events in splunk

The screenshot displays the Splunk web interface with the following components:

- Navigation Bar:** Search, Datasets, Reports, Alerts, Dashboards. The 'Alerts' tab is highlighted.
- Search Bar:** Contains the query `host=0b9f3995a92b "eventSource.objectName"="SYSTEM.ADMIN.PERFM.EVENT"`.
- Results Summary:** Shows 2 events for the time range 04/11/2016 12:10:34.000 to 04/11/2016 12:10:35.000. Includes a 'No Event Sampling' dropdown.
- Event List:** A table with columns 'Time' and 'Event'. The first event is expanded, showing its JSON data.
- Field List:** On the left, under 'Selected Fields', are `host`, `source`, and `sourcetype`. Under 'Interesting Fields', are `eventCreation`, `eventData.baseObjectName`, `eventData.highQueueDepth`, `eventData.msgDeqCount`, `eventData.msgEnqCount`, `eventData.queueMgrName`, `eventData.timeSinceReset`, `eventReason.name`, `eventReason.value`, `eventSource.objectName`, `eventSource.objectType`, `eventType.name`, and `eventType.value`.

The expanded event JSON data is as follows:

```
> 04/11/2016 12:10:34.000 {
  "eventSource" : { "objectName": "SYSTEM.ADMIN.PERFM.EVENT",
                    "objectType" : "Queue" },
  "eventType" : {
    "name" : "Perfm Event",
    "value" : 45
  },
  "eventReason" : {
    "name" : "Queue Full",
    "value" : 2053
  },
  "eventCreation" : "2016/11/04 12:10:24.29 GMT",
  "eventData" : {
    "queueMgrName" : "V9000_A",
    "baseObjectName" : "FULLEVT",
    "timeSinceReset" : 0,
    "highQueueDepth" : 4,
    "msgEnqCount" : 0,
    "msgDeqCount" : 0
  }
}
```

Below the JSON, the source information is displayed: `host = 0b9f3995a92b | source = /mqm/jsonevt.txt | sourcetype = _json`.

Using JSON event formatter with Activity Events

- Use the event formatter to output in JSON and then filter it further
 - Run via "service" if local qmgr
 - Could also use subscribe variant to obtain trace

```
amqsevt -m QM1 -q SYSTEM.ADMIN.TRACE.ACTIVITY.QUEUE -o json | jq -r -f jqFilt
```

- And then get one-line CSV output of key fields

```
"amqspu" ,"2018-07-11" ,"08:16:48" ,"Connx" ,0 ,"N/A"  
"amqspu" ,"2018-07-11" ,"08:16:48" ,"Open" ,0 ,"QL1"  
"amqspu" ,"2018-07-11" ,"08:16:48" ,"Put" ,0 ,"QL1" ,48 ,"414D512056393030305F4120202020205B2B77"  
"amqspu" ,"2018-07-11" ,"08:16:48" ,"Close" ,0 ,"QL1"  
"amqspu" ,"2018-07-11" ,"08:16:48" ,"Disc" ,0 ,"N/A"  
"amqsge" ,"2018-07-11" ,"08:16:48" ,"Connx" ,0 ,"N/A"  
"amqsge" ,"2018-07-11" ,"08:16:48" ,"Open" ,0 ,"QL1"  
"amqsge" ,"2018-07-11" ,"08:16:48" ,"Get" ,0 ,"QL1" ,38 ,"414D512056393030305F4120202020205B2B77"  
"amqsge" ,"2018-07-11" ,"08:16:48" ,"Get" ,2033 ,"QL1" ,250059 ,0  
"amqsge" ,"2018-07-11" ,"08:16:48" ,"Close" ,0 ,"QL1"  
"amqsge" ,"2018-07-11" ,"08:16:48" ,"Disc" ,0 ,"N/A"
```

<https://marketaylor.synology.me/?p=373>

A jq filter

```
select(.eventData.activityTrace != null) | .eventData.applName as $applName |
  (.eventData.activityTrace[] |
    [
      $applName, .operationDate, .operationTime, .operationId, .reasonCode.value,
      if (.objectName | length) > 0
      then
        .objectName
      else
        "N/A"
      end,
      if .operationId == "Get" or .operationId == "Put" or .operationId == "Put1"
      then
        .qmgrOpDuration, .msgId
      else
        empty
      end
    ]
  ) |
  @csv
```

Accounting records

```
{ "eventSource": { "objectName": "SYSTEM.ADMIN.ACCOUNTING.QUEUE", "objectType": "Queue" },
  "eventType": { "name": "Accounting MQI", "value": 167 }, "eventReason": { "name": "None", "value": 0 },
  "eventCreation": { "timeStamp": "2018-07-25T20:46:19Z", "epoch": 1532551579 },
  "eventData": { "queueMgrName": "V9000_A", "startDate": "2018-0725",
    "startTime": "21.46.16", "endDate": "2018-07-25", "endTime": "21.46.19", "commandLevel": 905,
    "connectionId": "414D514356393030305F4120202020205B2B779523E91D95", "sequenceNumber": 0,
    "applName": "amqsevt", "processId": 35389620, "threadId": 1, "userIdentifier": "metaylor", "connDate":
    "2018-07-25", "connTime": "21.46.16", "discDate": "2018-07-25", "discTime": "21.46.19",
    "opens": [0, 3, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
    "opensFailed": [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
    "closes": [0, 3, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
    "closesFailed": [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
    "puts": [24, 0],
    "putsFailed": 0,
    "put1s": [1, 0],
    "put1sFailed": 0,
    "putBytes": [2547972, 0],
    ...
  }
```

Multiple consumers for MQ events

- Traditional MQ events (queue full etc) are put to a specific named queue
- Makes it difficult to have multiple consumers for same event queue
 - Many monitors can be configured to "browse" but who does "get" and when?
- The MQ event queues can be redefined as topic aliases
- Monitor programs can then get independently from their own dedicated queues
 - I might then run Omegamon AND the JSON variant of amqsevt to different consoles

```
DELETE QLOCAL (SYSTEM.ADMIN.CHANNEL.EVENT)      PURGE
DELETE QLOCAL (SYSTEM.ADMIN.PERFM.EVENT)         PURGE

DEFINE QALIAS (SYSTEM.ADMIN.CHANNEL.EVENT) TARGET (SYSTEM.ADMIN.EVENT)  TARGTYPE (TOPIC)
DEFINE QALIAS (SYSTEM.ADMIN.PERFM.EVENT)    TARGET (SYSTEM.ADMIN.EVENT)  TARGTYPE (TOPIC)

DEFINE TOPIC (SYSTEM.ADMIN.EVENT)  TOPICSTR ('SYSTEM/ADMIN/EVENT')
DEFINE QLOCAL (SYSTEM.ADMIN.SUBSCRIBED.EVENT)
DEFINE SUB (SYSTEM.ADMIN.EVENT)  TOPICOBJ (SYSTEM.ADMIN.EVENT)  +
  DEST (SYSTEM.ADMIN.SUBSCRIBED.EVENT)
```

MQ REST Administration

- Enabling further management options
 - Easy access from any language
 - Scriptable via curl
- Many MQSC commands have REST equivalent
 - Others supported via generic command
 - V9.1.3 adds true JSON-formatted generic commands
- Can manage older qmgrs via proxy qmgr

```
C:\Program Files\IBM\Latest902\bin>curl -k "https://localh
{"queue": [{
  "name": "Q.LOCAL",
  "status": {
    "currentDepth": 0,
    "lastGet": "",
    "lastPut": "",
    "mediaRecoveryLogExtent": "",
    "monitoringRate": "off",
    "oldestMessageAge": -1,
    "onQueueTime": {
      "longSamplePeriod": -1,
      "shortSamplePeriod": -1
    },
    "openInputCount": 0,
    "openOutputCount": 0,
    "uncommittedMessages": 0
  },
  "type": "local"
}]}
```

From V9.0.5 "What's New and Changed"

Version 9.0.5 introduces various improvements to the management and output of error logs. The main changes are that you can:

- Log diagnostic messages, using additional file services and syslog on UNIX platforms, as well as AMQERR01.LOG.

- Use JSON for the description of the messages, as well as the existing format; see JSON format diagnostic messages.

- Reformat a log into another language or style; see mqrc.

For more information, see Diagnostic message services, and QMErrorLog service.

https://www.ibm.com/support/knowledgecenter/en/SSFKSJ_9.0.0/com.ibm.mq.pro.doc/q130630_.htm#q130630_errlog

JSON Error (aka DiagnosticMessage) Logs

- In same directory as classic error files
- Files AMQERRxx.json
- Unix systems can also direct entries to syslog
 - Which has a lot of backends and routing options

```
{
  "ibm_messageId": "AMQ5051I",
  "arith_insert_2": 1,
  "comment_insert_1": "LOGGER-IO",
  "ibm_datetime": "2017-11-16T09:54:26.331Z",
  "ibm_serverName": "QM1",
  "type": "mq_log",
  "host": "machine.somewhere.ibm.com",
  "loglevel": "INFO",
  "module": "amqzmut0.c:1650",
  "ibm_sequence": "1510826066_332014693",
  "ibm_processId": 7846,
  "ibm_threadId": 4,
  "ibm_version": "9.0.4.0",
  "ibm_processName": "amqzmuc0",
  "ibm_userName": "somebody",
  "ibm_installationName": "Installation3",
  "ibm_installationDir": "/opt/mqm",
  "message": "AMQ5051I: The queue manager task 'LOGGER-IO' has started."
}
```

Consuming tools

- See, for example, Elasticsearch module
 - <https://www.elastic.co/guide/en/beats/filebeat/master/filebeat-module-ibmmq.html#filebeat-module-ibmmq>

How to configure syslog with MQ 9.1

- This example from AIX. Other Unix platforms will be similar
- In /etc/syslog.conf

```
# MQ writes to the "user" facility
user.debug /var/mqm/errors/syslog.log rotate size 1m files 4 compress
```

- In queue manager's qm.ini

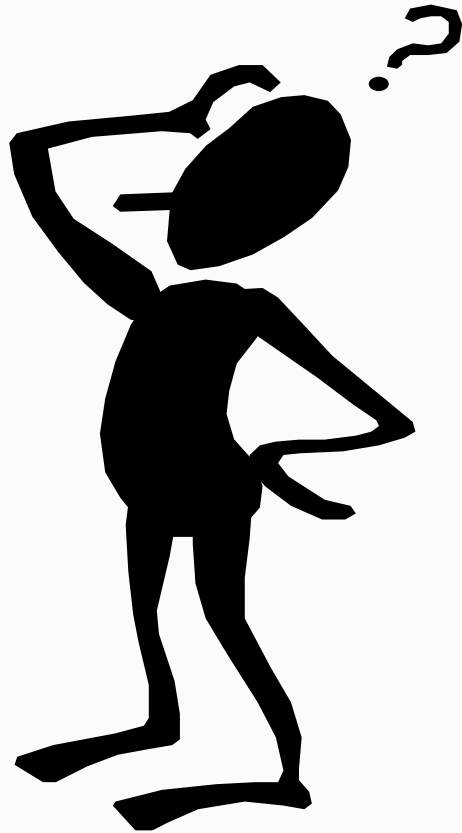
```
DiagnosticMessages:
  Name=DiagSyslog
  Service=Syslog
  Ident=mqseries
  Severities=I+
```

- Make sure syslog.log exists, then restart syslogd

```
Aug 14 15:58:50 example user:info mqseries: {"ibm_messageId":"AMQ9411I",
"ibm_arithInsert1":0, "ibm_arithInsert2":0,"ibm_datetime":"2018-08-14T14:58:50.250Z",
"ibm_serverName":"V9100_A", "type":"mq_log", "host":"example.hursley.ibm.com",
"loglevel":"INFO","module":"amqrrmfa.c:2108", "ibm_sequence":"1534258730_251676000",
"ibm_qmgrId":"V9100_A_2018-06-27_11.13.46", "ibm_version":"9.1.0.0", "ibm_processName":
"amqrrmfa", "ibm_userName":"metaylor", "ibm_installationDir":"/usr/mqm",
"message":"AMQ9411I: Repository manager ended normally."}
```

Summary

- MQ can be easily integrated with a variety of tools
- The pub/sub model for statistics makes it easy to add new consumers
 - Without disrupting any existing monitors
 - And makes it possible to add your own producers
- Using github for repository of code enables easy modification and sharing
github.com/ibm-messaging
- And blog posts for documenting what we have done
- Ability to use JSON as a common format for all operations



Any questions?