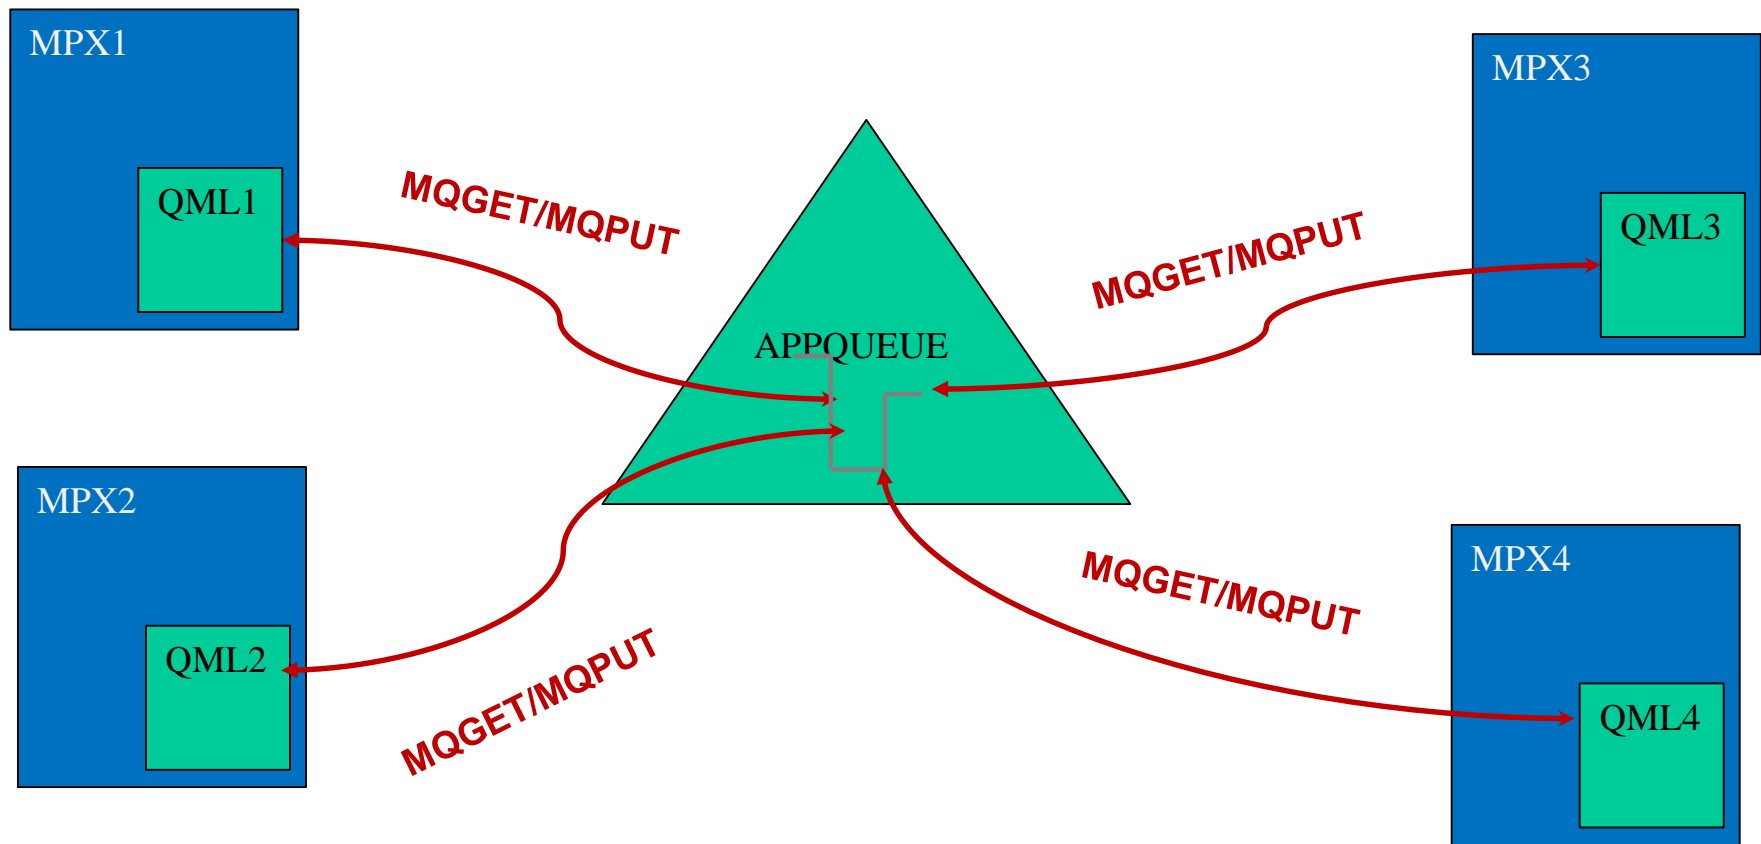


Shared Queues – where is my workload running

Agenda

- **What are shared queues – briefly**
- **What is workload skewing and why is it a problem?**
 - ▶ What are the symptoms and causes
 - Asymmetrical Sysplex
 - Connection Skewing
 - Put to Waiting Getter
 - ▶ ‘ Local’ favoritism
- **Mitigation Techniques:**
 - ▶ Queue Manager Clustering
 - ▶ Gateway queue managers
 - ▶ CICS CPSM options

Briefly – What is a Shared Queue




What is a shared queue?

- **Unique to z/OS**
 - ▶ Requires a coupling facility to host the queues
 - ▶ DB2 data sharing for the queue definitions
 - ▶ The gold standard for message availability
- **Treats a shared queues as local to each queue manager in the QSG**
 - ▶ Applications can PUT and GET
- **Messages are available as long as one QMGR in the QSG can access the Coupling Facility Structure.**
- **Nonpersistent messages are only 'lost' if the structure or CF itself are lost.**

What is MQ Workload Skewing?

- Workload skewing is detected when MQ driven work, typically transactions, is not close to being evenly distributed across the queue managers.



Why is MQ Workload Skewing a problem?

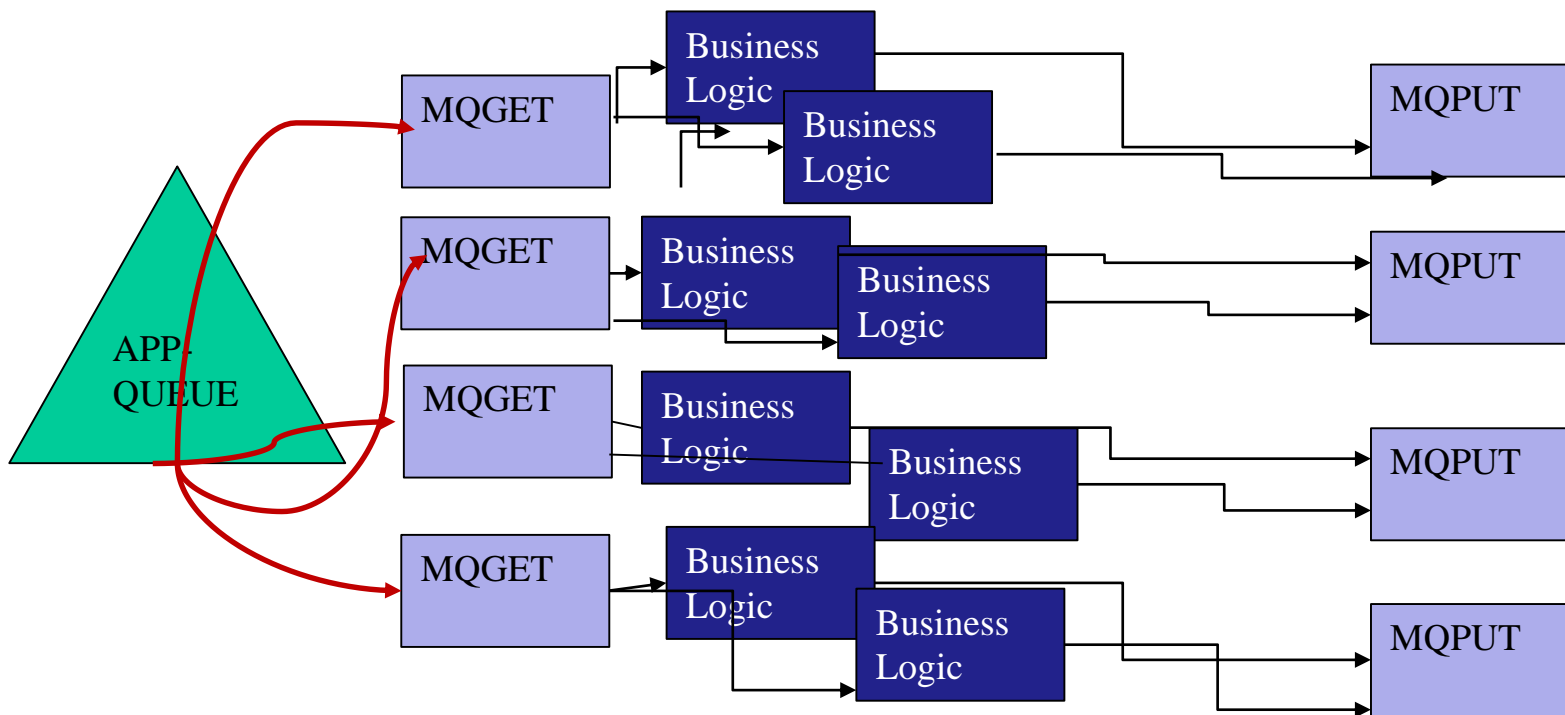
■ This is often less a technical problem, more of a pricing problem

- ▶ If the MLC 'rolling average' is taken from the LPAR that is heavily favored, usage pricing is not going to reflect reality
- ▶ Technical solutions to this problem may prove to be less efficient overall - lower throughput, slower response
- ▶ Using a VUE version of MQ can eliminate this issue

Why is MQ Workload Skewing a problem?

■ Can cause increased capacity demands in downstream workload

- ▶ Known to produce responsiveness problems
 - Overloading the processing programs
- ▶ Again this can contort MLC charges




MQ Workload Skewing Causes

- **Workload skewing in a QSG is often a result of the efficiencies of working locally**
 - ▶ z/OS, and all subsystems try to process requests locally to take advantage of CPU efficiency

MQ Workload Skewing Causes - Hardware

■ Asymmetric Sysplex


- ▶ When the LPARs in the Sysplex are not equally weighted
 - Examples include:
 - One LPAR is on an EC12, the others on older hardware
 - Two LPARs have 12 dedicated engines, two have 12 shared



MQ Workload Skewing Causes - Hardware

■ Asymmetric Sysplex

- ▶ Most common example - One LPAR is co-located with the primary coupling facility, the others are on different CPCs
- ▶ ICF links give much better service times than CBP



Physical Skewing – CF Activity Report

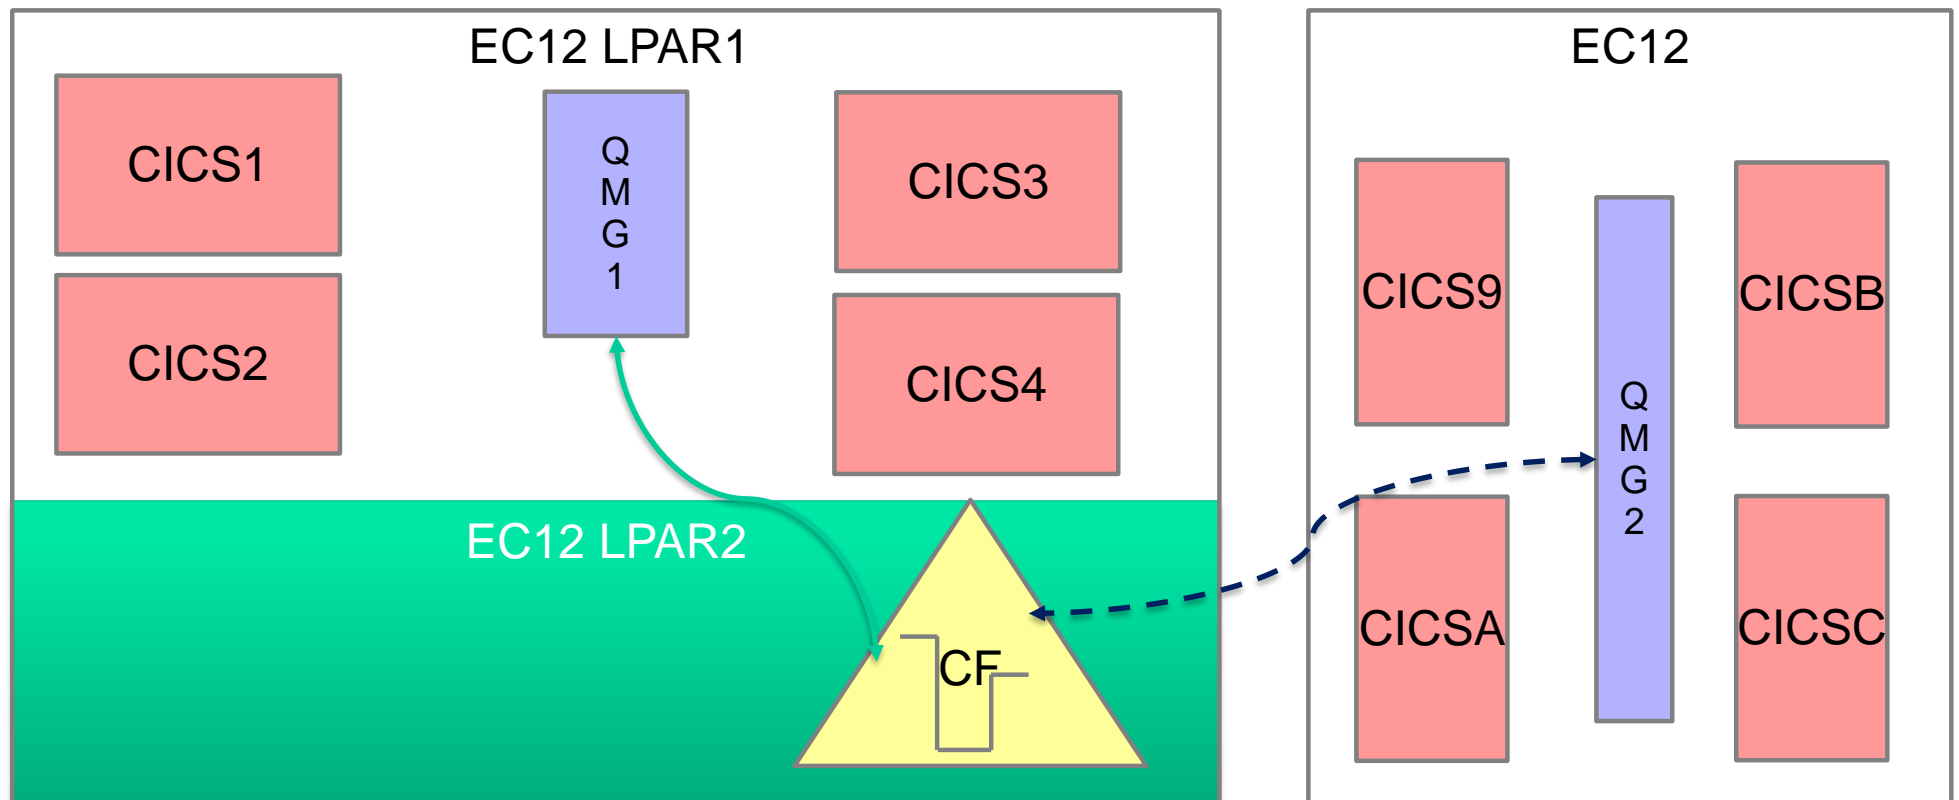
STRUCTURE NAME = QSGBUSER			TYPE = LIST			STATUS = ACTIVE			
SYSTEM	# REQ			REQUESTS					DE
NAME	TOTAL		#	% OF	-SERV TIME (MIC) -	REASON	#		%
	AVG/SEC		REQ	ALL	AVG	STD_DEV	REQ		R
MPX1	295K	SYNC	295K	26.9	4.3	1.2	NO SCH	0	0
	492.1	ASync	0	0.0	0.0	0.0	PR WT	0	0
		CHNGD	0	0.0	INCLUDED	IN ASync	PR CMP	0	0
		SUPPR	0	0.0			DUMP	0	0
MPX2	802K	SYNC	802K	73.1	17.8	2.5	NO SCH	0	0
	1339	ASync	0	0.0	0.0	0.0	PR WT	0	0
		CHNGD	0	0.0	INCLUDED	IN ASync	PR CMP	0	0
		SUPPR	0	0.0			DUMP	0	0

- We (the WSC) tend to use the CF Activity report rather than the MQ Statistics when looking at shared queue usage
- In the example shown above it is easy to see that the MPX2 LPAR is getting a much longer service time (almost 4 times!) than the MPX1 LPAR and that MPX2 is making many more requests.
 - In this particular case, this exposed some internal workload skewing that was not apparent to the customer - **except that they were missing SLAs consistently!**

MQ Workload Skewing Causes - Hardware

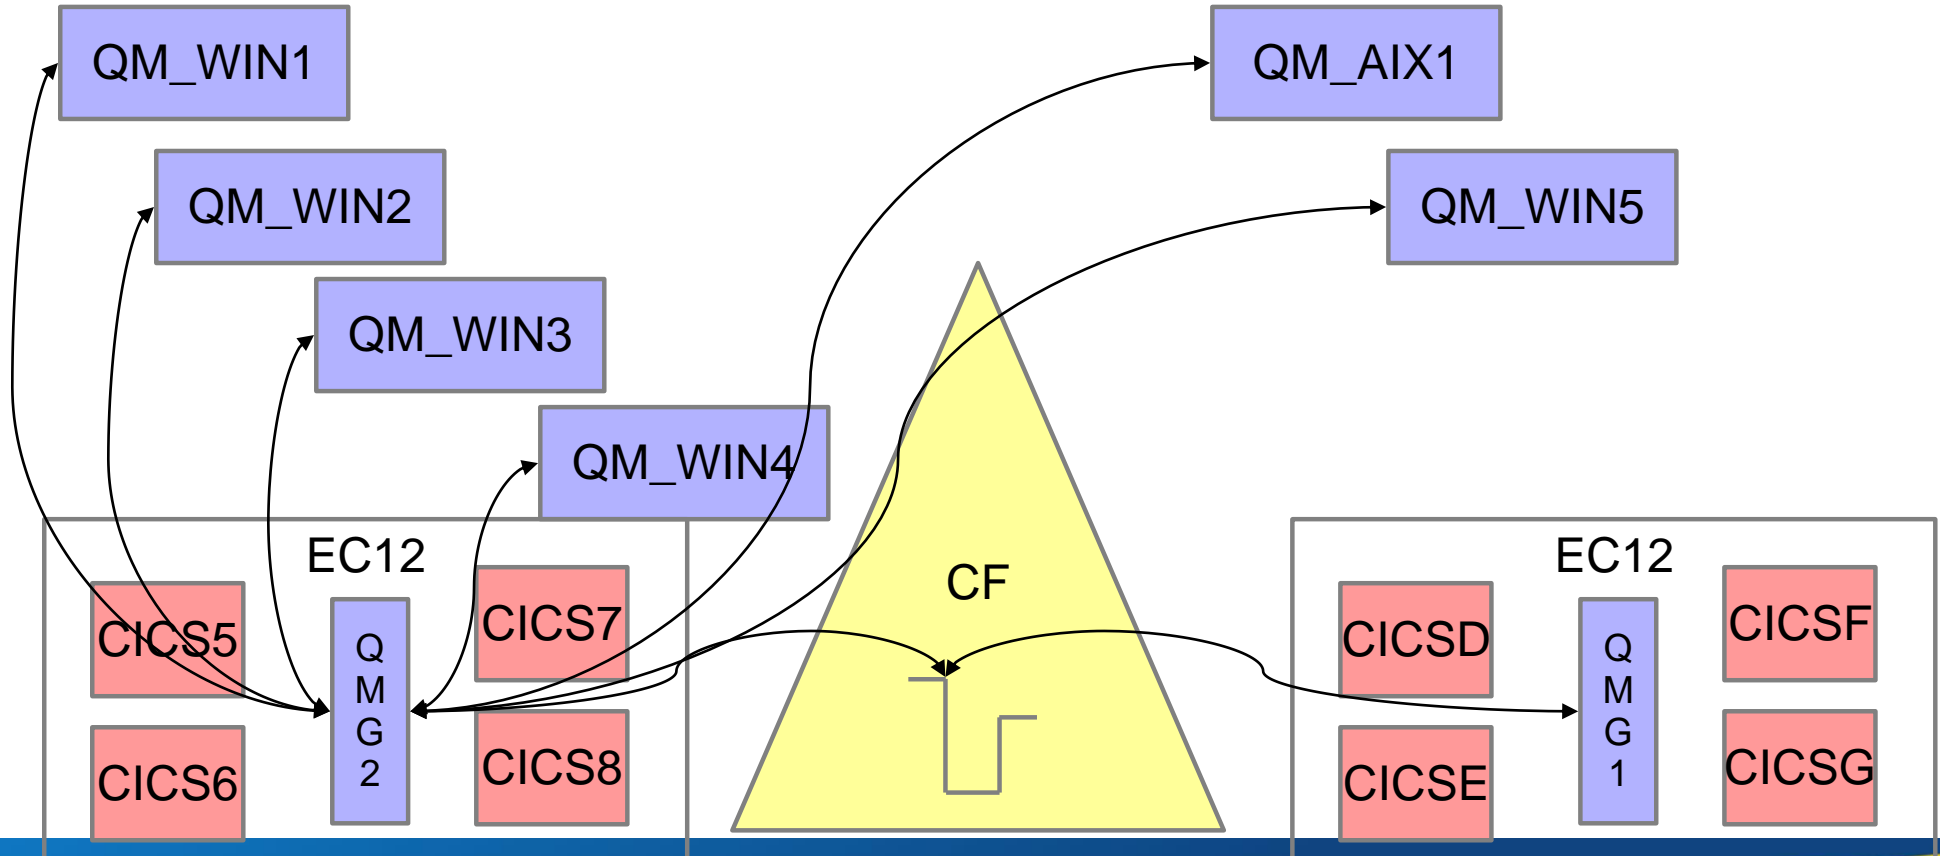
■ Location of the Coupling Facility

- ▶ When the coupling facility is internal, LPARs on the same CEC tend to get faster response
- ▶ When the coupling facility is external and one LPAR has more, faster, or less heavily used links it will get faster service



Connection Skewing

- **Connection skewing may be historical**
 - ▶ Hard-coded connections to specific queue managers
- **Connection skewing may be the result of a queue manager outage**
 - ▶ Connections to a QSG are routed to available queue managers



'Downstream' consequences

■ We've talked about the MLC impact

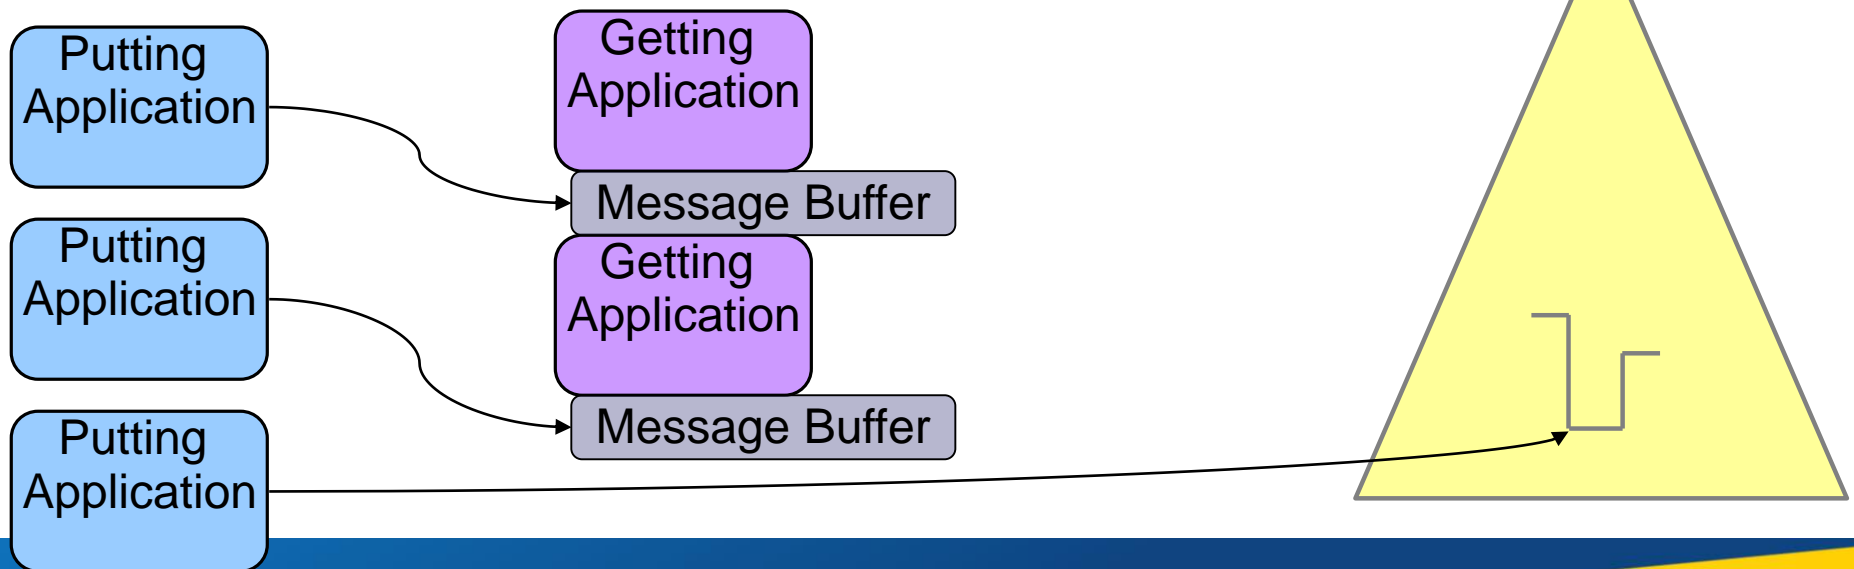
■ Resource use

- ▶ Not every queue manager is sized to absorb the entire workload
- ▶ Log impact of skewing has been seen
 - Rapid Log switches due to heavier workload – increasing I/O and CPU costs
- ▶ Bufferpool/Pageset impact
 - Filling the bufferpool, forced into I/O
- ▶ SMDS impact
 - One queue manager in QSG gets all offloaded messages

MQ Workload Skewing Causes

■ Put to waiting getter

- ▶ In V6 a performance feature was added called 'put to waiting getter'
- ▶ If a local put, from an application or message channel agent, is done and there is a getting application waiting the message is moved directly to the getting applications buffer
 - There is no posting to a shared queue
 - There is no notification to other available waiting applications
 - The CPU savings can be substantial
 - This works with connection skewing, and can maximize the effect



Put to Waiting Getter – SMF

- This shows messages flowing across a channel taking advantage of P2WG

Base_Name	CF Struct	Total_Val id_Gets	Total_Bytes _Put	Total_Val id_Puts	Total_Put2_Wa iting_Getter	Puts not to Waiting Getter
SYSTEM.QSG.CHANNEL.SYNCQ	CSQSYSAP	0	0	0	0	0
SHARED.INPUT.QUEUE	APP1	0	4501092223	2095814	2012394	83420

- The CPU comparison shows why it can be a good thing!

BASE_NAME	VALID_PUTS	PUT_ELAPSE D_TIME	PUT_CPU_TIME	PUT2_W AITING_G ETTER	Average PUT ET	Average PUT CT
QLOCAL.PUT2WG	14879	127753	117956	14793	8.59	7.93
QLOCAL.NO.PUT2WG	41547	1025028	1010038	0	24.67	24.31

- The CPU costs can be 3 times as high!

MQ Workload Skewing Causes

■ Local Favoritism

- ▶ When a message is posted to a shared queue, the queue manager where the message is put is typically notified FIRST about the availability.
- ▶ Normal processing by XCF, taking advantage of the efficiency of local processing.

Skewing Mitigation Techniques

■ Queue Manager Clusters

- ▶ Clusters provide workload balancing across queue managers
- ▶ Works with shared queues to distribute message 'puts' across queue managers in the QSG


■ Connection skewing mitigation

- ▶ Gateway queue managers
- ▶ Re-driving connections

■ CPSM mitigation

Queue Manager Clustering

- When messages are not bound to a specific queue manager ('bind not fixed'), the messages are routed evenly across the receiving queue managers
 - Black arrows show the first message put to the clustered queue
 - Green arrows show the second message

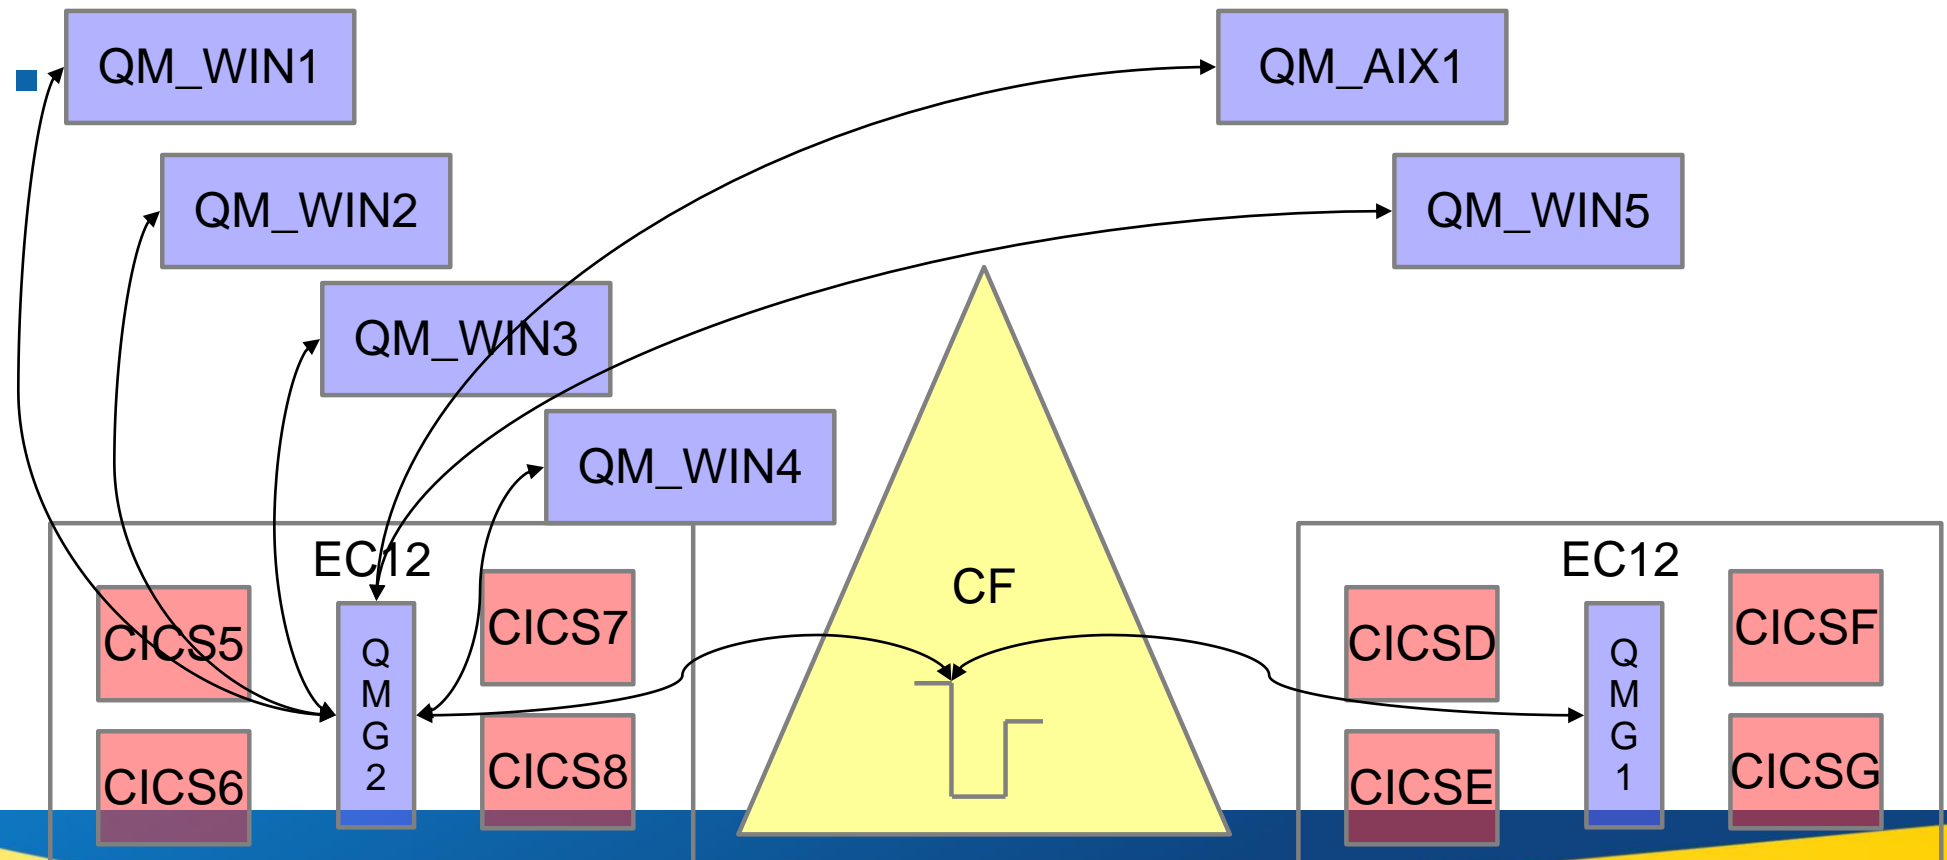


Connection Skewing Mitigation


- The slides that follow outline two mitigation techniques for connection skewing:
 - ▶ Gateway queue managers
 - ▶ Re-driving connections

Connection Skewing – No Gateway queue managers

- When external queue managers or clients are passing work directly to application hosting queue managers, every attempt is made to process the work locally
- Environments that use gateway queue managers into the Queue Sharing group often eliminate connection skewing.

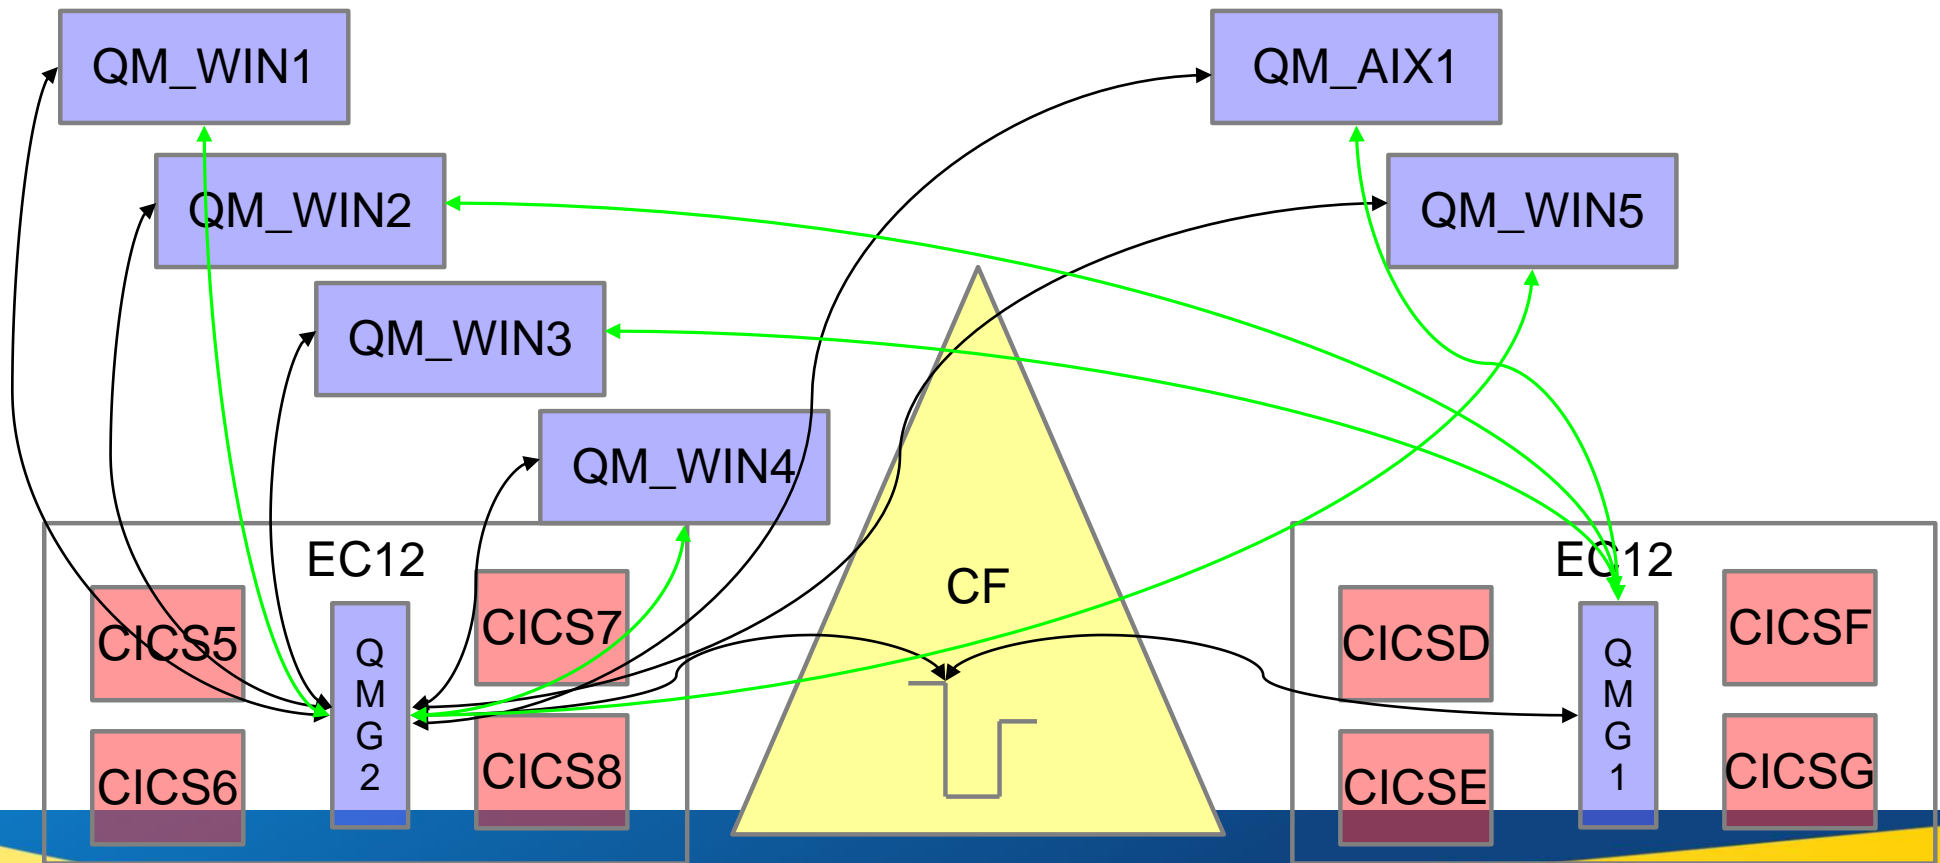


Gateway queue managers – the mitigation



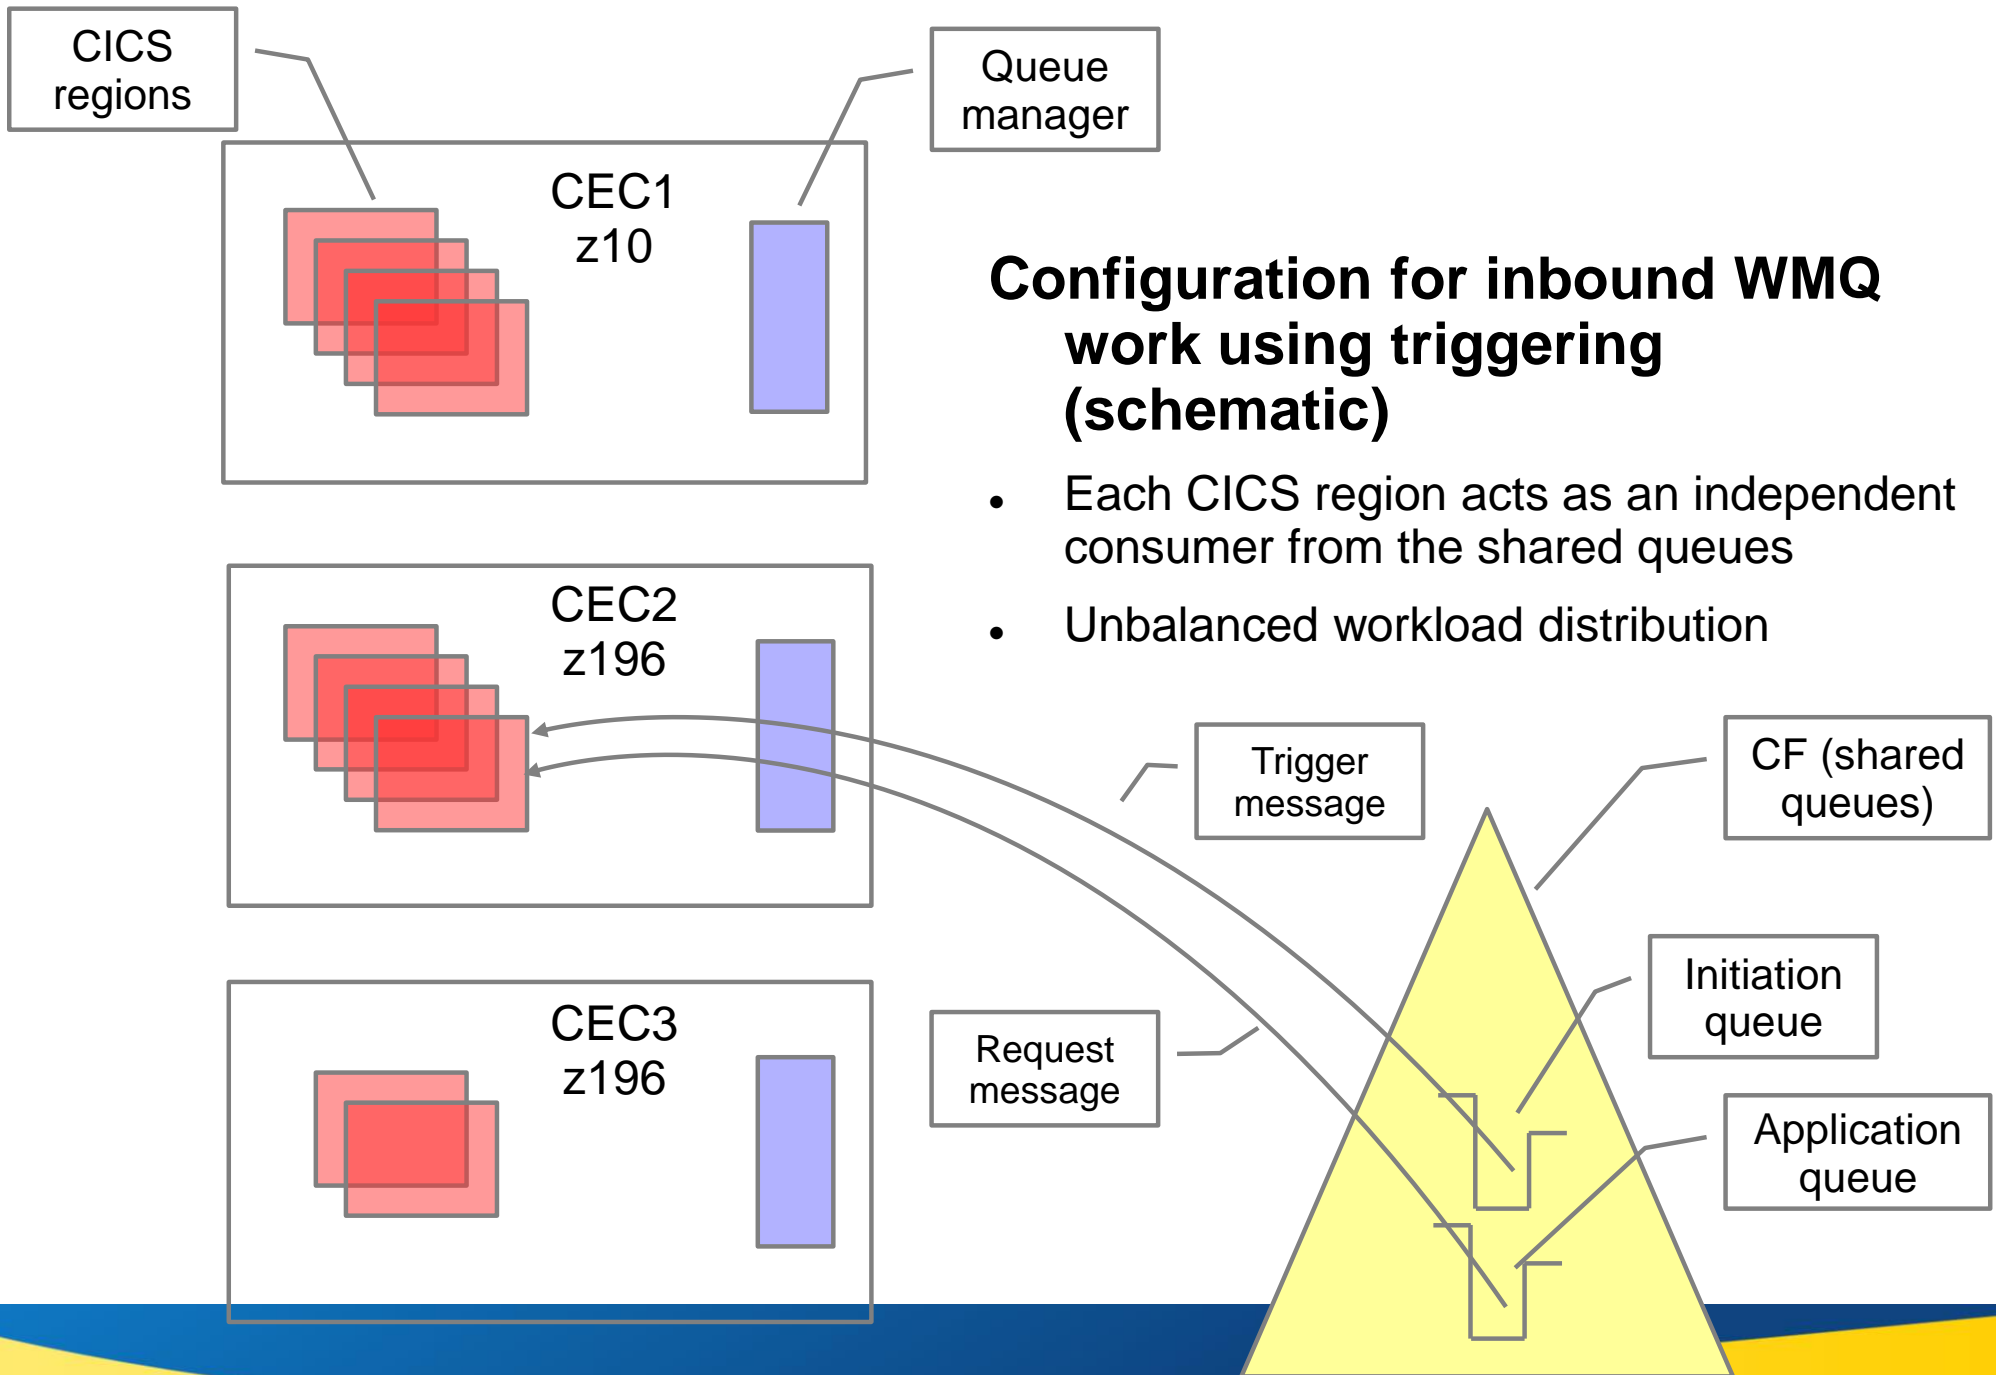
Re-driving Connections

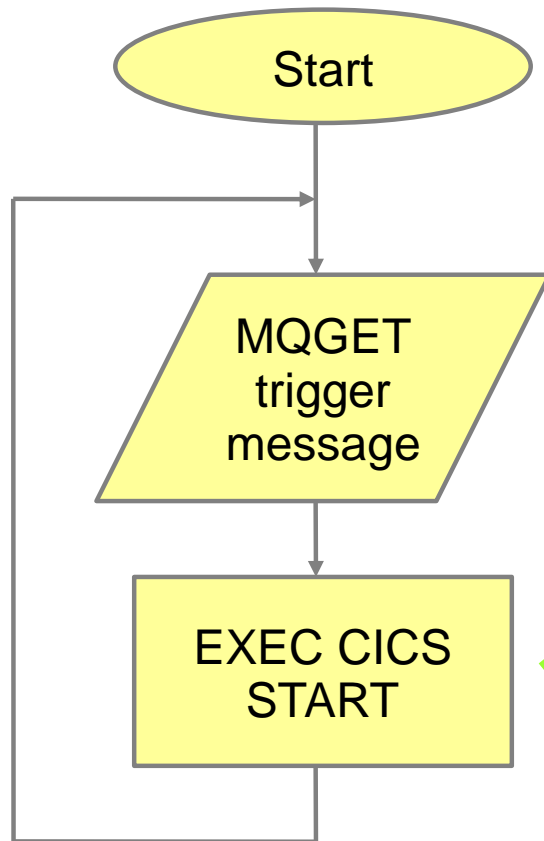
- When a queue manager is unavailable, inbound connections can get skewed to the other queue manager(s) in the group.
 - ▶ This is normal availability processing!
 - ▶ Once a connection is live and active, no attempt is made to balance the connections once all the queue managers are available.



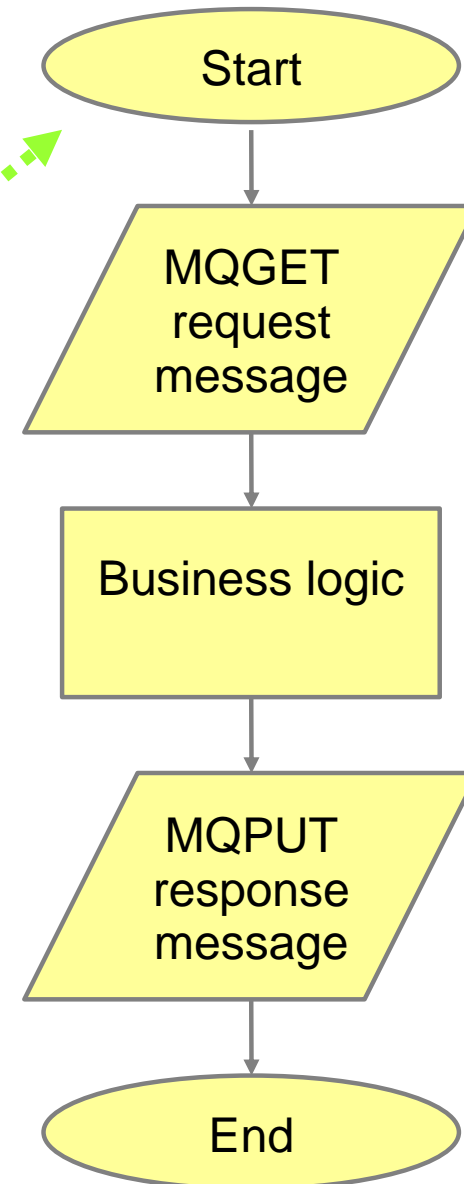
CICS – CPSM Mitigation

- The slides that follow outline a CPSM solution to the skewing problem based on the interaction between MQ triggering (CKTI) and CICS



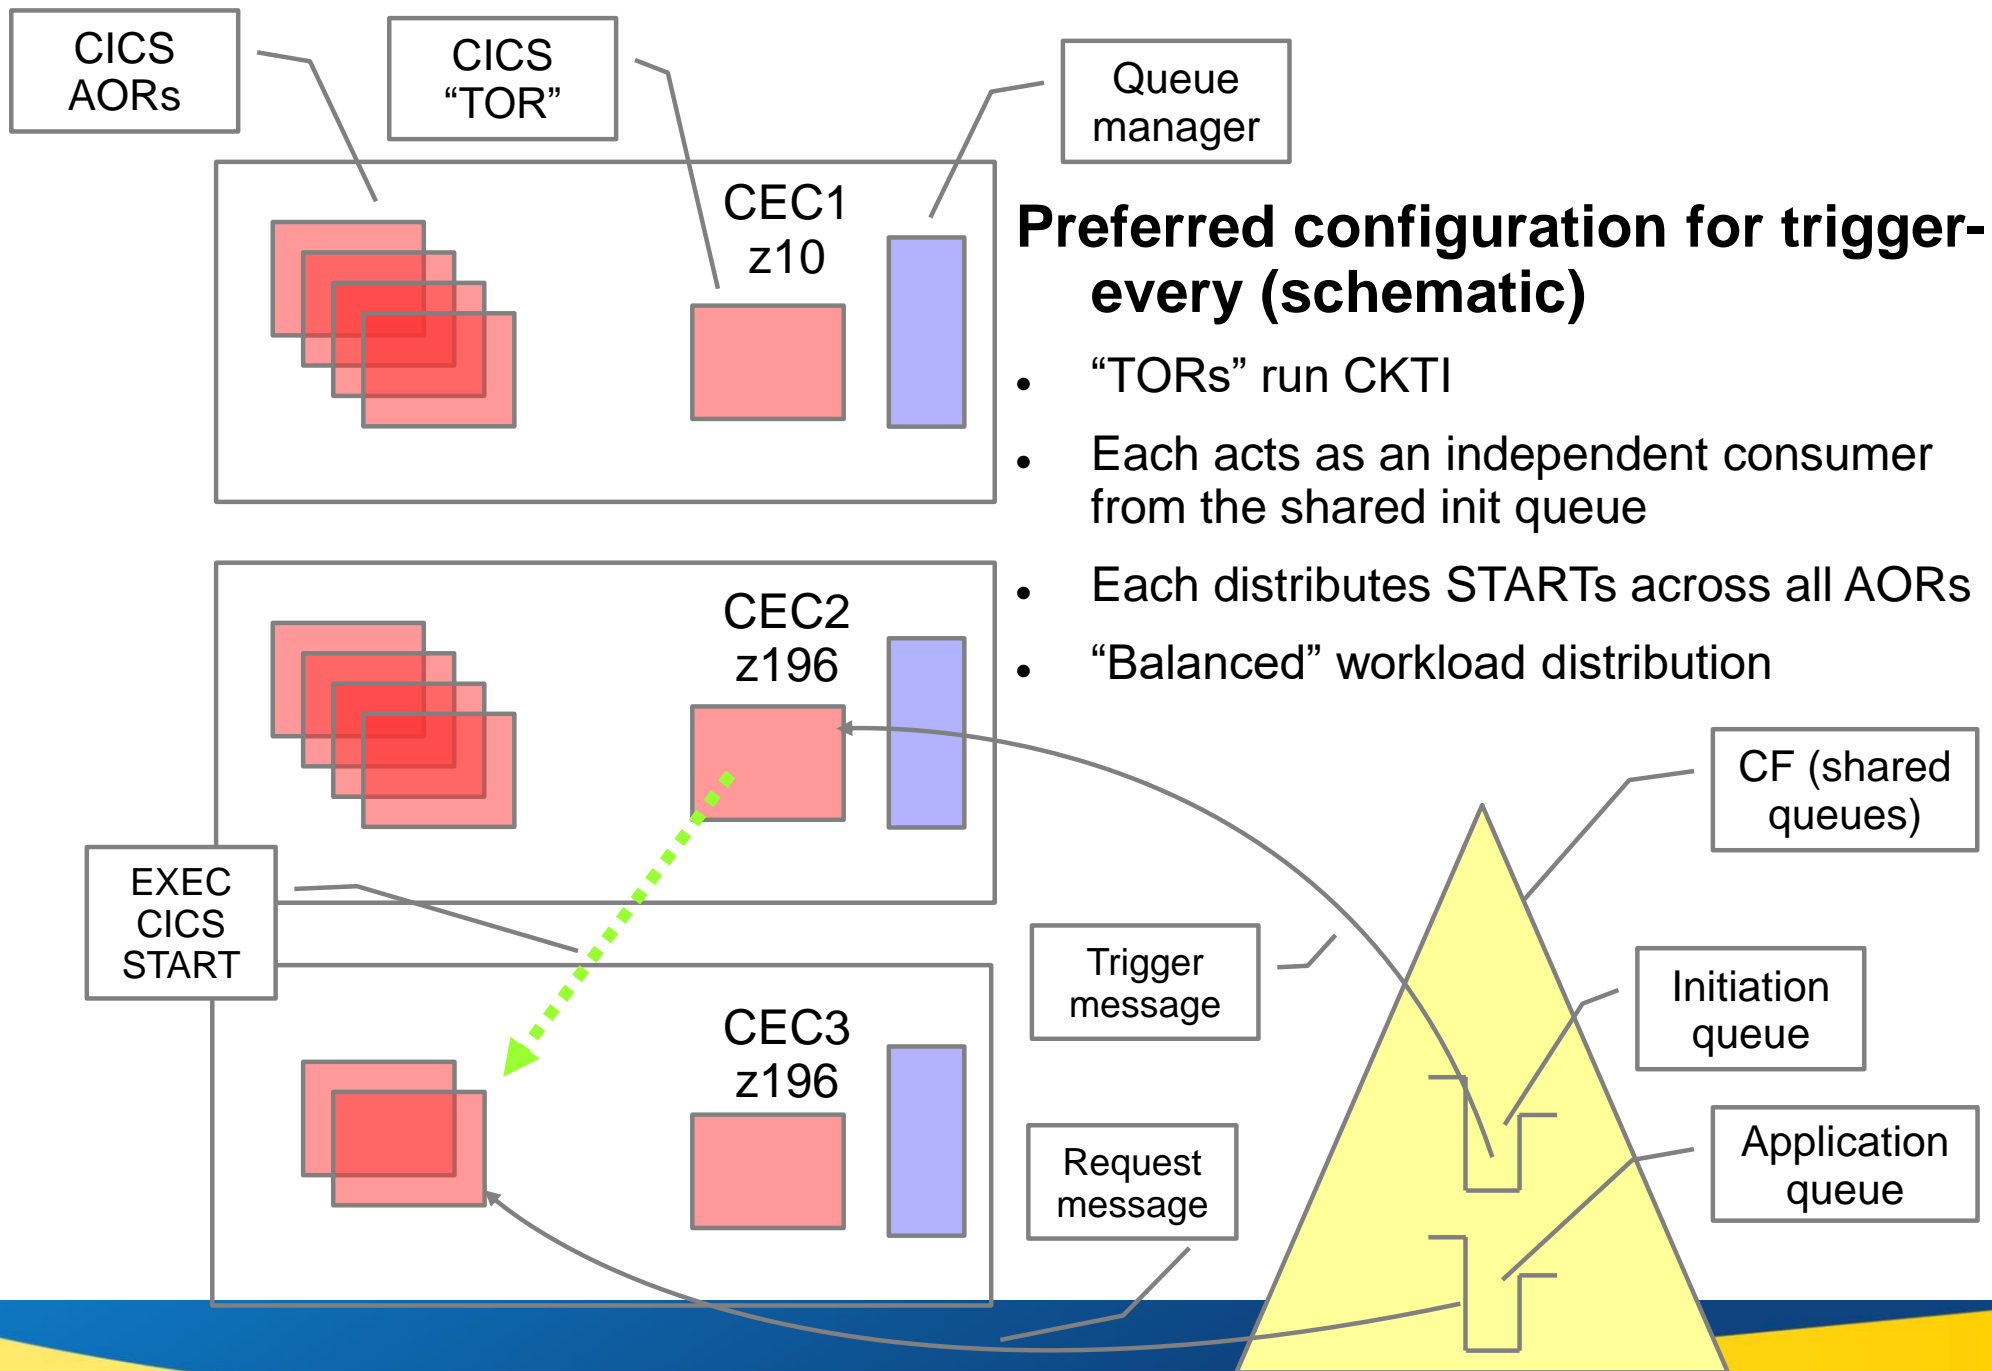


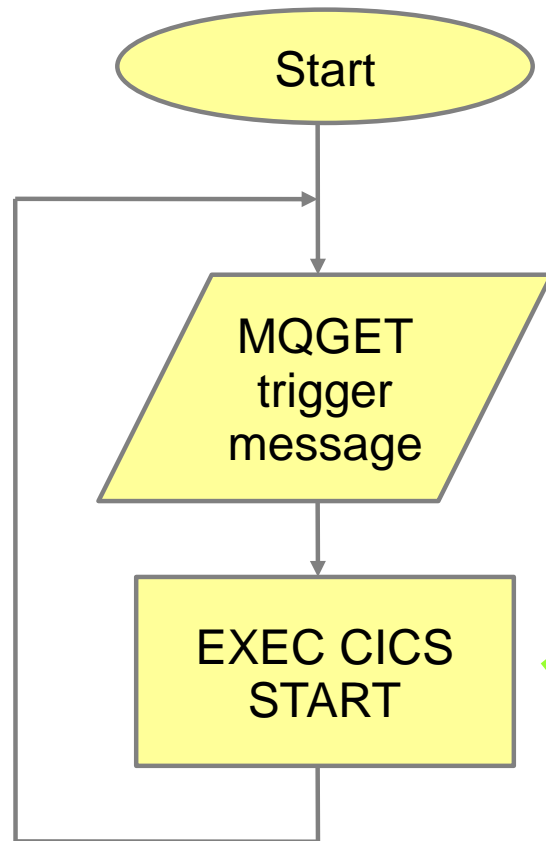
Business transaction




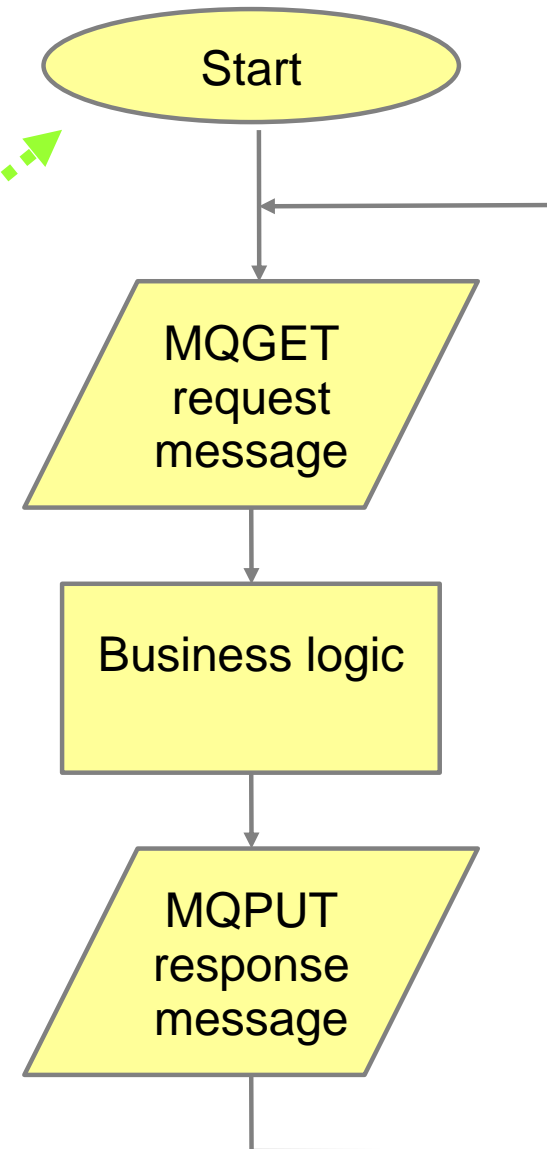
Trigger-every (schematic)

- Each business transaction processes few (~1) request messages
- Fastest CKTIs take lion's share of work





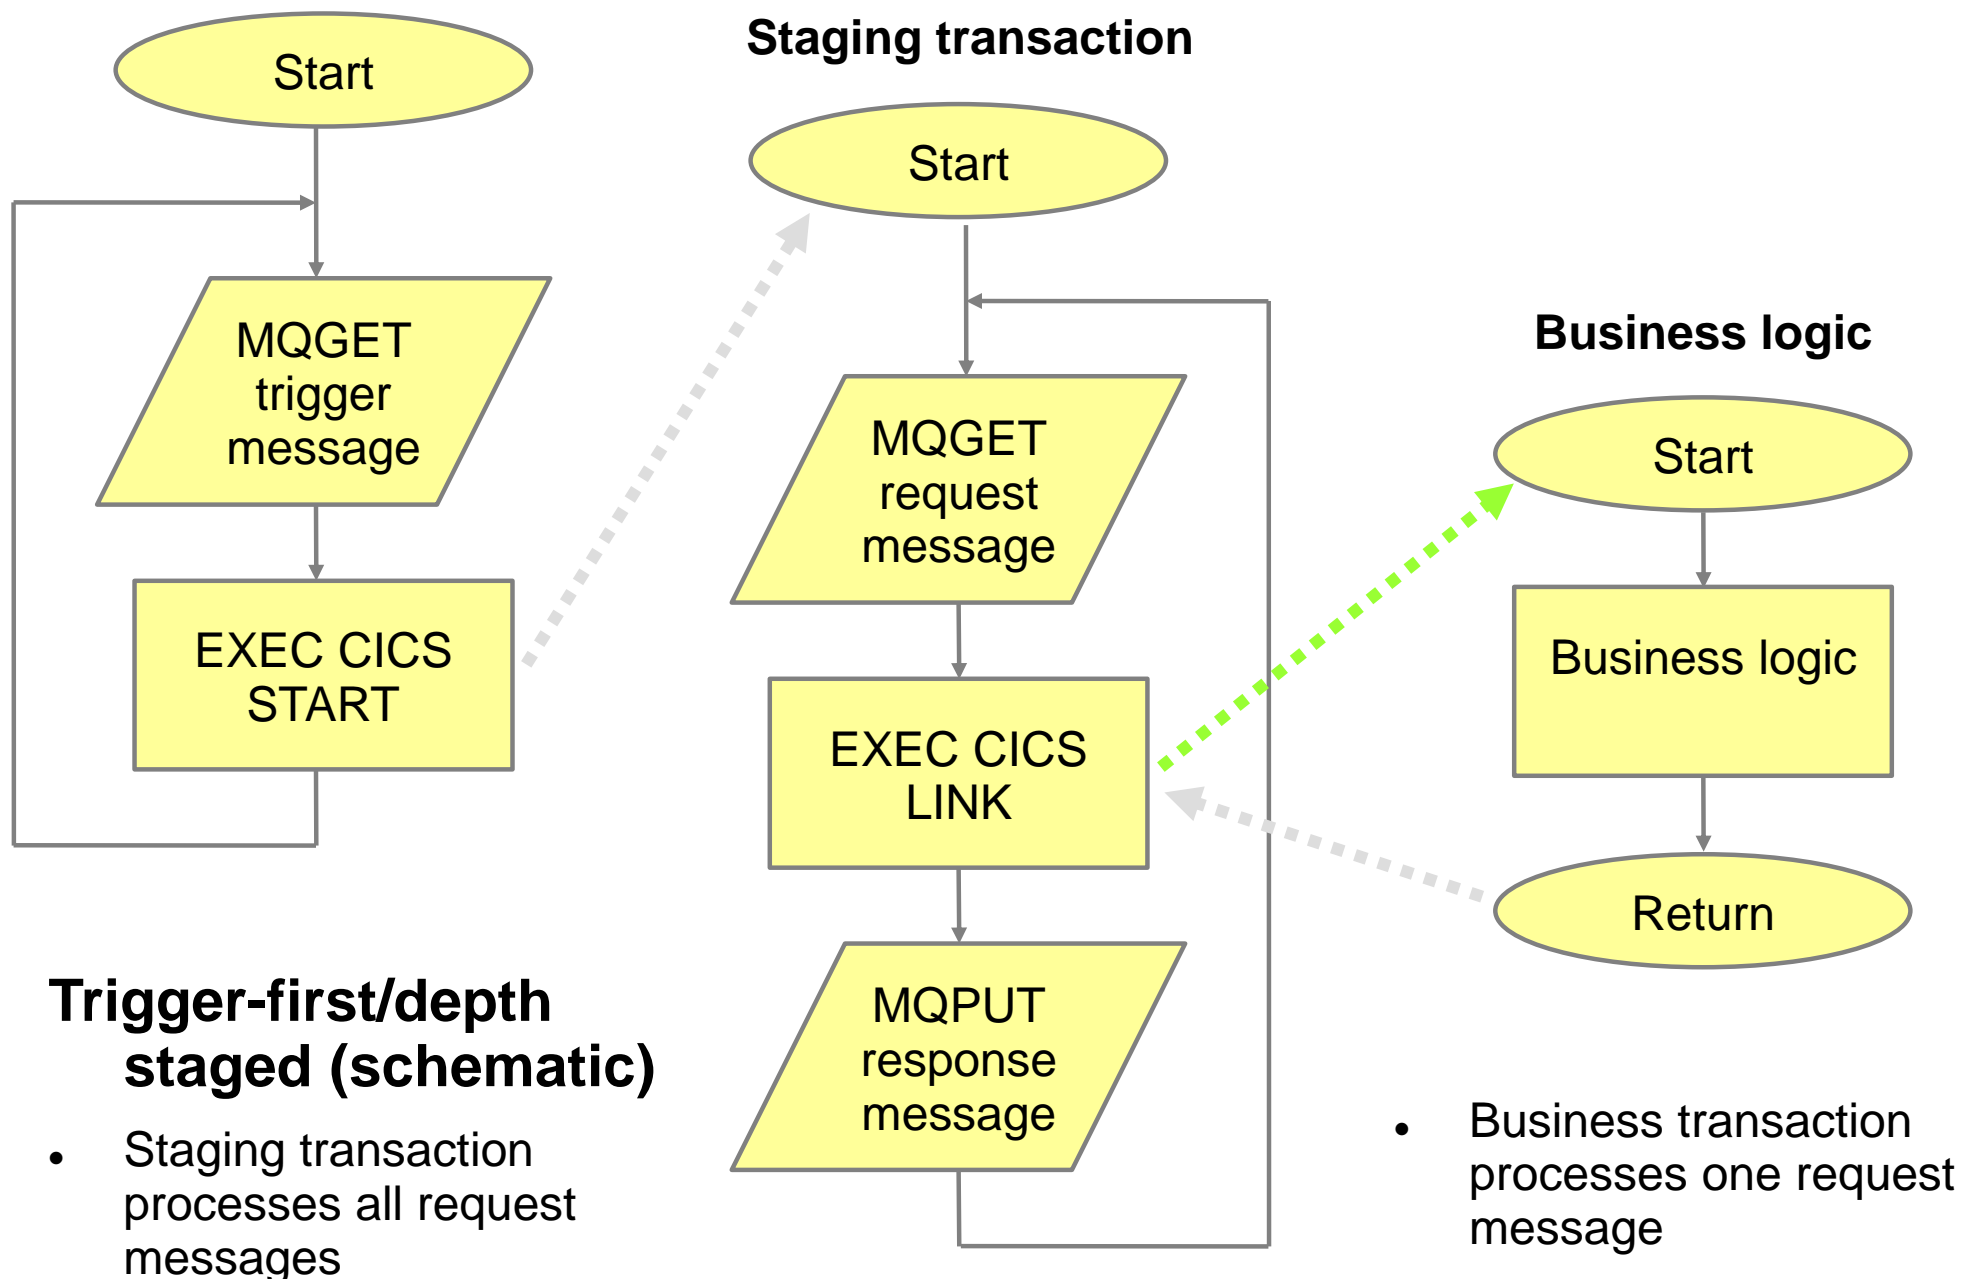
Business transaction




Trigger-first/depth (schematic)

- Each business transaction processes many (~all) request messages
- Fastest CKTI takes lion's share of work
- Corresponding business transaction takes lion's share of the work

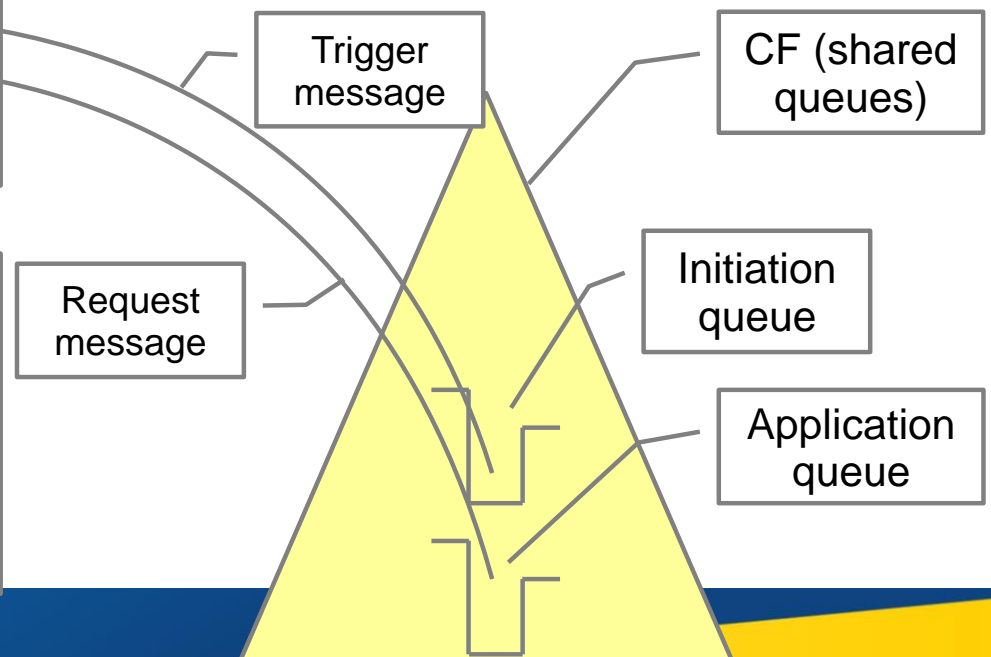
Trigger monitor (CKTI)





Preferred configuration for trigger-first/depth staged (schematic)

- "QORs" run CKTI and staging transaction
- Each acts as an independent consumer from the shared init queue and shared application queue
- Each distributes LINKs across all AORs
- "Balanced" workload distribution



Highlights

■ **Solution uses proven technology for CPSM routing:**

- Each TOR/QOR uses link-neutral goal algorithm
 - Selects target AOR based on AOR load and health
 - Does not “prefer” local (= same LPAR) AORs
 - Even distribution across AORs, but ...
 - ... responds to transient load/health variation
- XCF MRO for “remote” STARTs or LINKs
 - High-performance System z sysplex technology
 - Uses coupling facility (CF) instead of TCP/IP stack
- Sysplex-optimised workload routing
 - Highly responsive to transient variations
 - Uses CF to maintain current status for AORs

□ **Continuous operation and high availability through WMQ shared queues:**

- “Glitchless” recovery from region/LPAR/CEC outage
- “Instant” redistribution of workload
- In-flight messages backed-out, restart in another CICS region

□ **High throughput:**

- Exploits all available capacity
- Highly responsive to transient spare capacity

MQ Workload Balance Summary

- **MQ is a message delivery system, it does not try to balance workload**
- **Balancing the workload is attempting a technical solution for what is often a pricing problem**
 - ▶ Beware spending a lot of effort for a solution to a temporary problem as well!
 - ▶ Turning off performance improvements like put to waiting getter will impact all applications, not just the skewed ones
- **There are some mitigation techniques that can help the overall environment**
 - ▶ Clustering!
 - ▶ Gateway queue managers
 - ▶ Using CPSM to make appropriate routing decisions

Additional Resources

- **The following links are to additional information about WMQ**

- ▶ Queue Sharing Groups:

- http://publib.boulder.ibm.com/infocenter/wmqv7/v7r1/topic/com.ibm.mq.explorer.doc/e_qsg.htm

- ▶ Clustering:

- http://publib.boulder.ibm.com/infocenter/wmqv7/v7r1/topic/com.ibm.mq.doc/qc11220_.htm

- ▶ Intercommunication

- http://publib.boulder.ibm.com/infocenter/wmqv7/v7r1/topic/com.ibm.mq.doc/zx00011_.htm

- ▶ Redbooks:

- IBM WebSphere MQ V7.1 and V7.5 Features and Enhancements

- <http://www.redbooks.ibm.com/abstracts/sg248087.html?Open>

- High Availability in WebSphere Messaging Solutions

- <http://www.redbooks.ibm.com/abstracts/sg247839.html?Open>

- WebSphere MQ Queue Sharing Group in a Parallel Sysplex environment (dated, but still good basic information)

- <http://www.redbooks.ibm.com/redpieces/abstracts/redp3636.html?Open>

- ▶ Lyn's first YouTube video:

- <http://www.youtube.com/playlist?list=PL9N7JP2yU3T8JycrCOvEPM8c-0UdE97VT>

MQ Workload Balance - thanks

■ Many thanks to

- ▶ Steve Hobson for the CICS/CPSPM expertise and the wonderful graphics
- ▶ Mark Taylor for their patience and guidance on the rest of the foils
- ▶ Mark Taylor for providing the excellent editing and recording studio for the YouTube version of this pitch