


Integrating MQ monitoring data with common cloud and open source tooling

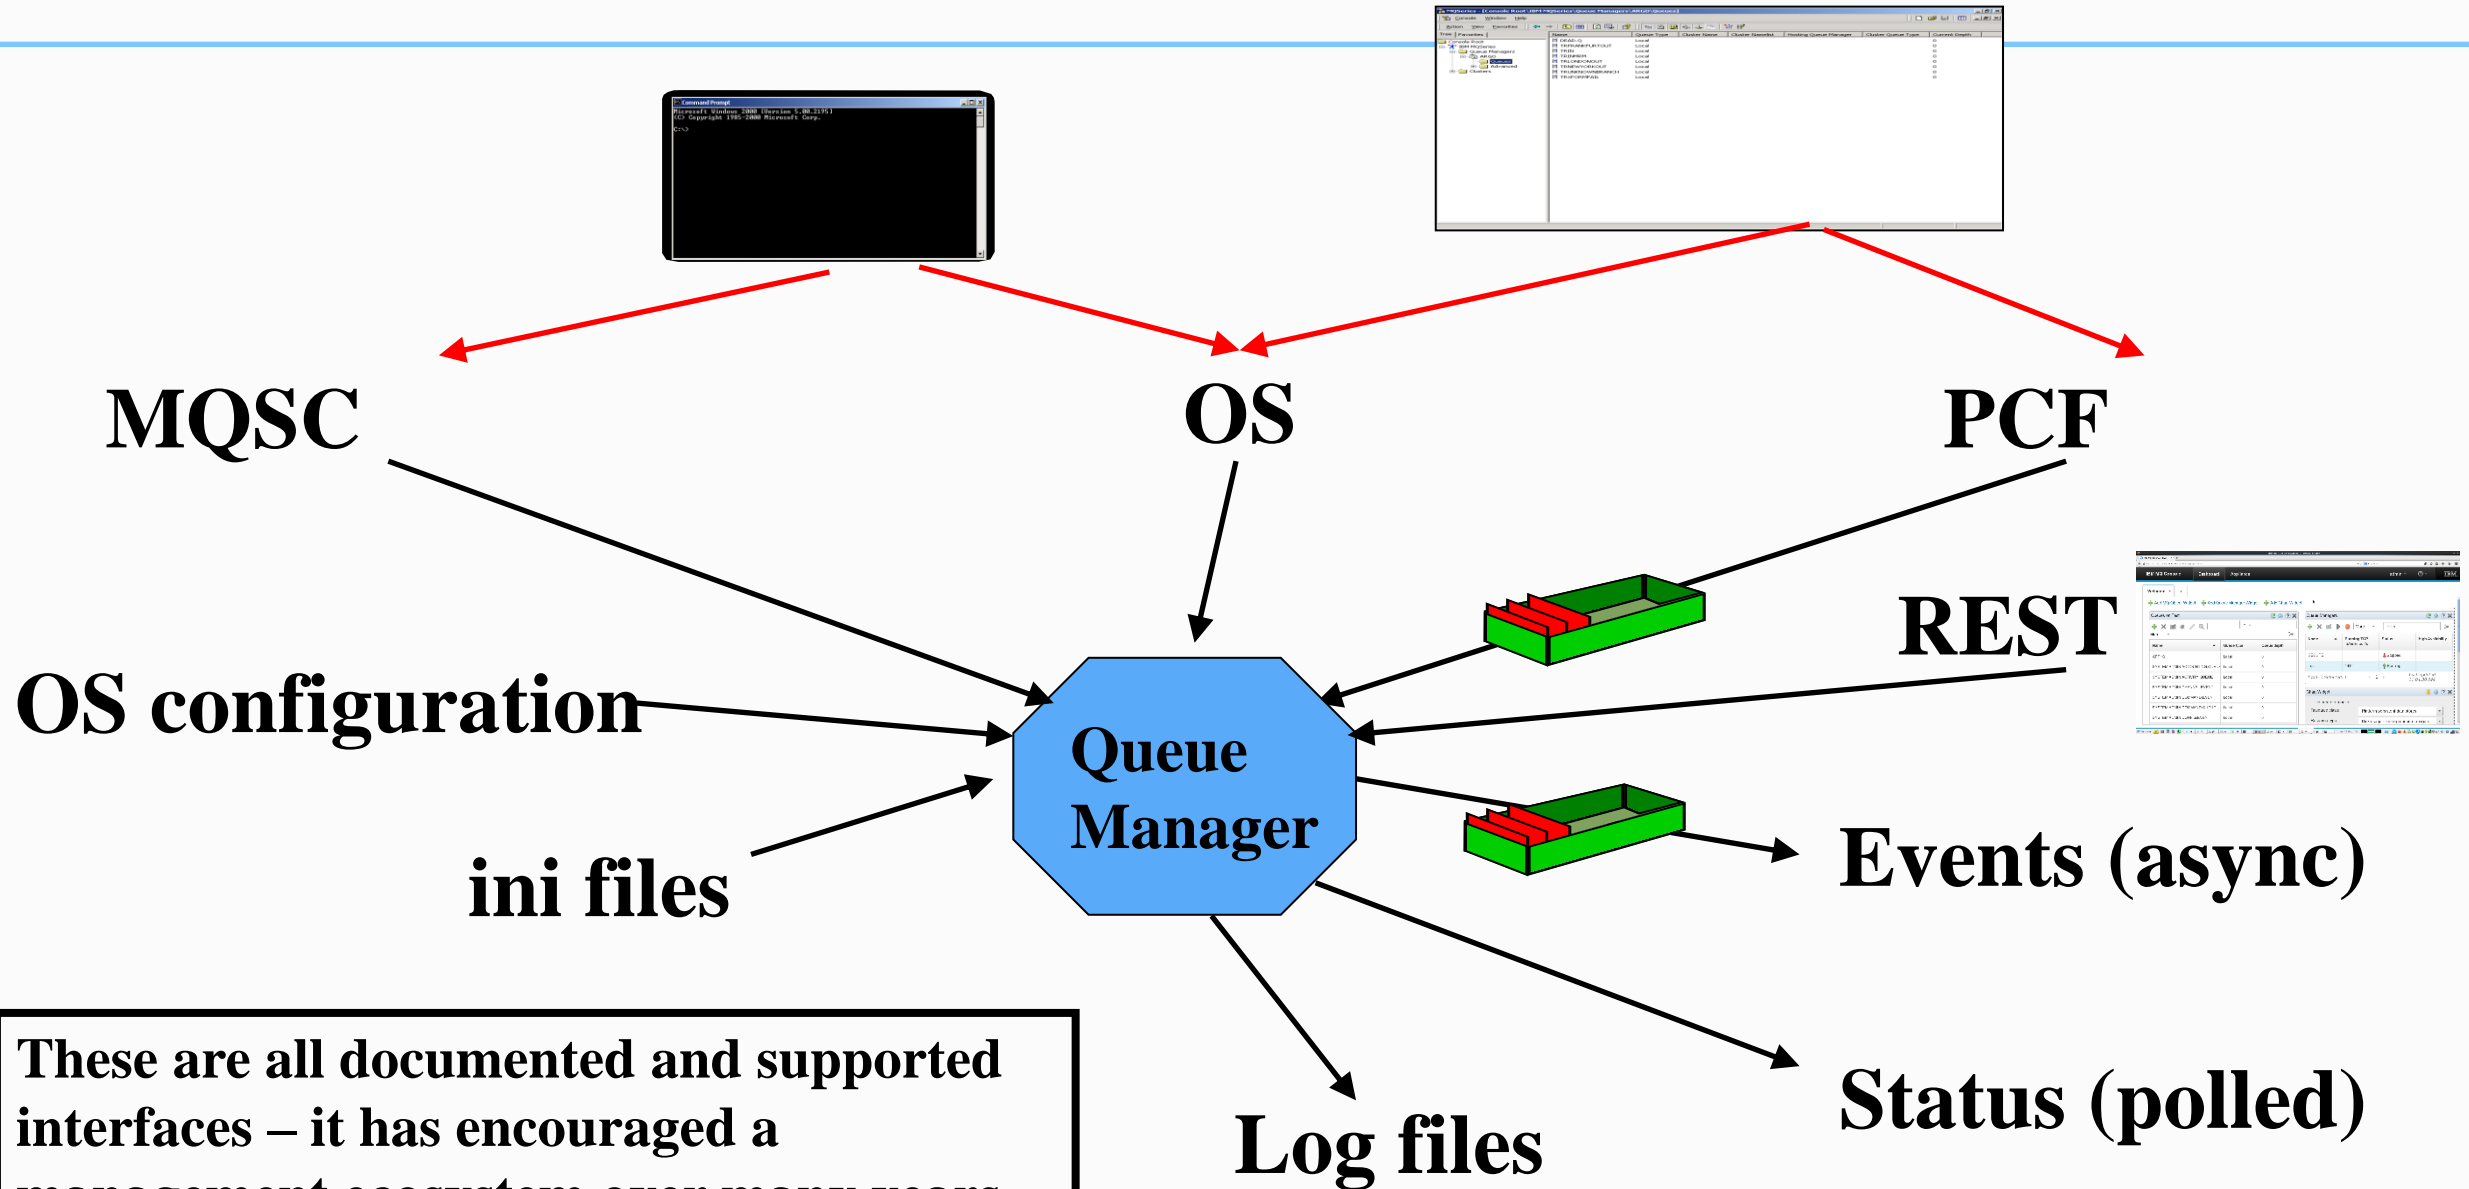


Mark Taylor

marke_taylor@uk.ibm.com

IBM Hursley

MQ Administration



These are all documented and supported interfaces – it has encouraged a management ecosystem over many years

IBM MQ - MQSC

Old Example (1995): AIX smit panels

- Command line interface
- V8 enhanced runmqsc
 - Make it world-executable
 - Enable direct client-connection
- MQSC intended for human consumption
 - Parsable by eye, less easy in programs
 - For example, **DESCR('This is 'a' description with quote & paren(')**
 - No guaranteed ordering in runmqsc, two-column output
- Despite awkwardness, basis for many script-based admin tools
 - echo "DISPLAY Q(X) IPPROCS" | runmqsc QM1
- Same commands – different front-end (CSQUTIL) – for z/OS

IBM MQ - PCF

- A "self-describing" MQ message used for administrative operations
- Your programs can send commands and get responses using PCF
 - Equivalent to "DISPLAY QSTATUS" or "ALTER CHANNEL"
- MQ emits events in PCF format
 - "Queue is getting full"
- PCF intended for programs – usually C or Java
 - Can tell exactly what the parameter is for, its length and value
 - But cannot easily be scripted
- Approximately one-one mapping between MQSC commands and PCF
- Remember that PCF invented before formats like JSON or XML
 - And there are many MQ apps that are built on PCF

An event message

**** Message length - 300 of 300 bytes ***

```
00000000: 0000 0007 0000 0024 0000 0003 0000 0063 '.....$......c'
00000010: 0000 0001 0000 0001 0000 0000 0000 096C '.....1'
00000020: 0000 0002 0000 0014 0000 0010 0000 1F41 '.....A'
00000030: 0000 0004 0000 0004 0000 0020 0000 0BE5 '.....å'
00000040: 0000 0333 0000 000C 6D65 7461 796C 6F72 '...3...metaylor'
00000050: 2020 2020 0000 0003 0000 0010 0000 03F3 '.....ó'
00000060: 0000 0001 0000 0004 0000 0044 0000 0BE7 '.....D...ç'
00000070: 0000 0333 0000 0030 5638 3030 335F 4120 '...3...0V8003_A '
00000080: 2020 2020 2020 2020 2020 2020 2020 2020 ' '
00000090: 2020 2020 2020 2020 2020 2020 2020 2020 ' '
000000A0: 2020 2020 2020 2020 0000 0003 0000 0010 '.....'
000000B0: 0000 03FD 0000 005A 0000 0014 0000 0010 '...ý...Z.....'
000000C0: 0000 1F42 0000 0004 0000 0004 0000 0018 '...B.....'
000000D0: 0000 0BFB 0000 0000 0000 0001 5800 0000 '...û.....X...'
000000E0: 0000 0003 0000 0010 0000 03F8 0000 0001 '.....ø.....'
000000F0: 0000 0006 0000 0024 0000 0BF9 0000 0000 '.....$....ù....'
00000100: 0000 0001 0000 0008 6D65 7461 796C 6F72 '.....metaylor'
00000110: 0000 0000 0000 0005 0000 0018 0000 045C '.....\'
00000120: 0000 0002 0000 000B 0000 0009 '.....'
```

An event message

**** Message length - 300 of 300 bytes ***

00000000:	0000	0007	0000	0024	0000	0003	0000	0063	'.....\$......c'
00000010:	0000	0001	0000	0001	0000	0000	0000	096C	'.....1'
00000020:	0000	0002	0000	0014	0000	0010	0000	1F41	'.....A'
00000030:	0000	0004	0000	0004	0000	0020	0000	0BE5	'.....å'
00000040:	0000	0333	0000	000C	6D65	7461	796C	6F72	'...3...metaylor'
00000050:	2020	2020	0000	0003	0000	0010	0000	03F3	'.....ó'
00000060:	0000	0001	0000	0004	0000	0044	0000	0BE7	'.....D...ç'
00000070:	0000	0333	0000	0030	5638	3030	335F	4120	'...3...0V8003_A '
00000080:	2020	2020	2020	2020	2020	2020	2020	2020	'.....'
00000090:	2020	2020	2020	2020	2020	2020	2020	2020	'.....'
000000A0:	2020	2020	2020	2020	0000	0003	0000	0010	'.....'
000000B0:	0000	03FD	0000	005A	0000	0014	0000	0010	'...ý...Z.....'
000000C0:	0000	1F42	0000	0004	TYPE (cfst) LEN (24)				'...B.....'
000000D0:	PARM (MQCA...)		CCSID (0)		LEN (1)		DATA		'...û.....X...'
000000E0:	0000	0003	0000	0010	0000	03F8	0000	0001	'.....ø.....'
000000F0:	0000	0006	0000	0024	0000	0BF9	0000	0000	'.....\$....ù.....'
00000100:	0000	0001	0000	0008	6D65	7461	796C	6F72	'.....metaylor'
00000110:	0000	0000	0000	0005	0000	0018	0000	045C	'.....\'
00000120:	0000	0002	0000	000B	0000	0009			'.....'

Event formatting C sample introduced in V8

- No sample previously shipped to format all "standard" events
 - Authorisation, queue full, service interval, command/config etc
 - Other product samples **are** available for acct/stats, activity reports
 - Several SupportPacs but product only has out-of-date source code in the KC
- The **amqsevt** program formats events into readable English-ish text
 - Option to stay with full MQI constant name instead of making it look nice
 - Uses MQCB to read from multiple event queues. No polling required
 - Can connect as client to any remote queue manager including z/OS
 - Source code included
- Includes C header file to help convert MQI numbers to strings
 - Similar to Java MQConstants.lookup() capability for all sets of constants

```
printf("Error is %s\n",MQRC_STR(2035));
```

An event message decoded

Event Type : Command Event
Reason : Command MQSC
Event created : 2015/06/03 13:28:20.51 GMT
Correlation ID : 414D512056383030335F412020202020556F00F120001E05

COMMAND CONTEXT


Event User Id : metaylor
Event Origin : Console
Event Queue Mgr : V8003_A
Command : Set Auth Rec

COMMAND DATA

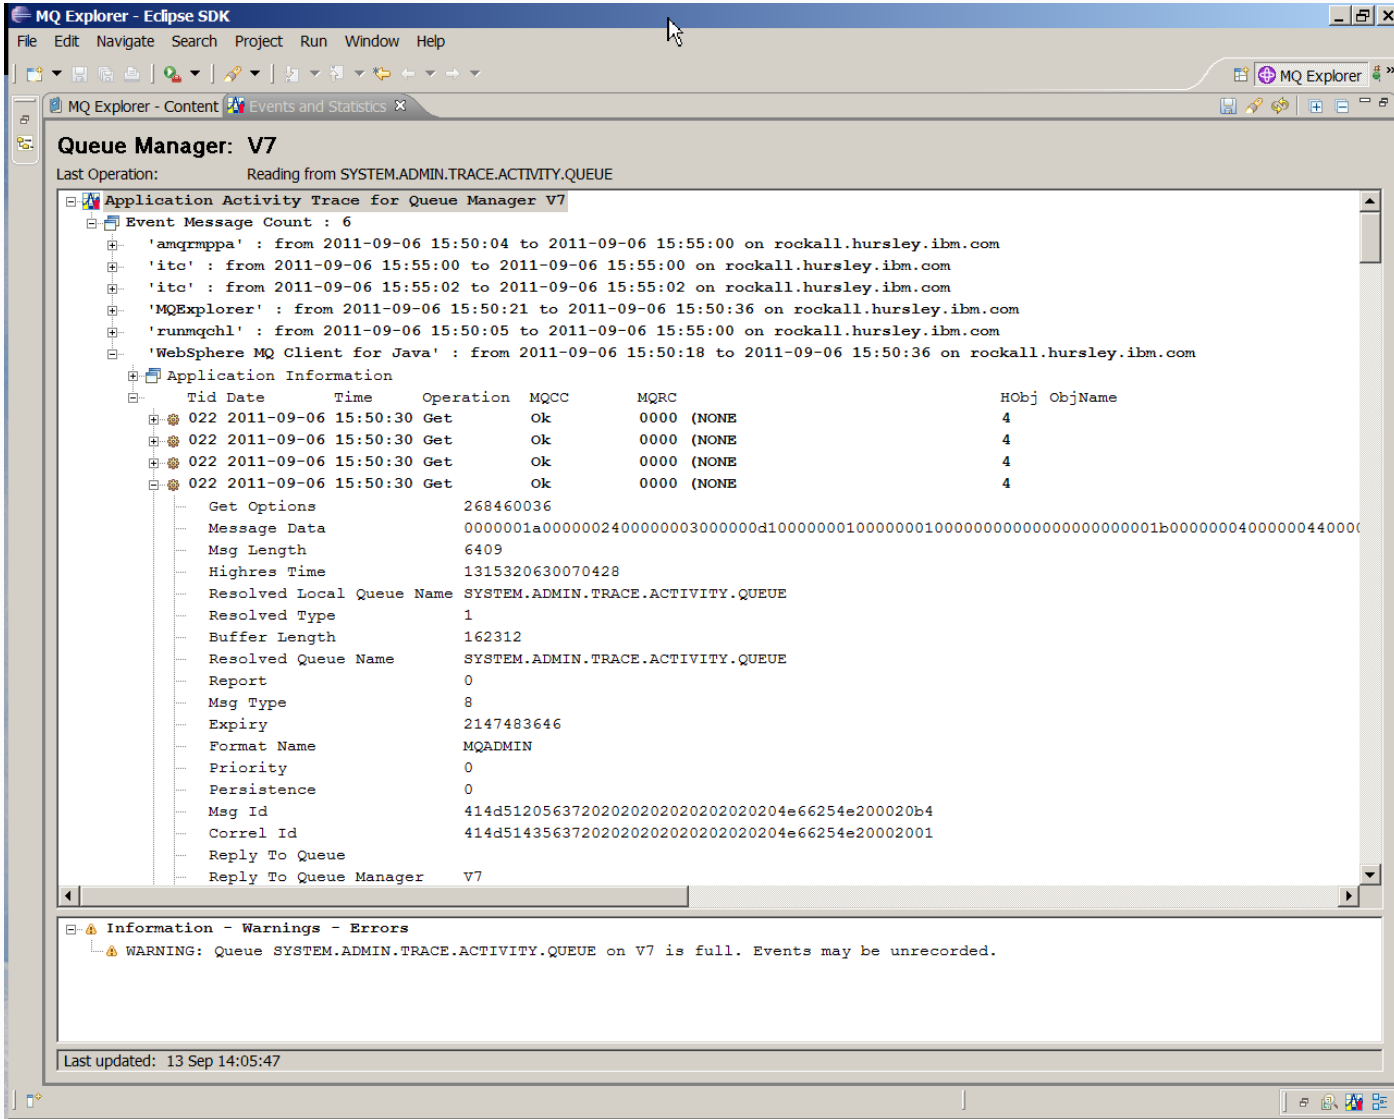
Auth Profile Name : X
Object Type : Queue
Principal Entity Names: metaylor
Auth Add Auths : Output
: Input

Third-party solutions

- Many vendor products – this screenshot from ITCAM/Omegamon

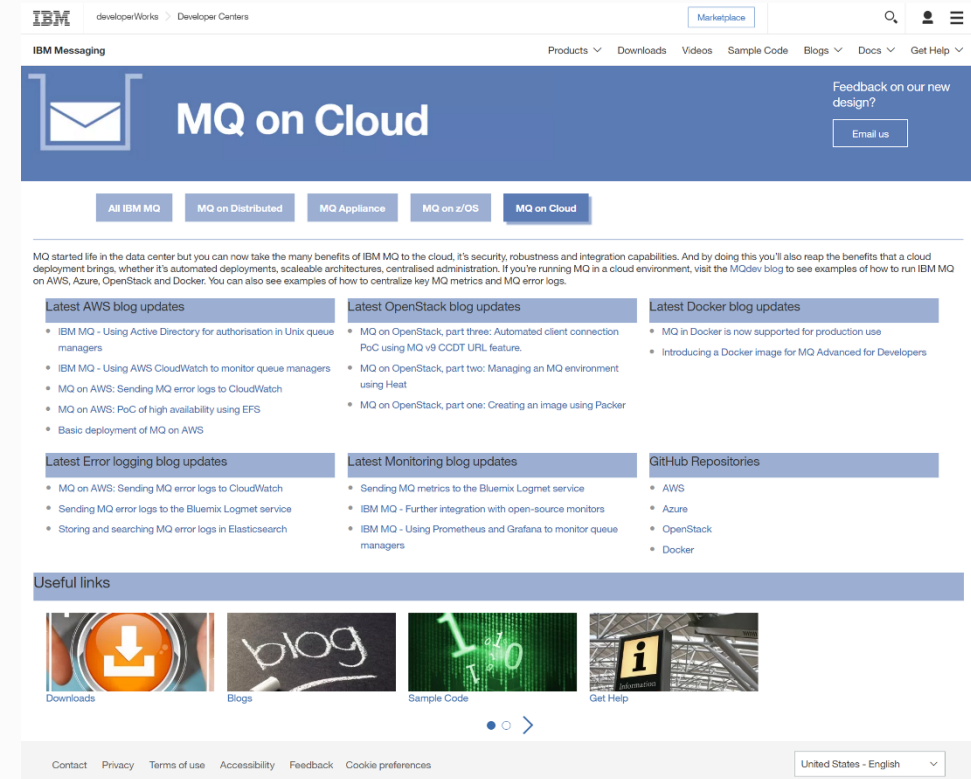


Application Activity inside MQ Explorer using MSOP



Many people now using different tools

- Because they are using those tools for other products
- And because MQ is being used in more environments
- Therefore MQ has to be able to be integrated with them
- You are unlikely to buy Tivoli if other parts of your infrastructure are being monitored via Grafana



<https://developer.ibm.com/messaging/mq-on-cloud/>



Decided to demonstrate MQ monitoring integration

- Using the V9 resource statistics data
- Feeding a variety of monitoring tools
- And doing it in public – Github, blog articles etc
 - See github.com/ibm-messaging/mq-golang
 - Video at youtube.com/watch?v=Pi_jHCiqTgU
- Other integration aspects – availability, security, deployment – also demonstrated

System Monitoring with V9

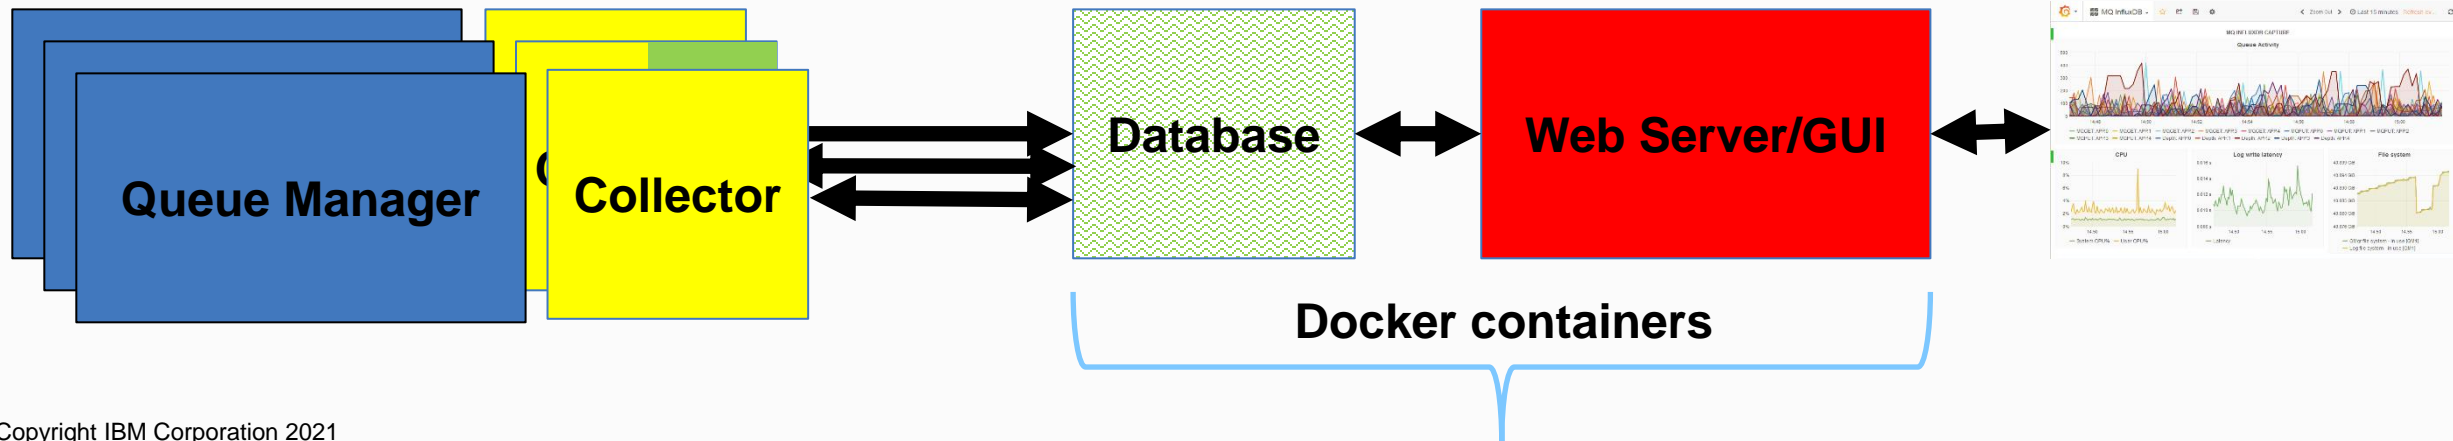
- More statistics available via a pub/sub model
- Includes CPU and Disk usage
 - As well as many MQ statistics
 - Not full replacement for accounting/statistics events but many key values
- Subscribe to meta-topic to learn which classes of statistics are available
 - **\$SYS/MQ/INFO/QMGR/<qmgr>/Monitor/METADATA/CLASSES**
 - Then subscribe to specific topics
 - See amqsrua sample program
- Distributed platforms only
- User applications can generate their own monitoring data in this style
 - The MQ Bridge to Salesforce contributes statistics

System Monitoring Example

```
$ amqsrua -m V9000_A
CPU : Platform central processing units
DISK : Platform persistent data stores
STATMQI : API usage statistics
STATQ : API per-queue usage statistics
Enter Class selection
==> CPU
SystemSummary : CPU performance - platform wide
QMgrSummary : CPU performance - running queue manager
Enter Type selection
==> SystemSummary
Publication received PutDate:20160411 PutTime:10465573
User CPU time percentage 0.01%
System CPU time percentage 1.30%
CPU load - one minute average 8.00
CPU load - five minute average 7.50
CPU load - fifteen minute average 7.30
RAM free percentage 2.02%
RAM total bytes 8192MB
```

Monitoring Architecture

- Architecture is split – database and user interface
 - The database is usually a "time-series" DB, not traditional SQL
 - Designed and optimised for {timestamp, metric, value} storage and queries
- These databases include Prometheus, InfluxDB, OpenTSDB
- Collection architecture may have intermediate layers – collectd



Started with Prometheus

- Seemed to be one of the most popular
- Which does have its own limited GUI
- Model is "pull" – calls a collector program at intervals via http
 - Most other DBs are "push" where collector sends to DB at interval
- Standard API for getting data to Prometheus is in Go
 - And we had no Go API for MQ ...

The Go API for MQ

- So first off, I had to create a new language binding
 - Based on full MQI rather than a "simplified" version
 - But not all function implemented
 - Trying to make it look natural to Go programmers

```
if err == nil {  
    putmqmd := ibmmq.NewMQMD()  
    pmo      := ibmmq.NewMQPMO()  
    pmo.Options = ibmmq.MQPMO_SYNCPOINT | ibmmq.MQPMO_NEW_MSG_ID  
    putmqmd.Format = "MQSTR"  
    msgData := "Hello from Go"  
    buffer := []byte(msgData)  
    err = qObject.Put(putmqmd, pmo, buffer)  
    if err != nil {  
        fmt.Println(err)  
    }  
}
```

Collector configurations

- Collector subscribes to all data for qmgr (cpu, disk etc) and nominated queues
 - Command line parameters name the queues with wildcards
- Started via MQ Service definition and shell script
- Can connect as client to remote queue managers including MQ appliance
 - Any system that supports the resource statistics
 - One collector instance per queue manager

```
/usr/local/bin/mqgo/mq_prometheus -ibmmq.queueManager=QM1  
-ibmmq.monitoredQueues=APP.*,MYQ.*  
-ibmmq.httpListenPort=9157  
-log.level=error
```

Grafana

- Although Prometheus has a GUI it is not very sophisticated
- Instead, prefer to use Grafana as visualisation tool
 - Supports many different backend databases
 - Understands the metric names, query capabilities etc of each

Add data source

Config Dashboards

Name My data source name ⓘ Default ☐

Type Prometheus

Http settings

Url http://localhost:9090 ⓘ

Access proxy ⓘ

Http Auth Basic Auth ☐ With Credentials ☐

Add data source

Name My data source name ⓘ Default ☐

Type CloudWatch

CloudWatch details

Credentials profile name default ⓘ

Default Region ⓘ

Custom Metrics namespace Namespace1, Namespace2 ⓘ

Assume Role ARN arn:aws:iam:* ⓘ

Add data source

Config Dashboards

Name My data source name ⓘ Default ☐

Type Graphite

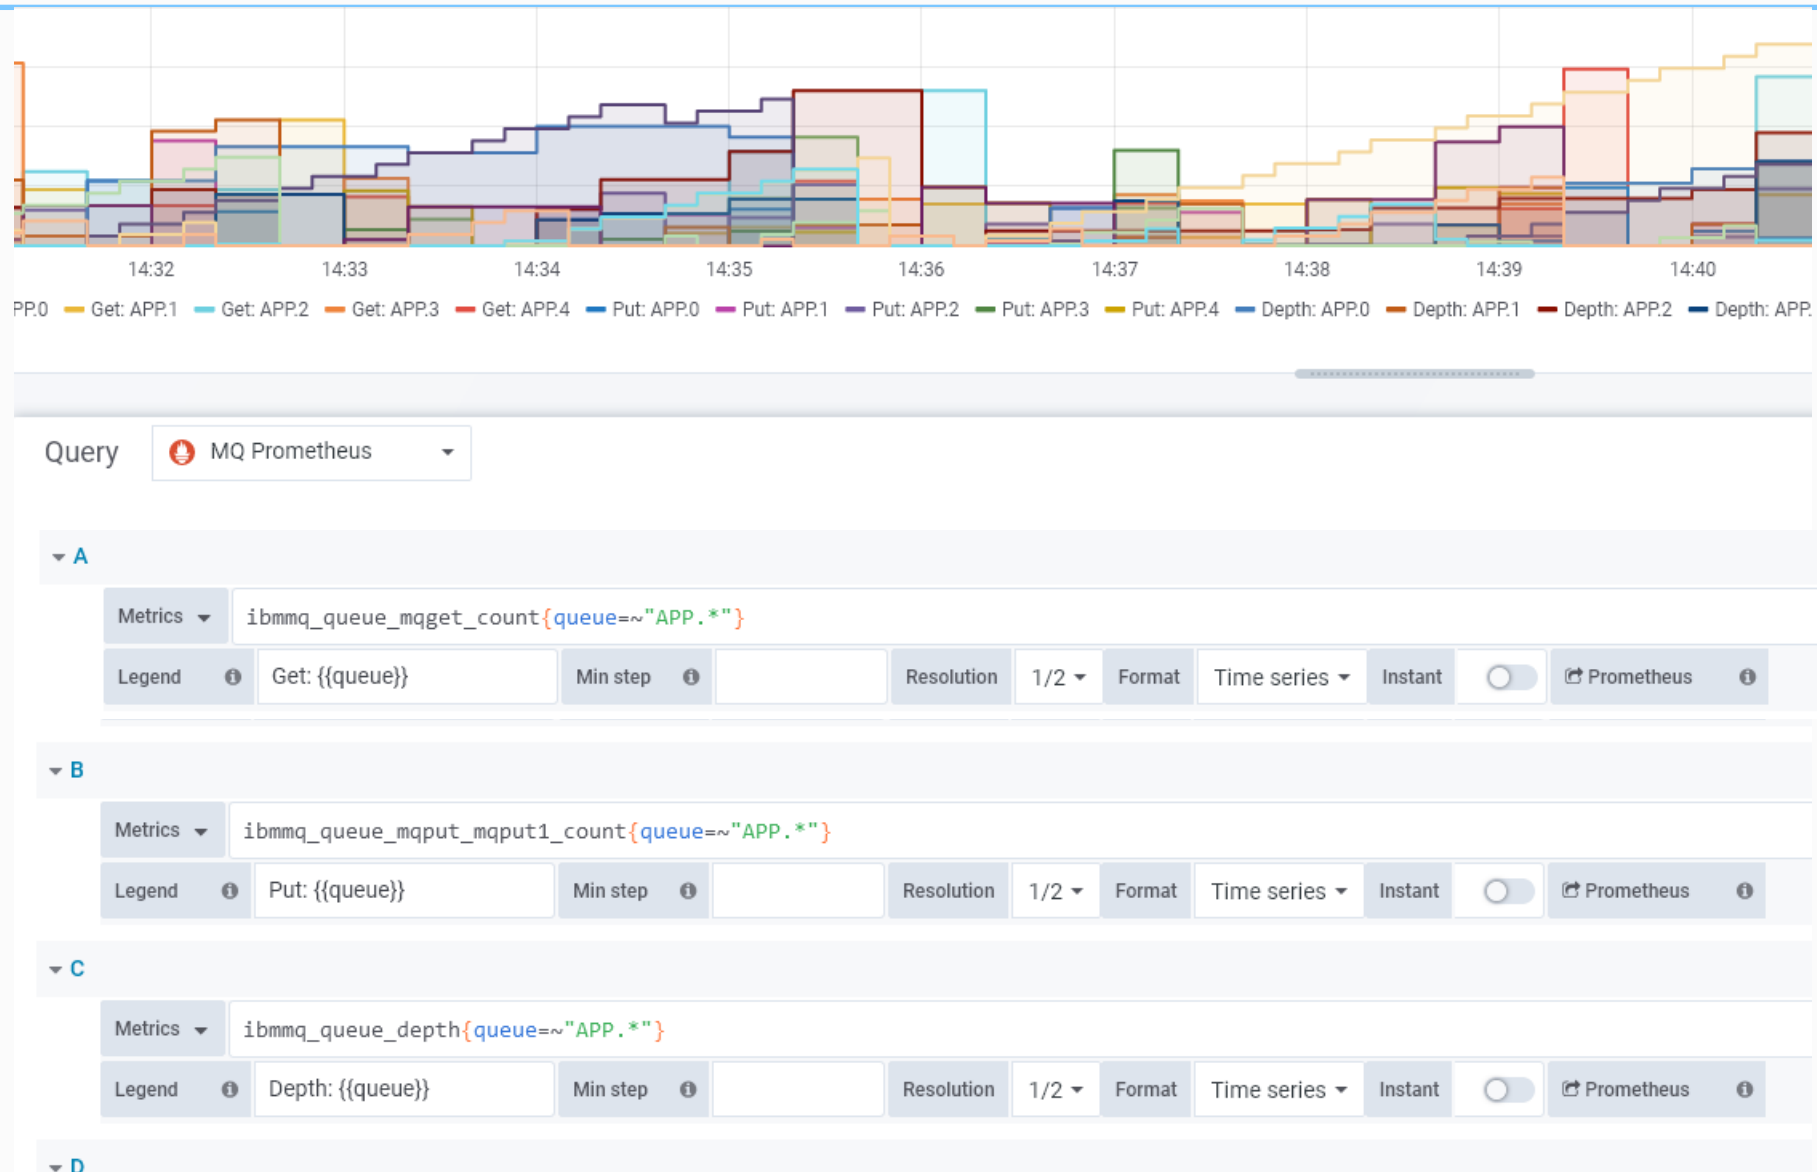
Http settings

Url http://localhost:8080 ⓘ


Access proxy ⓘ

Http Auth Basic Auth ☐ With Credentials ☐

Accessing queue stats from Prometheus in Grafana



Grafana dashboard

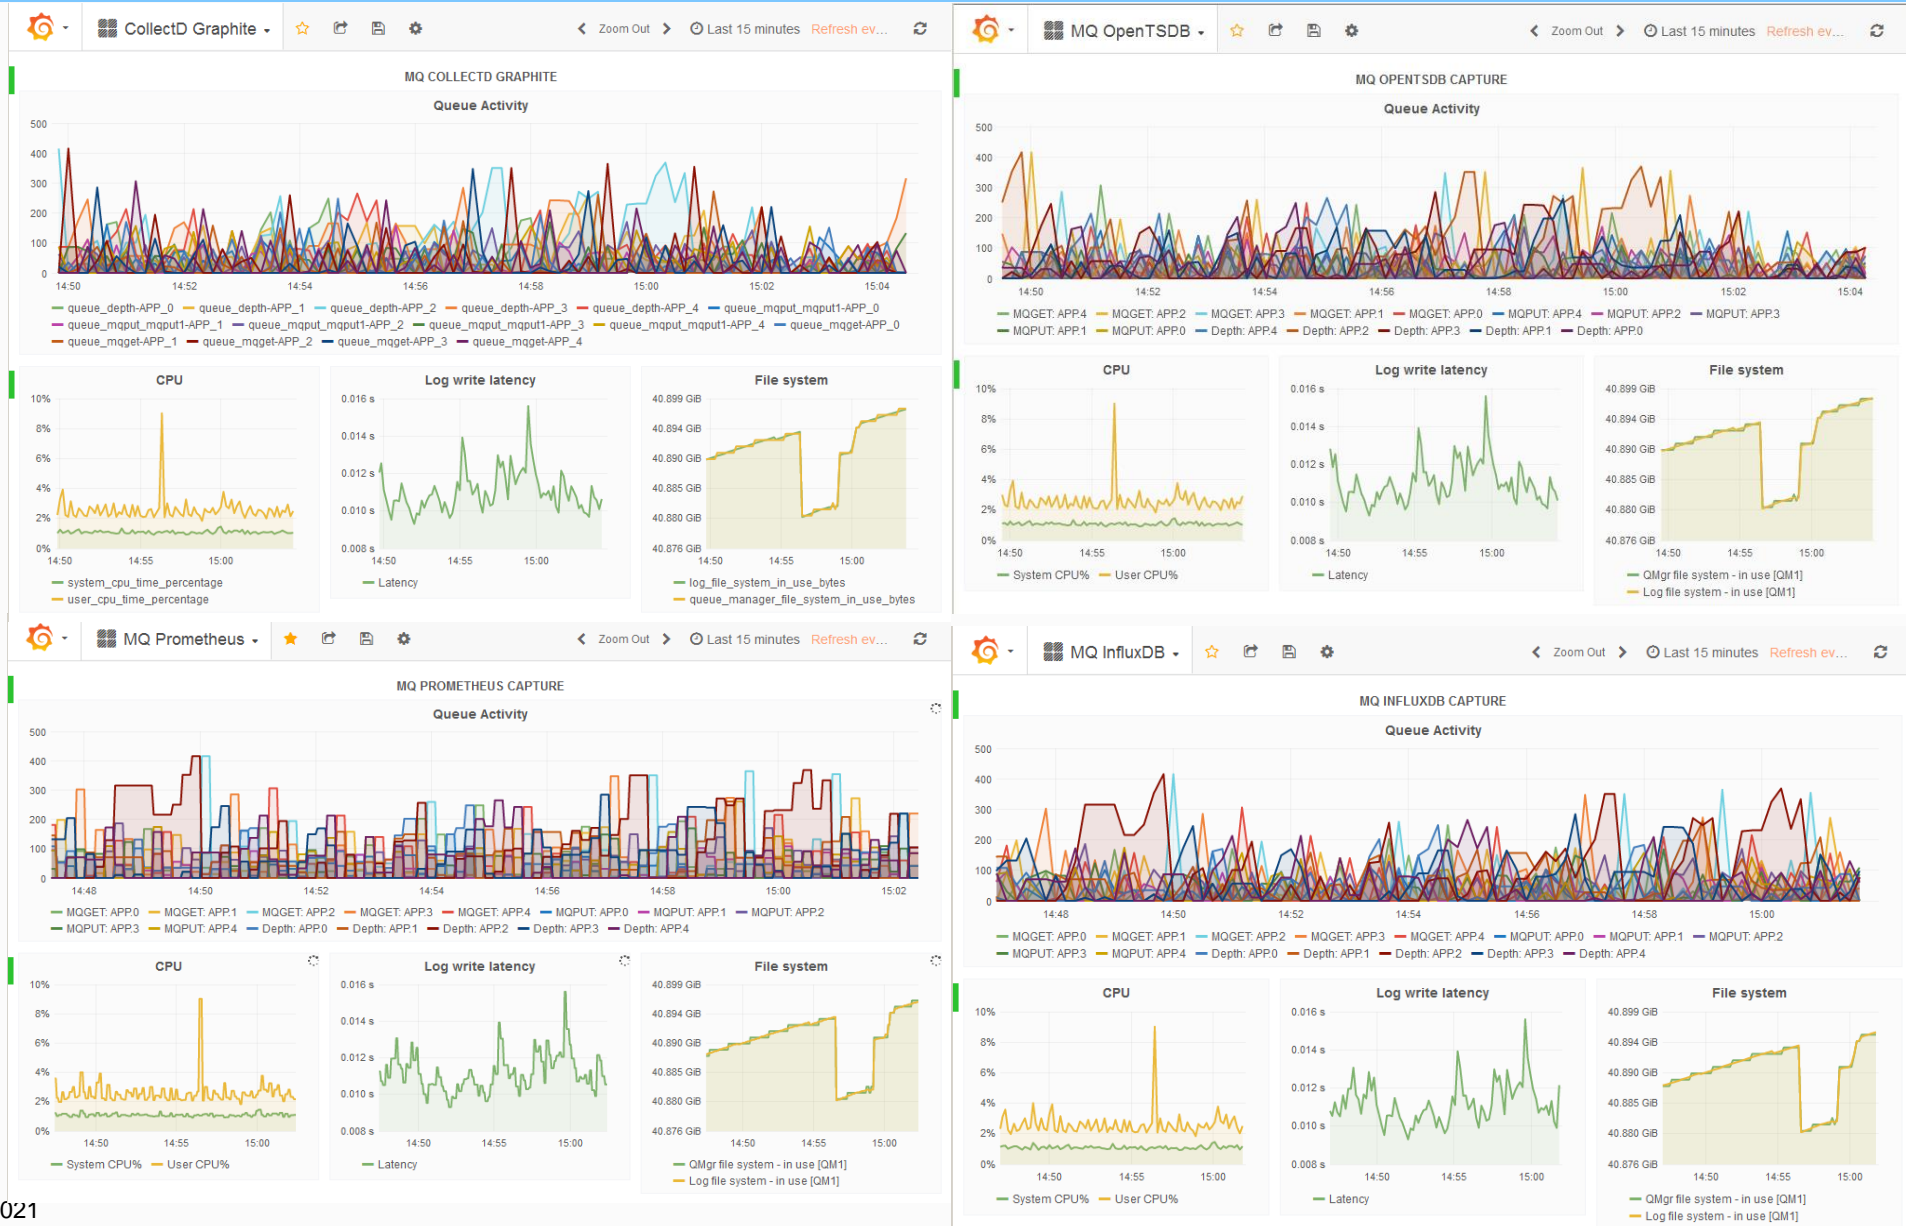


Then added more variants

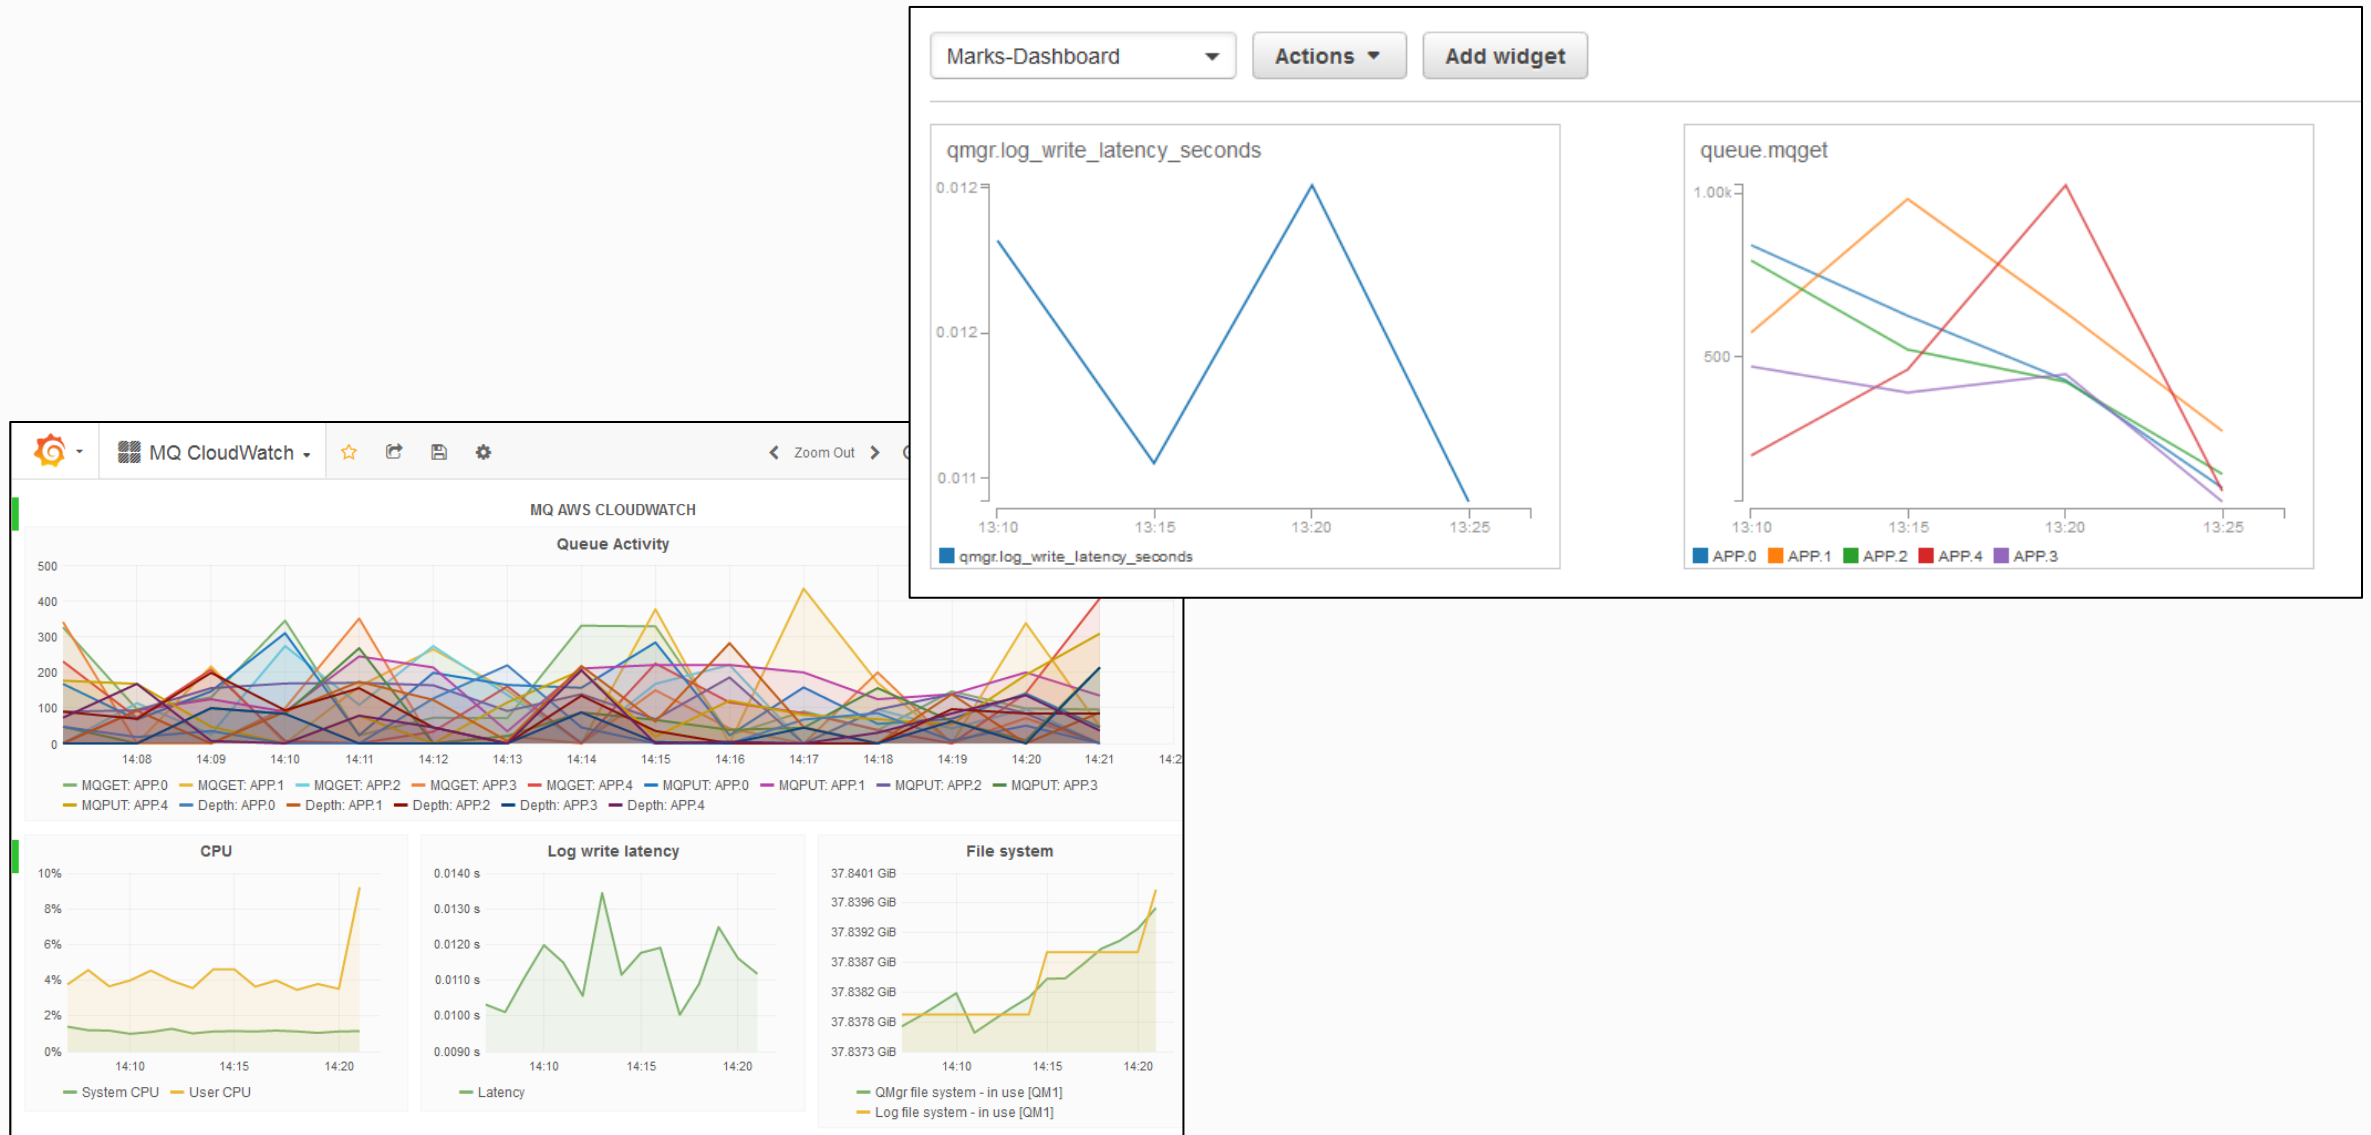
- Rapidly added support for Influx, OpenTSDB
 - Different collectors with slightly different parameters
- Graphite is another database, but fed via collectd
- Also added an AWS collector for CloudWatch
- Generic JSON formatting

```
{ "collectionTime" : {  
    "timeStamp" : "2016-11-07-T15:00:55Z"  
    "epoch" : 1478527255    },  
  "points" : [  
    { "queueManager" : "QM1", "ramTotalBytes" : 15515735206 },  
    { "queueManager" : "QM1", "userCpuTimePercentage" : 1.33 }  
  ]  
}
```

Four equivalent Grafana dashboards



AWS Cloudwatch



Requirement: Monitor channels


Monitoring - channels

Channels ▾

Channel Status

Channel Name ▲	Connection Name	Local QMgr	Remote QMgr	Type	State	Status
TO.QM1	127.0.0.1	QM1	QM2	RECEIVER	Running	3
TO.QM2	127.0.0.1(1415)	QM1	QM2	SENDER	Running	3
TO.QM2.2	localhost(1415)	QM1		SENDER	Transition	5

Channel Messages




Requirement: Monitor z/OS

Monitoring for z/OS



z/OS status - QML3

Last 15 minutes Refresh every 10s



Channel Status

Channel Name	Connection Name	Local QMgr	Type	State	Status
SYSTEM.ADMIN.SVRCONN	192.168.0.121	QML3	SVRCONN	Running	3
SYSTEM.ADMIN.SVRCONN	192.168.0.126	QML3	SVRCONN	Running	3
SYSTEM.ADMIN.SVRCONN	192.168.0.141	QML3	SVRCONN	Running	3
SYSTEM.ADMIN.SVRCONN	192.168.0.99	QML3	SVRCONN	Running	3
SYSTEM.ADMIN.SVRCONN	192.168.215.140	QML3	SVRCONN	Running	3
SYSTEM.ADMIN.SVRCONN	192.168.215.204	QML3	SVRCONN	Running	3






z/OS


z14

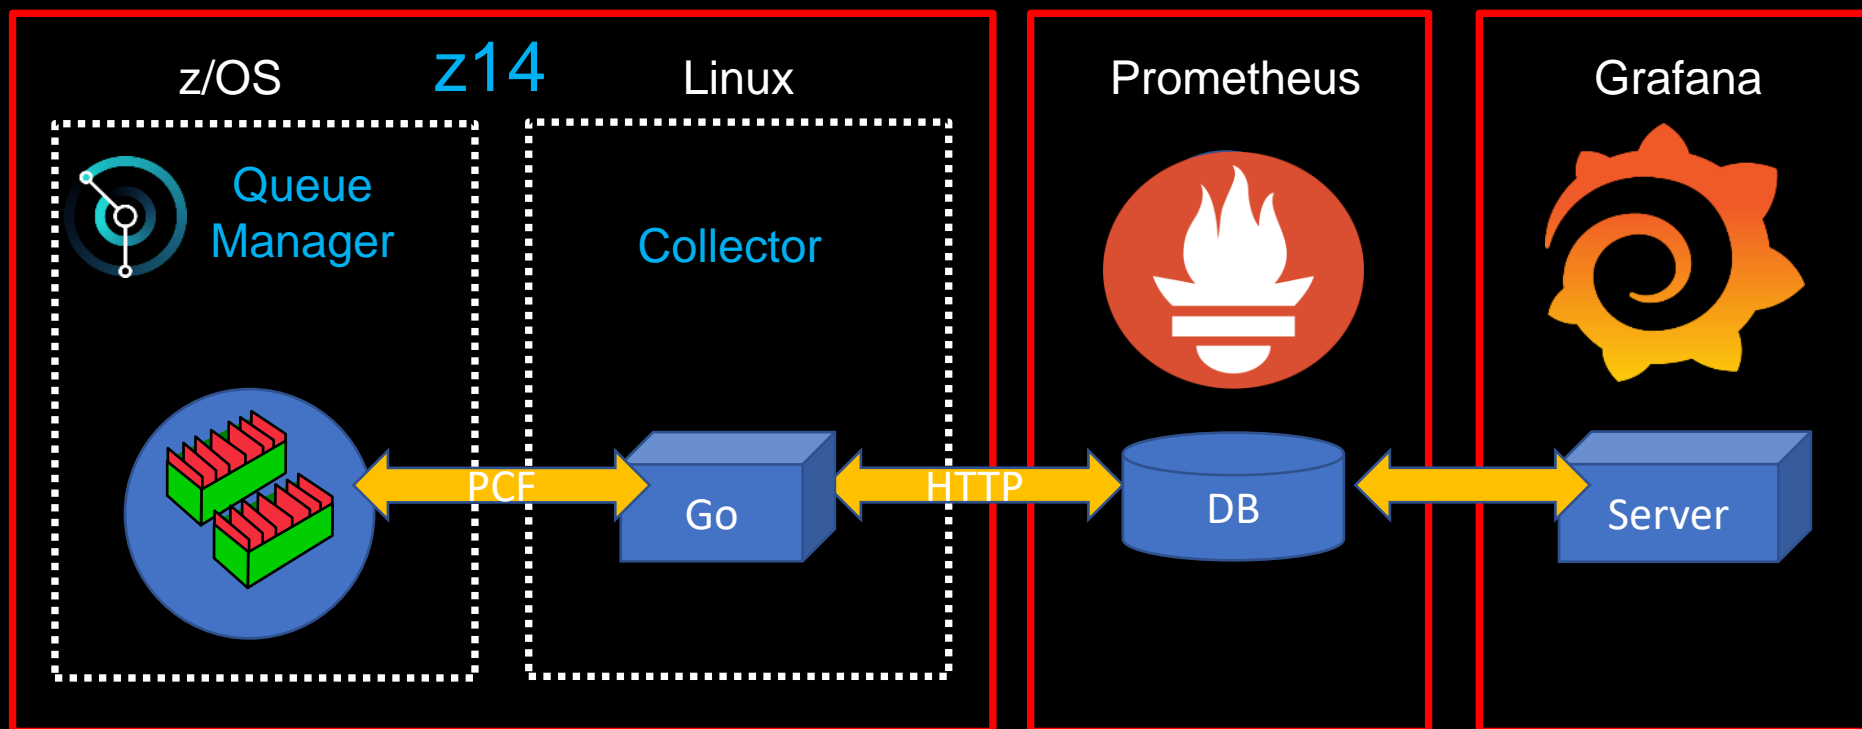


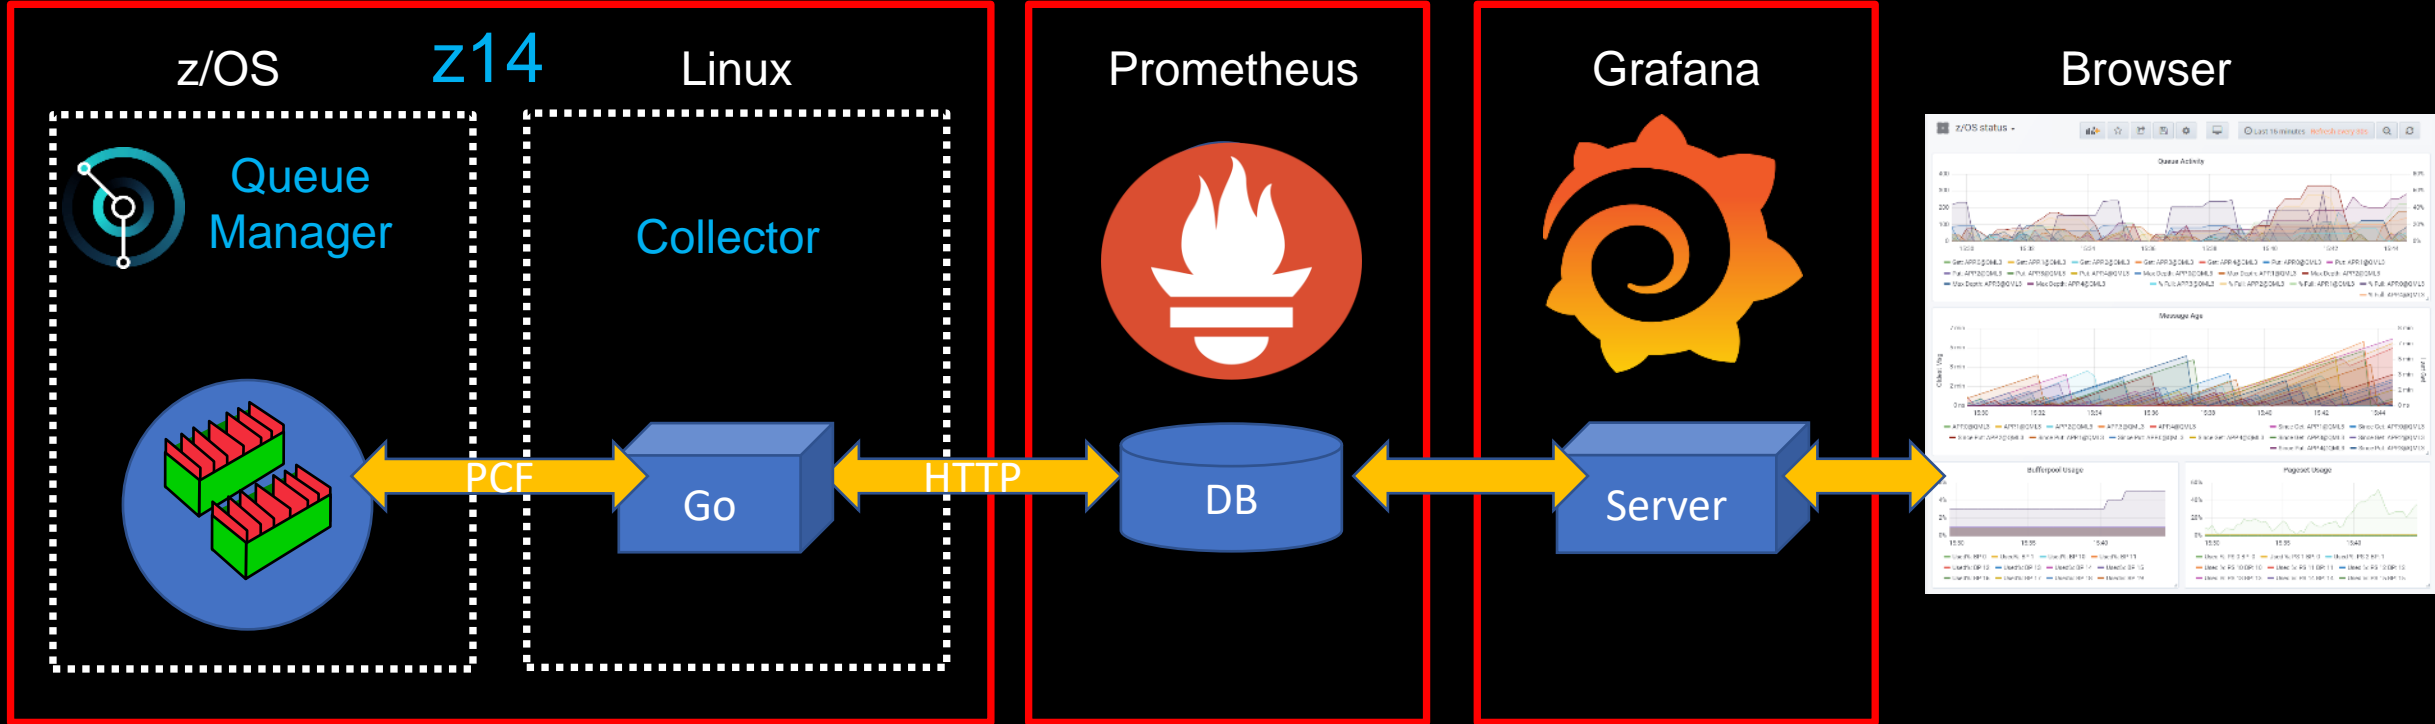
Queue
Manager














Queue Activity




Message Age



Bufferpool Usage




Pageset Usage




Channel Status

Channel Name	Connection Name	Local QMgr	Type	State	Status
SYSTEM.ADMIN.SVRCONN	192.168.0.254	QML3	SVRCONN	Running	3
SYSTEM.ADMIN.SVRCONN	192.168.17.254	QML3	SVRCONN	Running	3
TO.QML5	192.168.17.252	QML3	CLUSSDR	Transition	5

Channel Messages



Unique Channels



Unique Channels

46

BP/PS Usage

Bufferpool	Page Set	Queue Manager	PS Used	BP Used
3	4	QML3	35.6%	5.0%
0	0	QML3	1.6%	1.0%
0	1	QML3	1.0%	1.0%
2	3	QML3	0.3%	1.0%
1	2	QML3	0.2%	1.0%

No longer reliant on resource publications

- The monitors now extract information from STATUS commands
 - Queues, Channels, QMgr, Topics, Subscriptions
- And USAGE for Bufferpools and Pagesets
- Giving key information for ALL platforms including z/OS

June 2020 update

- All collectors brought to same functional level
- Can use YAML file as alternative to command line options

```
prometheus:
  port: 9157
  metricsPath: "/metrics"
  namespace: ibmmq
global:
  useObjectStatus: true
  useResetQStats: false
  logLevel: INFO
  rediscoverInterval: 1h
connection:
  queueManager: QM1
  replyQueue: SYSTEM.DEFAULT.MODEL.QUEUE
objects:
  queues:
    - APP.*
    - "!SYSTEM.*"
  queueSubscriptionSelector:
    - PUT
    - GET
```

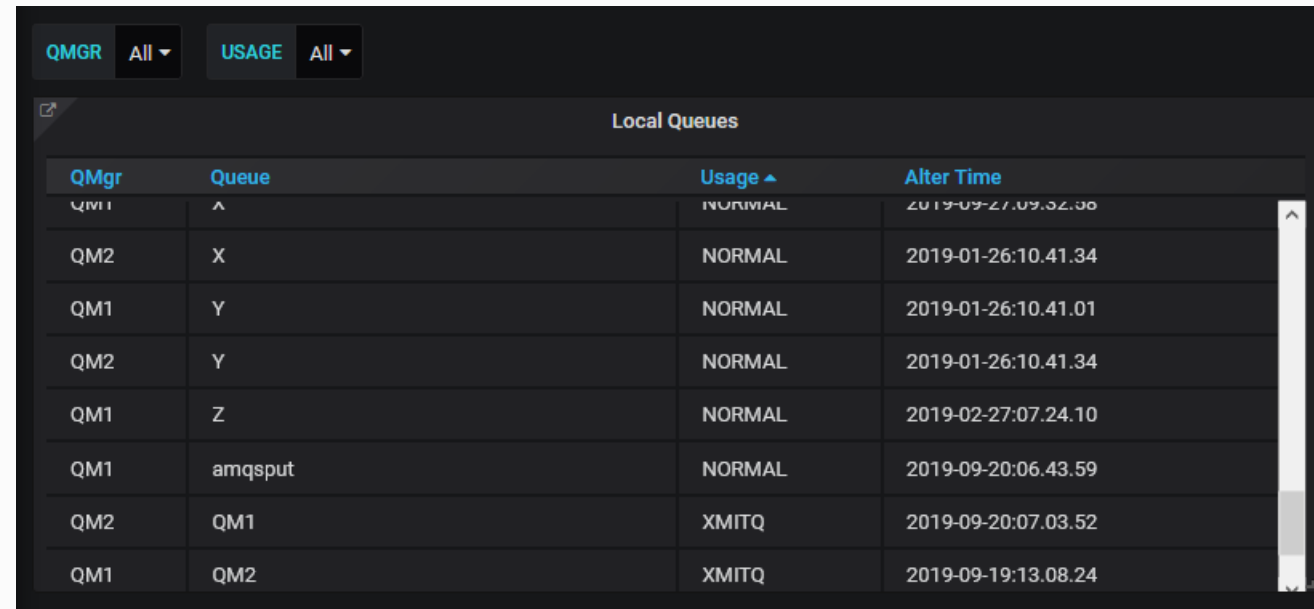
Additional info

- GitHub source: <https://github.com/ibm-messaging/mq-metric-samples>
- z/OS monitoring video (2019): <https://youtu.be/dCTTZWh5NDw>
- Original features video (2016): https://youtu.be/Pi_jHCiqTgU

Connecting metrics to configuration

- Grafana can also display information from some SQL databases
- It is possible to populate tables with MQ configuration information
 - For example, which queues are associated with particular bufferpools and pagesets
- The MQ REST API gets responses in JSON to be directly inserted
 - Databases are able to do SELECT on JSON-formatted columns

See <https://marketaylor.synology.me/?p=532>



The screenshot shows a web interface for monitoring MQGR USAGE. It features a table titled 'Local Queues' with four columns: 'QMGr', 'Queue', 'Usage', and 'Alter Time'. The table contains eight rows of data. The first row is partially obscured by a header row. The 'Usage' column shows values like 'NORMAL', 'XMITQ', and 'XMITQ'. The 'Alter Time' column shows timestamps in YYYY-MM-DD:HH:MM:SS format.

QMGr	Queue	Usage	Alter Time
QM1	^	NORMAL	2019-09-27:09:52.30
QM2	X	NORMAL	2019-01-26:10:41.34
QM1	Y	NORMAL	2019-01-26:10:41.01
QM2	Y	NORMAL	2019-01-26:10:41.34
QM1	Z	NORMAL	2019-02-27:07:24.10
QM1	amqsput	NORMAL	2019-09-20:06:43.59
QM2	QM1	XMITQ	2019-09-20:07:03.52
QM1	QM2	XMITQ	2019-09-19:13:08.24

QMGR

All ▾

USAGE

XMITQ ▾

Local Queues

QMgr	Queue	Usage ▴	Alter Time
QM2	QM1	XMITQ	2019-09-20:07.03.52
QM1	QM2	XMITQ	2019-09-19:13.08.24
QM1	QM2.RETRY	XMITQ	2019-09-30:14.49.41
QM1	QM2.STOPPED	XMITQ	2019-06-11:09.53.22

Query

MySQL ▾

▼ A




```
SELECT doc->>"$.queueManager" as 'QMgr',
       doc->>"$.a.queue" as 'Queue',
       doc->>"$.a.usage" as 'Usage',
       concat(doc->>"$.a.altdate", ":", doc->>"$.a.alttime") as "Alter Time"
FROM MQCFG.QLOCAL
WHERE doc->>"$.a.queue" NOT LIKE 'SYSTEM.%' AND (
  CASE WHEN $USAGE = 'ALL' THEN
    doc->>"$.a.usage" LIKE '%'
  ELSE
    doc->>"$.a.usage" = $USAGE
  END \ AND /
```

Format as

Table ▾

Show Help ▶

Generated SQL ▶



Adding resource statistics to your own applications

- Article showing how to publish similar statistics from your own applications
 - And therefore have monitors such as these showing status
 - Even if your apps are connecting to a z/OS queue manager
- Based on the MQ Salesforce Bridge code
 - Shows how to construct the PCF metadata describing your resources
- See <https://marketaylor.synology.me/?p=380>

What are differences? Which is best?

- Differences are generally in
 - The names and formats of metrics ("ibmmq_queue_mqget")
 - Naming for individual resources such as the queue name
 - Query capabilities to select and display chosen metrics
 - Can you use wildcards on object names
 - Creating labels on graphs
 - Can it be automatic based on the query?
 - Alerting capabilities
- The best is going to be whatever you are already using!
 - But I found the Prometheus/Grafana combination to be flexible and usable
- No easy way to report as string (eg "STARTED", "STOPPED" status)
 - Have to do a mapping via an integer or label

Squashing Metrics

Type	State	Status
RECEIVER	Running	3
SENDER	Running	3
SENDER	Transition	5

ibmmq_channel_status_squash+0	
Legend format	legend format
ibmmq_channel_status+0	
Legend format	legend format

Value #A	Options	Type	Value Mappings	Thresholds
channel	Apply to columns named Value #A	Type String	Type Value to text	Thresholds 1,2
qmgr	Column Header State	Sanitize HTML	0 Stopped	Color Mode Cell
Value #B	Render value as link	Preserve Formatting	1 Transition	Colors Invert
Value #C			2 Running	
type			+	

Latest Go/Prometheus status

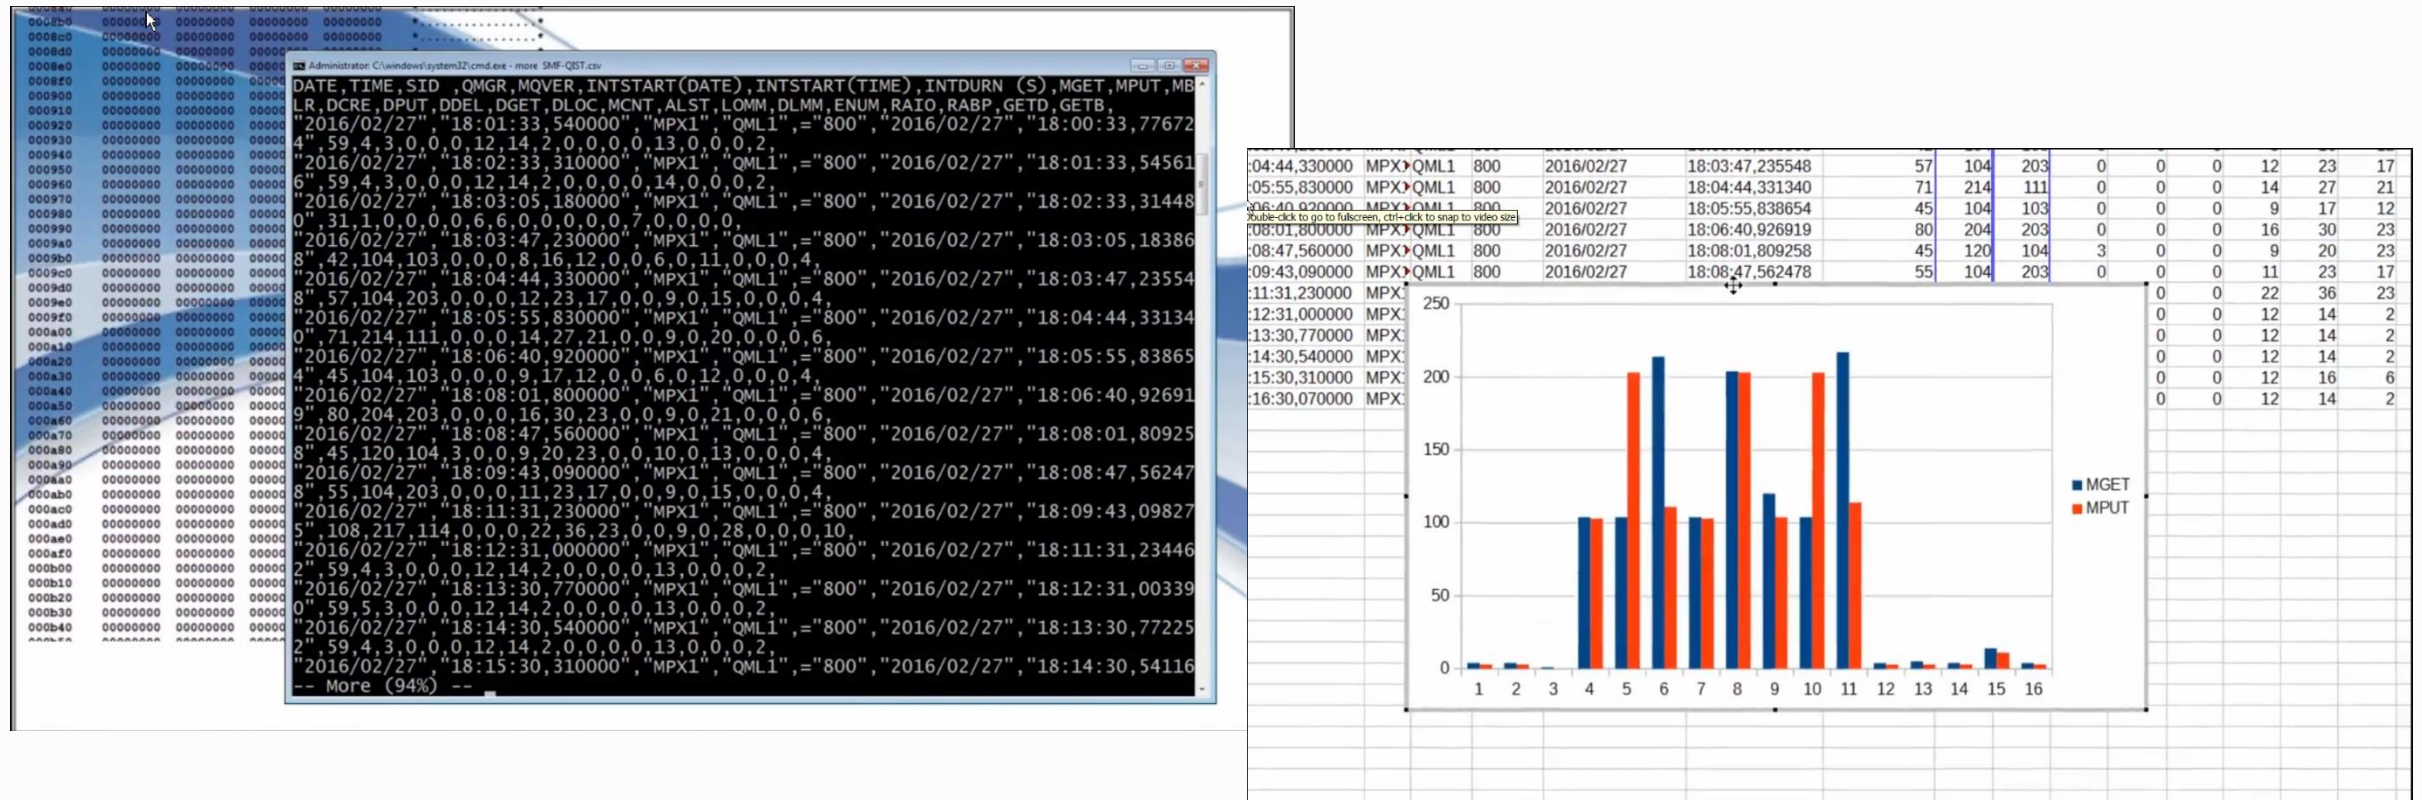
- All repositories under <https://github.com/ibm-messaging>
- **mq-golang** has the core MQI and PCF packages
 - Some sample code to demonstrate use of most functions
 - Assumes you already know the MQI principles from another language
- **mq-metric-samples** has Prometheus, Cloudwatch etc monitor programs
 - Along with a "vendor" tree
- Both have scripts to compile programs in Docker containers for ease of build

Formally supported metric collection & reporting

- The certified MQ container images include metrics reporting option
- Only some metrics are available: essentially, only the queue manager level
- The MQ managed service has same capability, with reporting through sysdig
- It is possible (though not generally recommended) to use both the supported container collector and my collectors
 - See blog post on adding Queue Depth monitoring to CP4I
 - <https://community.ibm.com/community/user/integration/blogs/matt-roberts1/2021/05/03/monitoring-mq-qdepth-cp4i>

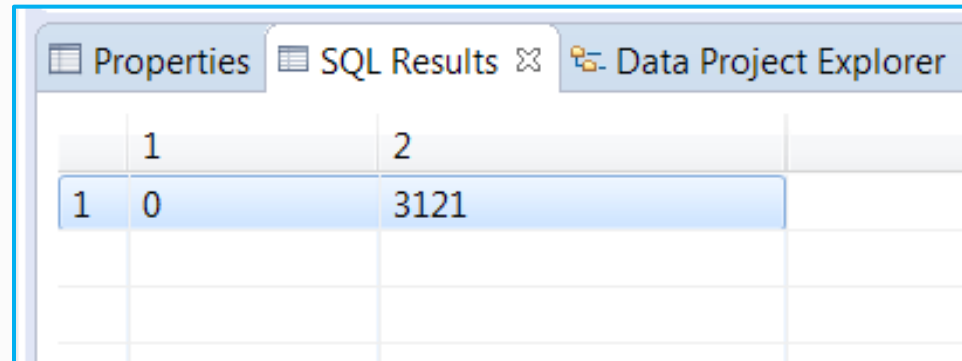
Similar resource data available on z/OS but via SMF

- mqsmfcsv ... open source tool to format MQ z/OS SMF records for easy import to spreadsheets and databases
 - <http://github.com/ibm-messaging/mq-smf-csv>



Example queries

- What was my largest message size retrieved for this queue?
 - `SELECT MAX(Get_Max_Msg_Size) from MQSMF.WQ where (Base_Name='LYNS.TEST.QUEUE');`
 - Result was 11,189 (application people insisted it was 3,800)
- How many MQPUTs and MQPUT1s were completed?
 - `SELECT SUM (Put_Count), SUM (Put1_Count) from MQSMF.WQ where (Base_Name = 'LYNS.TEST.QUEUE');`
 - Results:



The screenshot shows a SQL Results window with three tabs: Properties, SQL Results, and Data Project Explorer. The SQL Results tab is active, displaying a table with two columns and one row of data. The first column is labeled '1' and the second column is labeled '2'. The data row shows '0' in the first column and '3121' in the second column.

1	2
0	3121

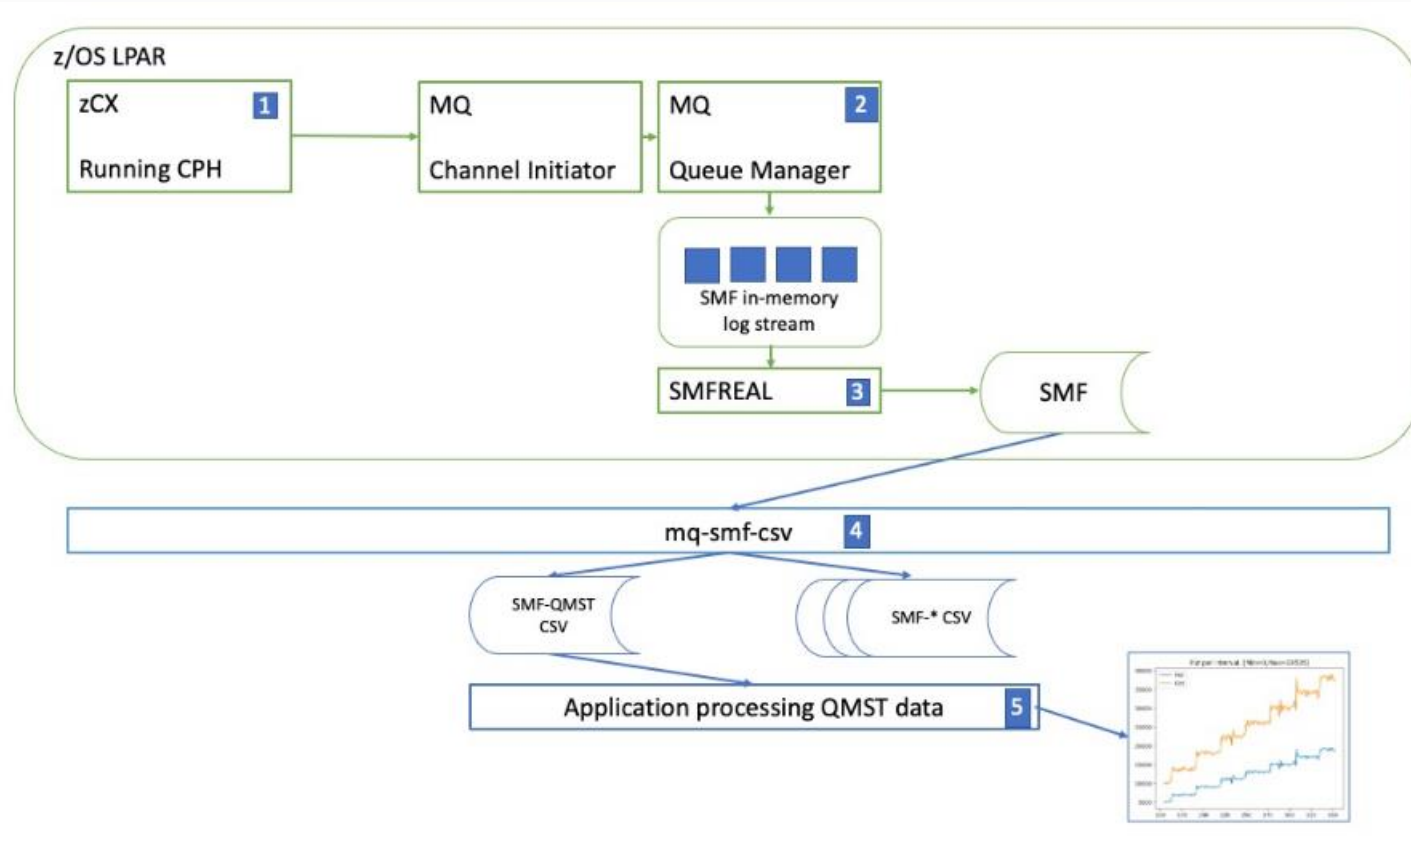
And we can now do it in JSON

- `mqsmfcsv -i <input file> -f json`

```
{
  "recordType" : 116,
  "recordSubType" : 0,
  "structure" : "QMAC",
  "date" : "2015/11/23",
  "time" : "11:00:00.020000",
  "lpar" : "H019",
  "qmgr" : "MQPC",
  "mqVersion" : "800",
  "authorisationId" : "IMS      ",
  "correlId" : "F0F2F3F6C2C3F1E4C4D6C340",
  "connectionName" : "PRDC      ",
  "operatorId" : "PLN1231 ",
  "applicationType" : "IMS MPP/BMP",
  "accountingToken" :
  "0000000000000000000000000000000000000000000000000000000000000000",
  "networkId" :
  "D7D9C4C340404040044E0A0800000001",
  ...
}
```

Real-time monitoring of SMF

- Using in-memory SMF log streams + mqsmfcsv



<https://community.ibm.com/community/user/integration/blogs/anthony-sharkey1/2021/11/30/whats-new-in-mq-for-zos-924-smf-enhancements?CommunityKey=183ec850-4947-49c8-9a2e-8e7c7fc46c64&tab=recentcommunityblogsdashboard>

Processing other MQ events

- Already shown amqsevt as shipped in MQ V8
- It now also supports JSON output option
 - Included from V9.1
- Can be used to feed JSON consumers such as splunk
- MQ 9.2.4 improved JSON output to create single-line messages
 - More convenient for some consumers; no need for additional processing
 - Trivial script enables events to be read via REST messaging API
 - See <https://marketaylor.synology.me/?p=959>

MQ events in splunk

The screenshot displays the Splunk web interface with the following components:

- Navigation Bar:** Search, Datasets, Reports, Alerts, Dashboards. The 'Alerts' tab is highlighted.
- Search Bar:** Contains the query `host=0b9f3995a92b "eventSource.objectName"="SYSTEM.ADMIN.PERFM.EVENT"`.
- Results Summary:** Shows 2 events for the time range 04/11/2016 12:10:34.000 to 04/11/2016 12:10:35.000. Includes a 'No Event Sampling' indicator.
- Event List:** A table with columns 'Time' and 'Event'. The first event is expanded, showing its JSON data.
- Fields List:** A sidebar on the left showing 'Selected Fields' (host, source, sourcetype) and 'Interesting Fields' (eventCreation, eventData.baseObjectName, eventData.highQueueDepth, eventData.msgDeqCount, eventData.msgEnqCount, eventData.queueMgrName, eventData.timeSinceReset, eventReason.name, eventReason.value, eventSource.objectName, eventSource.objectType, eventType.name, eventType.value, index, linecount, punct, splunk_server, timestamp).

Expanded Event Data (Event 1):

```
> 04/11/2016 12:10:34.000 {
  "eventSource" : { "objectName": "SYSTEM.ADMIN.PERFM.EVENT",
                    "objectType" : "Queue" },
  "eventType" : {
    "name" : "Perfm Event",
    "value" : 45
  },
  "eventReason" : {
    "name" : "Queue Full",
    "value" : 2053
  },
  "eventCreation" : "2016/11/04 12:10:24.29 GMT",
  "eventData" : {
    "queueMgrName" : "V9000_A",
    "baseObjectName" : "FULLEVT",
    "timeSinceReset" : 0,
    "highQueueDepth" : 4,
    "msgEnqCount" : 0,
    "msgDeqCount" : 0
  }
}
```

host = 0b9f3995a92b | source = /mqm/jsonevt.txt | sourcetype = _json

Using JSON event formatter with Activity Events

- Use the event formatter to output in JSON and then filter it further
 - Run via "service" if local qmgr
 - Could also use subscribe variant to obtain trace

```
amqsevt -m QM1 -q SYSTEM.ADMIN.TRACE.ACTIVITY.QUEUE -o json | jq -r -f jqFilt
```

- And then get one-line CSV output of key fields

```
"amqspu" ,"2018-07-11" ,"08:16:48" ,"Connx" ,0 ,"N/A"  
"amqspu" ,"2018-07-11" ,"08:16:48" ,"Open" ,0 ,"QL1"  
"amqspu" ,"2018-07-11" ,"08:16:48" ,"Put" ,0 ,"QL1" ,48 ,"414D512056393030305F4120202020205B2B77"  
"amqspu" ,"2018-07-11" ,"08:16:48" ,"Close" ,0 ,"QL1"  
"amqspu" ,"2018-07-11" ,"08:16:48" ,"Disc" ,0 ,"N/A"  
"amqsge" ,"2018-07-11" ,"08:16:48" ,"Connx" ,0 ,"N/A"  
"amqsge" ,"2018-07-11" ,"08:16:48" ,"Open" ,0 ,"QL1"  
"amqsge" ,"2018-07-11" ,"08:16:48" ,"Get" ,0 ,"QL1" ,38 ,"414D512056393030305F4120202020205B2B77"  
"amqsge" ,"2018-07-11" ,"08:16:48" ,"Get" ,2033 ,"QL1" ,250059 ,0  
"amqsge" ,"2018-07-11" ,"08:16:48" ,"Close" ,0 ,"QL1"  
"amqsge" ,"2018-07-11" ,"08:16:48" ,"Disc" ,0 ,"N/A"
```

<https://marketaylor.synology.me/?p=373>

A jq filter

```
select(.eventData.activityTrace != null) | .eventData.applName as $applName |
  (.eventData.activityTrace[] |
    [
      $applName, .operationDate, .operationTime, .operationId, .reasonCode.value,
      if (.objectName | length) > 0
      then
        .objectName
      else
        "N/A"
      end,
      if .operationId == "Get" or .operationId == "Put" or .operationId == "Put1"
      then
        .qmgrOpDuration, .msgId
      else
        empty
      end
    ]
  ) |
  @csv
```

Accounting records

```
{ "eventSource": { "objectName": "SYSTEM.ADMIN.ACCOUNTING.QUEUE", "objectType": "Queue" },
  "eventType": { "name": "Accounting MQI", "value": 167 }, "eventReason": { "name": "None", "value": 0 },
  "eventCreation": { "timeStamp": "2018-07-25T20:46:19Z", "epoch": 1532551579 },
  "eventData": { "queueMgrName": "V9000_A", "startDate": "2018-0725",
    "startTime": "21.46.16", "endDate": "2018-07-25", "endTime": "21.46.19", "commandLevel": 905,
    "connectionId": "414D514356393030305F4120202020205B2B779523E91D95", "sequenceNumber": 0,
    "applName": "amqsevt", "processId": 35389620, "threadId": 1, "userIdentifier": "metaylor", "connDate":
    "2018-07-25", "connTime": "21.46.16", "discDate": "2018-07-25", "discTime": "21.46.19",
    "opens": [0, 3, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
    "opensFailed": [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
    "closes": [0, 3, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
    "closesFailed": [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0],
    "puts": [24, 0],
    "putsFailed": 0,
    "put1s": [1, 0],
    "put1sFailed": 0,
    "putBytes": [2547972, 0],
    ...
```

Multiple consumers for MQ events

- Traditional MQ events (queue full etc) are put to a specific named queue
- Makes it difficult to have multiple consumers for same event queue
 - Many monitors can be configured to "browse" but who does "get" and when?
- The MQ event queues can be redefined as topic aliases
- Monitor programs can then get independently from their own dedicated queues
 - I might then run Omegamon AND the JSON variant of amqsevt to different consoles

```
DELETE QLOCAL (SYSTEM.ADMIN.CHANNEL.EVENT)      PURGE
DELETE QLOCAL (SYSTEM.ADMIN.PERFM.EVENT)         PURGE

DEFINE QALIAS (SYSTEM.ADMIN.CHANNEL.EVENT) TARGET (SYSTEM.ADMIN.EVENT)  TARGTYPE (TOPIC)
DEFINE QALIAS (SYSTEM.ADMIN.PERFM.EVENT)    TARGET (SYSTEM.ADMIN.EVENT)  TARGTYPE (TOPIC)

DEFINE TOPIC (SYSTEM.ADMIN.EVENT)  TOPICSTR ('SYSTEM/ADMIN/EVENT')
DEFINE QLOCAL (SYSTEM.ADMIN.SUBSCRIBED.EVENT)
DEFINE SUB (SYSTEM.ADMIN.EVENT)  TOPICOBJ (SYSTEM.ADMIN.EVENT)  +
  DEST (SYSTEM.ADMIN.SUBSCRIBED.EVENT)
```

MQ REST Administration

- Enabling further management options
 - Easy access from any language
 - Scriptable via curl
- Many MQSC commands have REST equivalent
 - Others supported via generic command
 - V9.1.3 adds true JSON-formatted generic commands
- Can manage older qmgrs via proxy qmgr

```
C:\Program Files\IBM\Latest902\bin>curl -k "https://localh
{"queue": [{
  "name": "Q.LOCAL",
  "status": {
    "currentDepth": 0,
    "lastGet": "",
    "lastPut": "",
    "mediaRecoveryLogExtent": "",
    "monitoringRate": "off",
    "oldestMessageAge": -1,
    "onQueueTime": {
      "longSamplePeriod": -1,
      "shortSamplePeriod": -1
    },
    "openInputCount": 0,
    "openOutputCount": 0,
    "uncommittedMessages": 0
  },
  "type": "local"
}]}
```

From V9.0.5 "What's New and Changed"

Version 9.0.5 introduces various improvements to the management and output of error logs. The main changes are that you can:

- Log diagnostic messages, using additional file services and syslog on UNIX platforms, as well as AMQERR01.LOG.

- Use JSON for the description of the messages, as well as the existing format; see JSON format diagnostic messages.

- Reformat a log into another language or style; see mqrc.

For more information, see Diagnostic message services, and QMErrorLog service.

https://www.ibm.com/support/knowledgecenter/en/SSFKSJ_9.0.0/com.ibm.mq.pro.doc/q130630.htm#q130630_errlog

JSON Error (aka DiagnosticMessage) Logs

- In same directory as classic error files
- Files AMQERRxx.json
- Unix systems can also direct entries to syslog
 - Which has a lot of backends and routing options

```
{
  "ibm_messageId": "AMQ5051I",
  "arith_insert_2": 1,
  "comment_insert_1": "LOGGER-IO",
  "ibm_datetime": "2017-11-16T09:54:26.331Z",
  "ibm_serverName": "QM1",
  "type": "mq_log",
  "host": "machine.somewhere.ibm.com",
  "loglevel": "INFO",
  "module": "amqzmut0.c:1650",
  "ibm_sequence": "1510826066_332014693",
  "ibm_processId": 7846,
  "ibm_threadId": 4,
  "ibm_version": "9.0.4.0",
  "ibm_processName": "amqzmuc0",
  "ibm_userName": "somebody",
  "ibm_installationName": "Installation3",
  "ibm_installationDir": "/opt/mqm",
  "message": "AMQ5051I: The queue manager task 'LOGGER-IO' has started."
}
```

Consuming tools

- See, for example, Elasticsearch module
 - <https://www.elastic.co/guide/en/beats/filebeat/master/filebeat-module-ibmmq.html#filebeat-module-ibmmq>

How to configure syslog with MQ 9.1

- This example from AIX. Other Unix platforms will be similar
- In /etc/syslog.conf

```
# MQ writes to the "user" facility
user.debug /var/mqm/errors/syslog.log rotate size 1m files 4 compress
```

- In queue manager's qm.ini


```
DiagnosticMessages:
  Name=DiagSyslog
  Service=Syslog
  Ident=mqseries
  Severities=I+
```

- Make sure syslog.log exists, then restart syslogd

```
Aug 14 15:58:50 example user:info mqseries: {"ibm_messageId":"AMQ9411I",
"ibm_arithInsert1":0, "ibm_arithInsert2":0,"ibm_datetime":"2018-08-14T14:58:50.250Z",
"ibm_serverName":"V9100_A", "type":"mq_log", "host":"example.hursley.ibm.com",
"loglevel":"INFO","module":"amqrrmfa.c:2108", "ibm_sequence":"1534258730_251676000",
"ibm_qmgrId":"V9100_A_2018-06-27_11.13.46", "ibm_version":"9.1.0.0", "ibm_processName":
"amqrrmfa", "ibm_userName":"metaylor", "ibm_installationDir":"/usr/mqm",
"message":"AMQ9411I: Repository manager ended normally."}
```

Summary

- MQ can be easily integrated with a variety of tools
- The pub/sub model for statistics makes it easy to add new consumers
 - Without disrupting any existing monitors
 - And makes it possible to add your own producers
- Using github for repository of code enables easy modification and sharing
github.com/ibm-messaging
- And blog posts for documenting what we have done
- Ability to use JSON as a common format for all operations



Any questions?